



# EECE 4353 Image Processing

Lecture Notes: Light Scattering, Image  
Formation, Digital Images, and MATLAB

Richard Alan Peters II

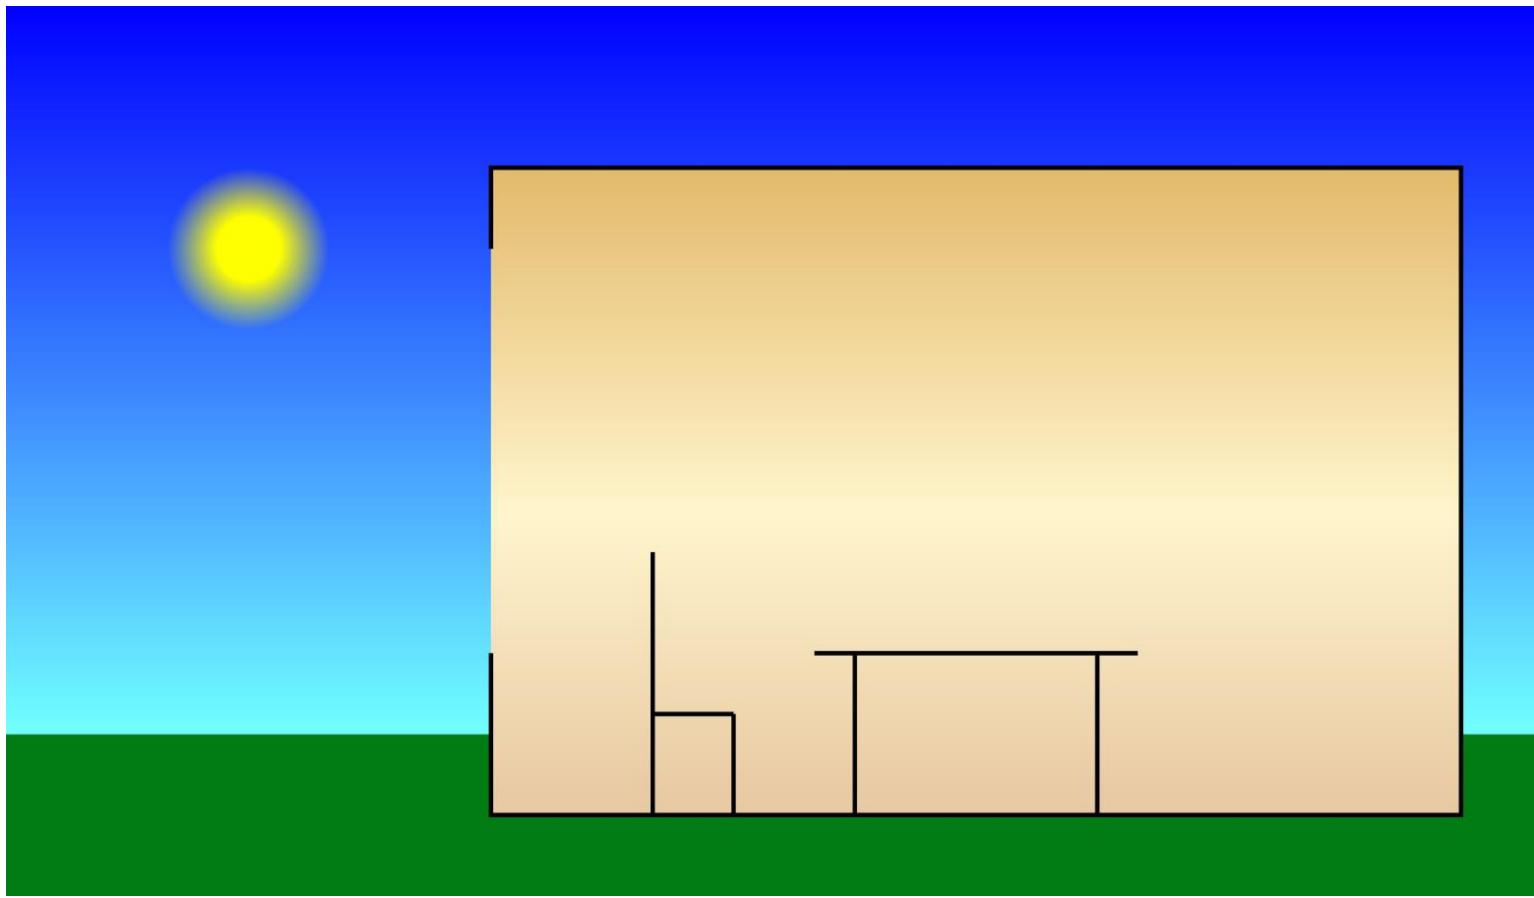
Department of Electrical Engineering and  
Computer Science

Fall Semester 2016



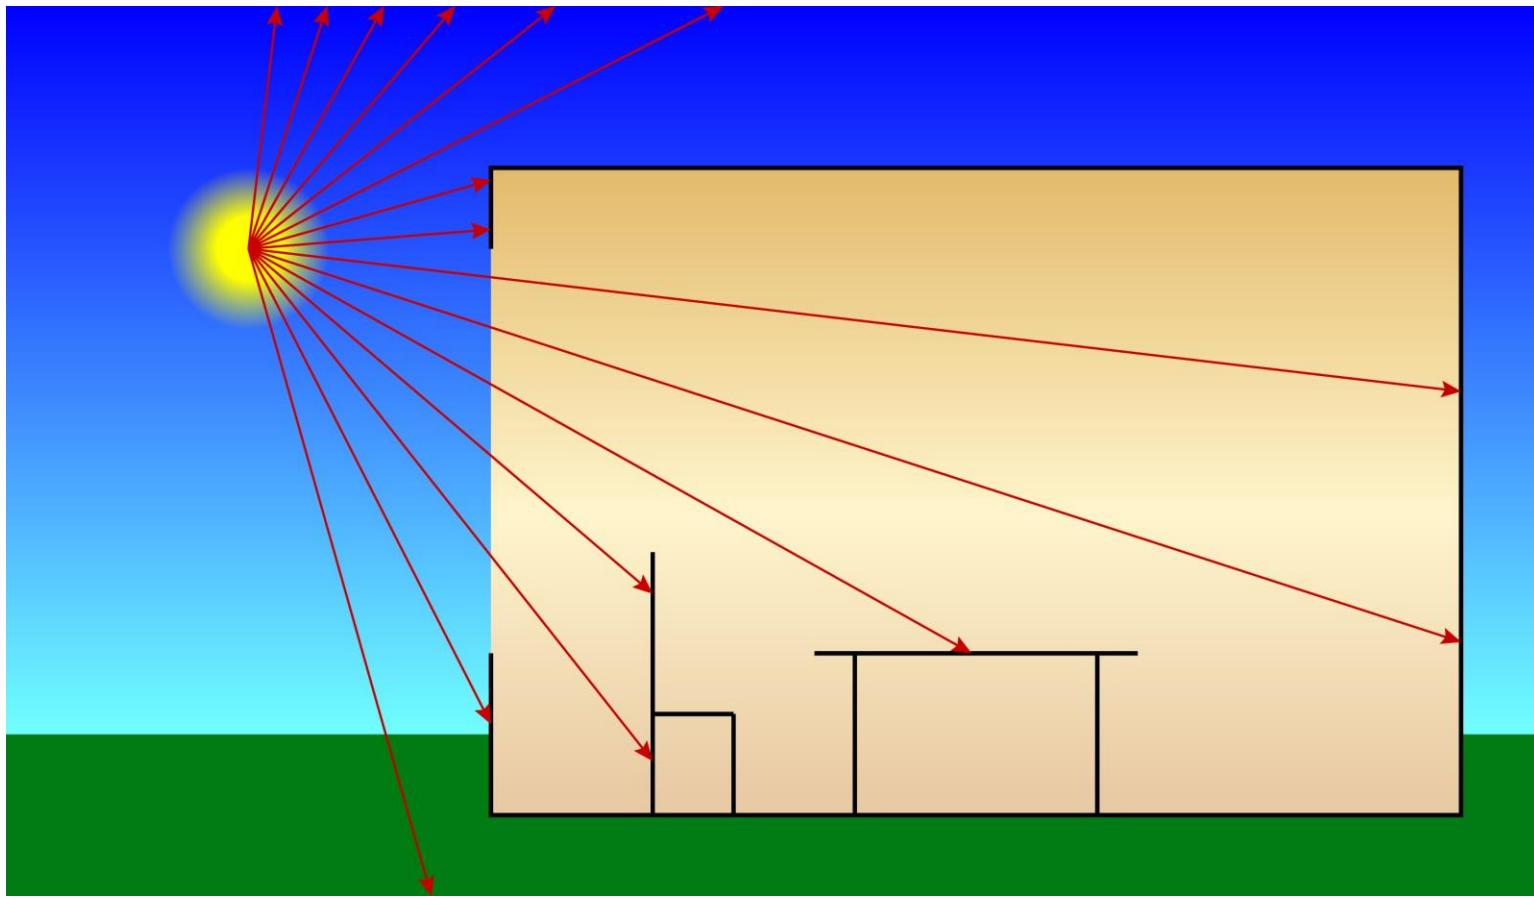


# Light from a source ...



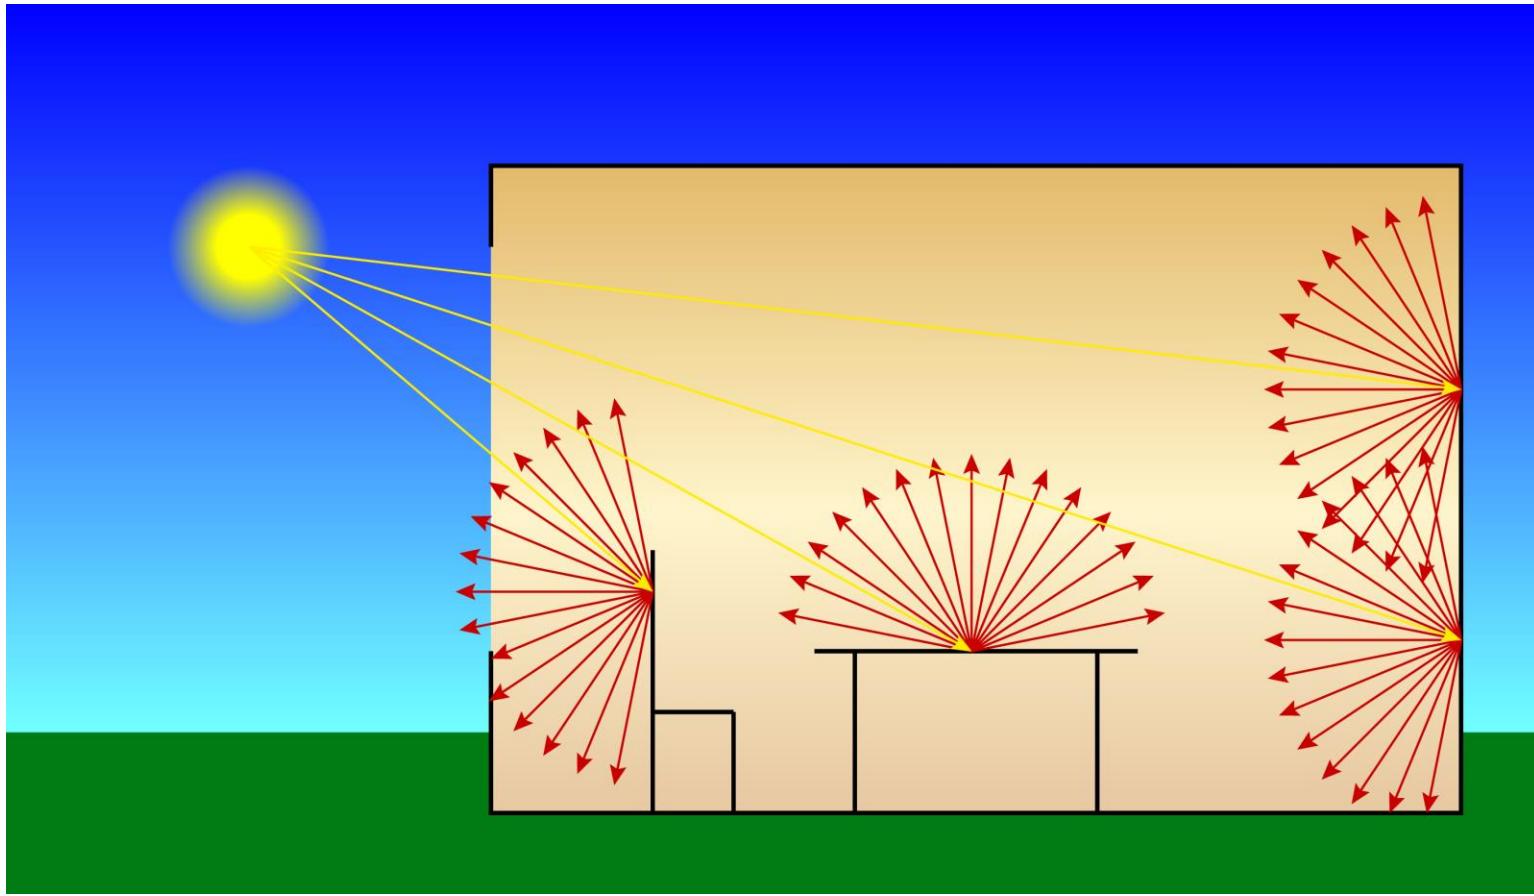


... illuminates a scene.



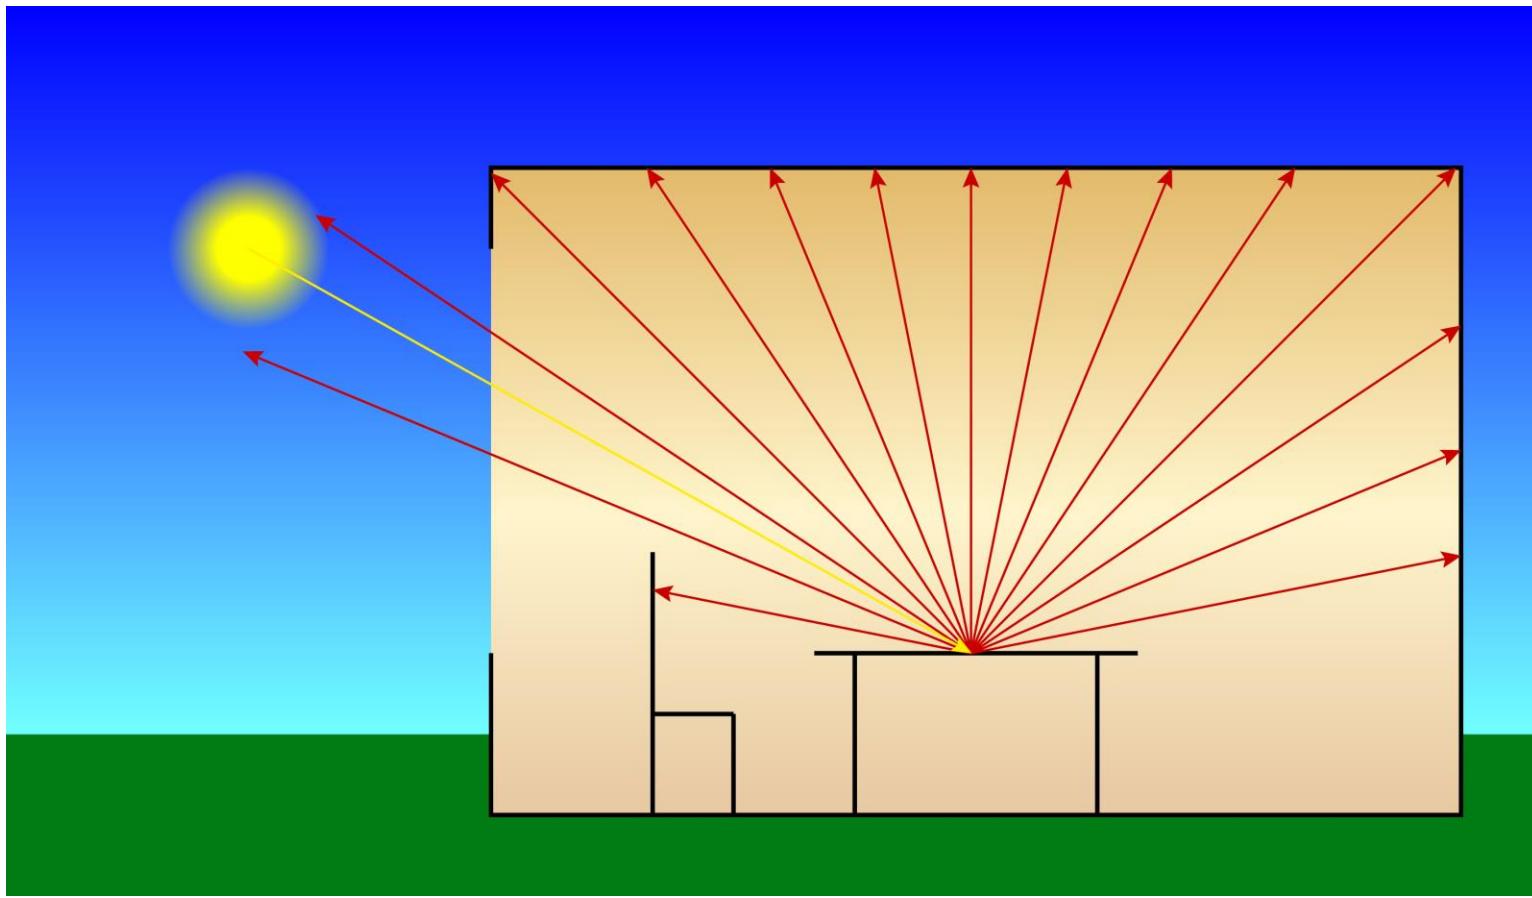


# Each surface scatters light ...



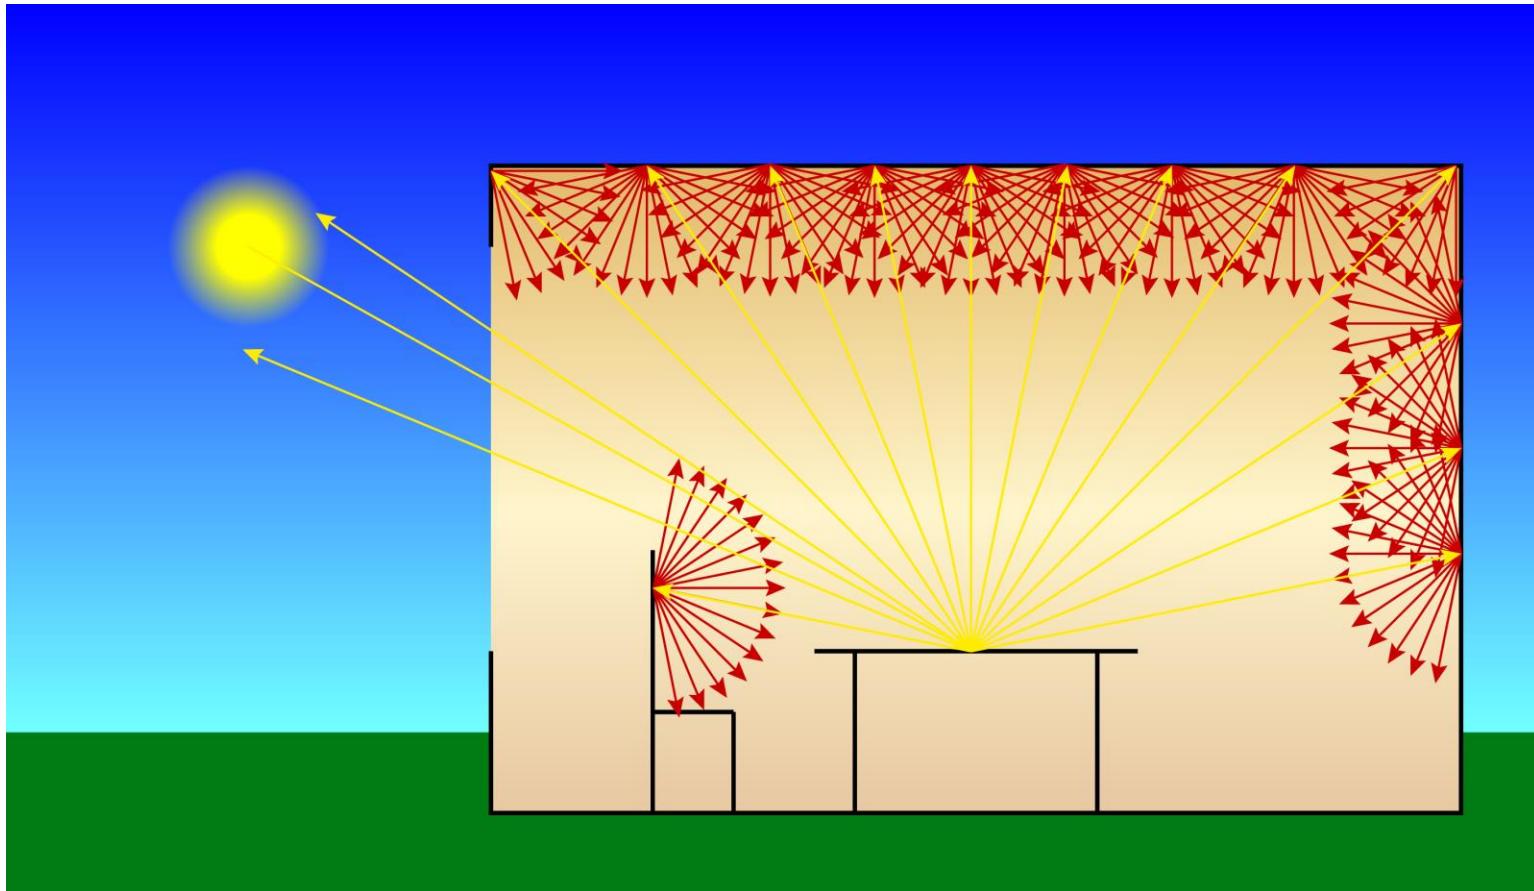


... that illuminates other surfaces ...



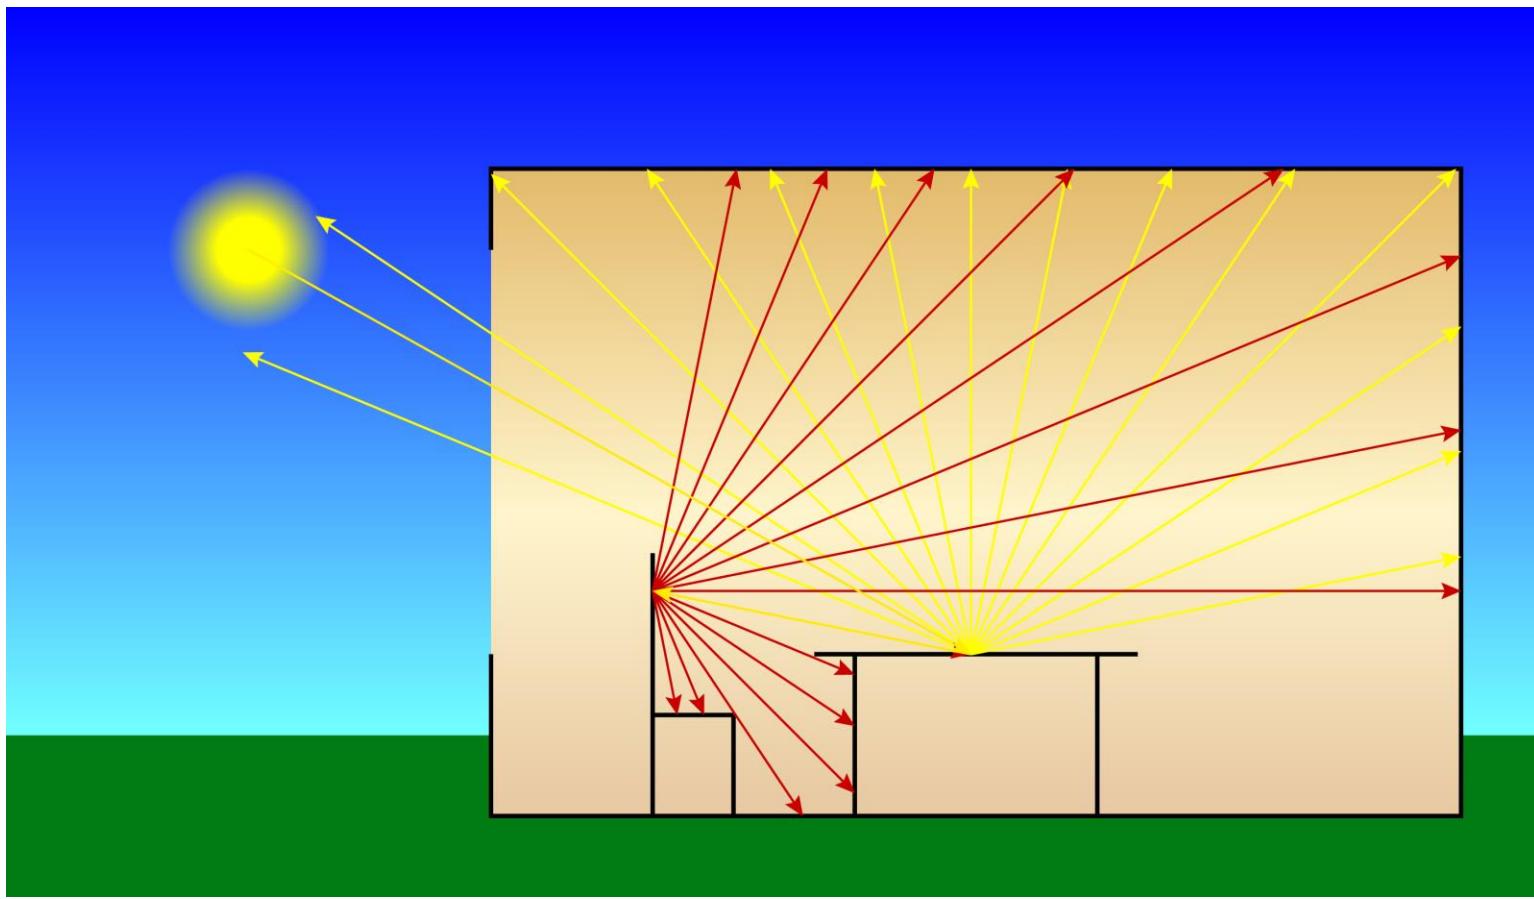


.. that scatter light.



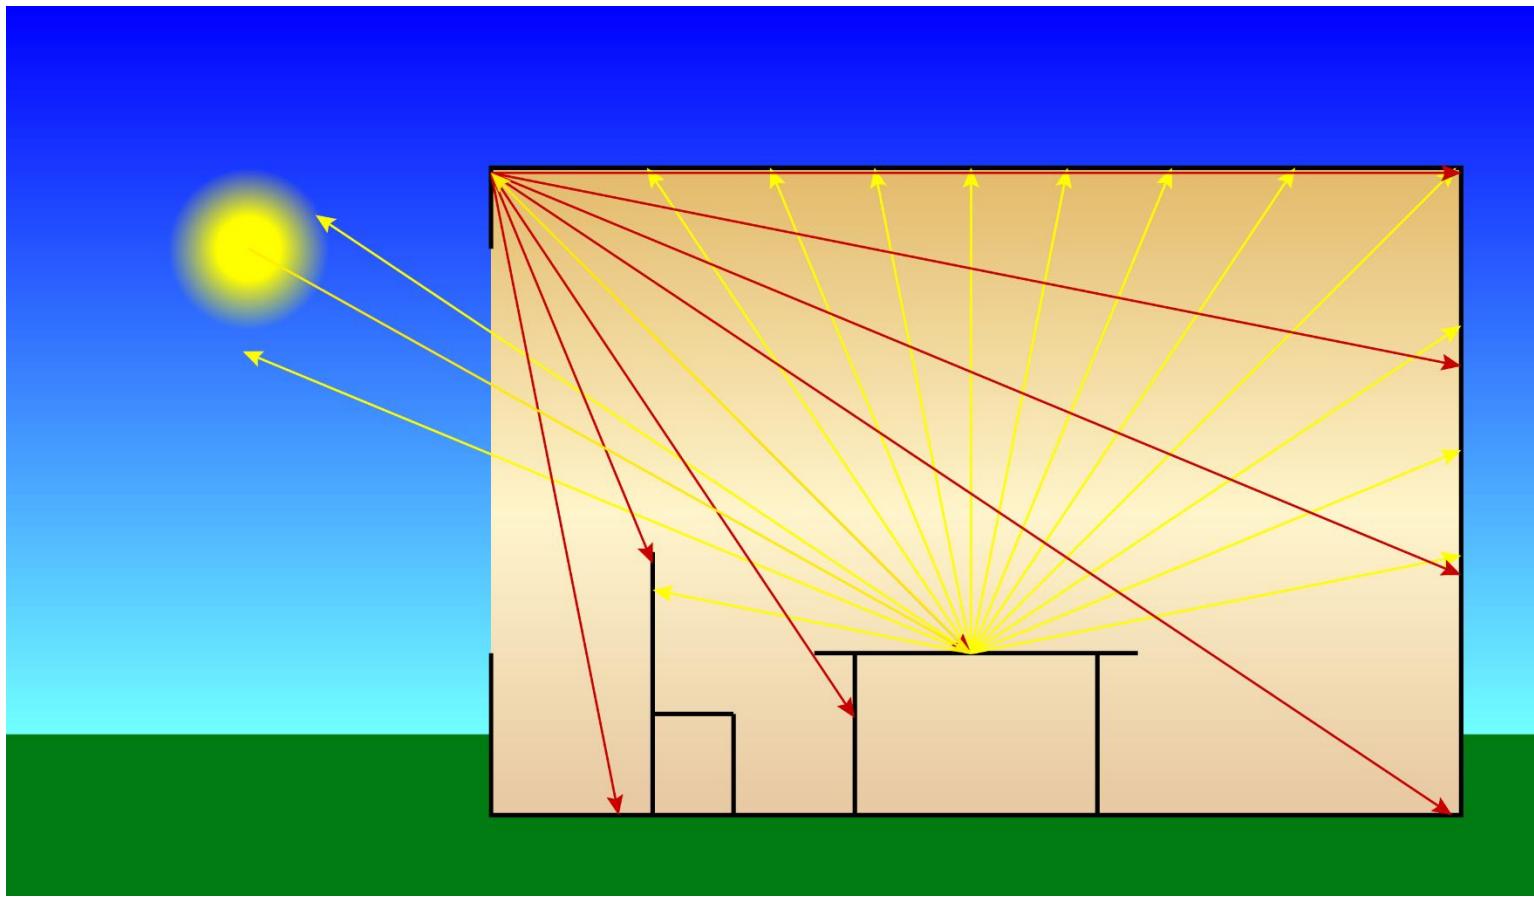


# The degree of interreflection ...



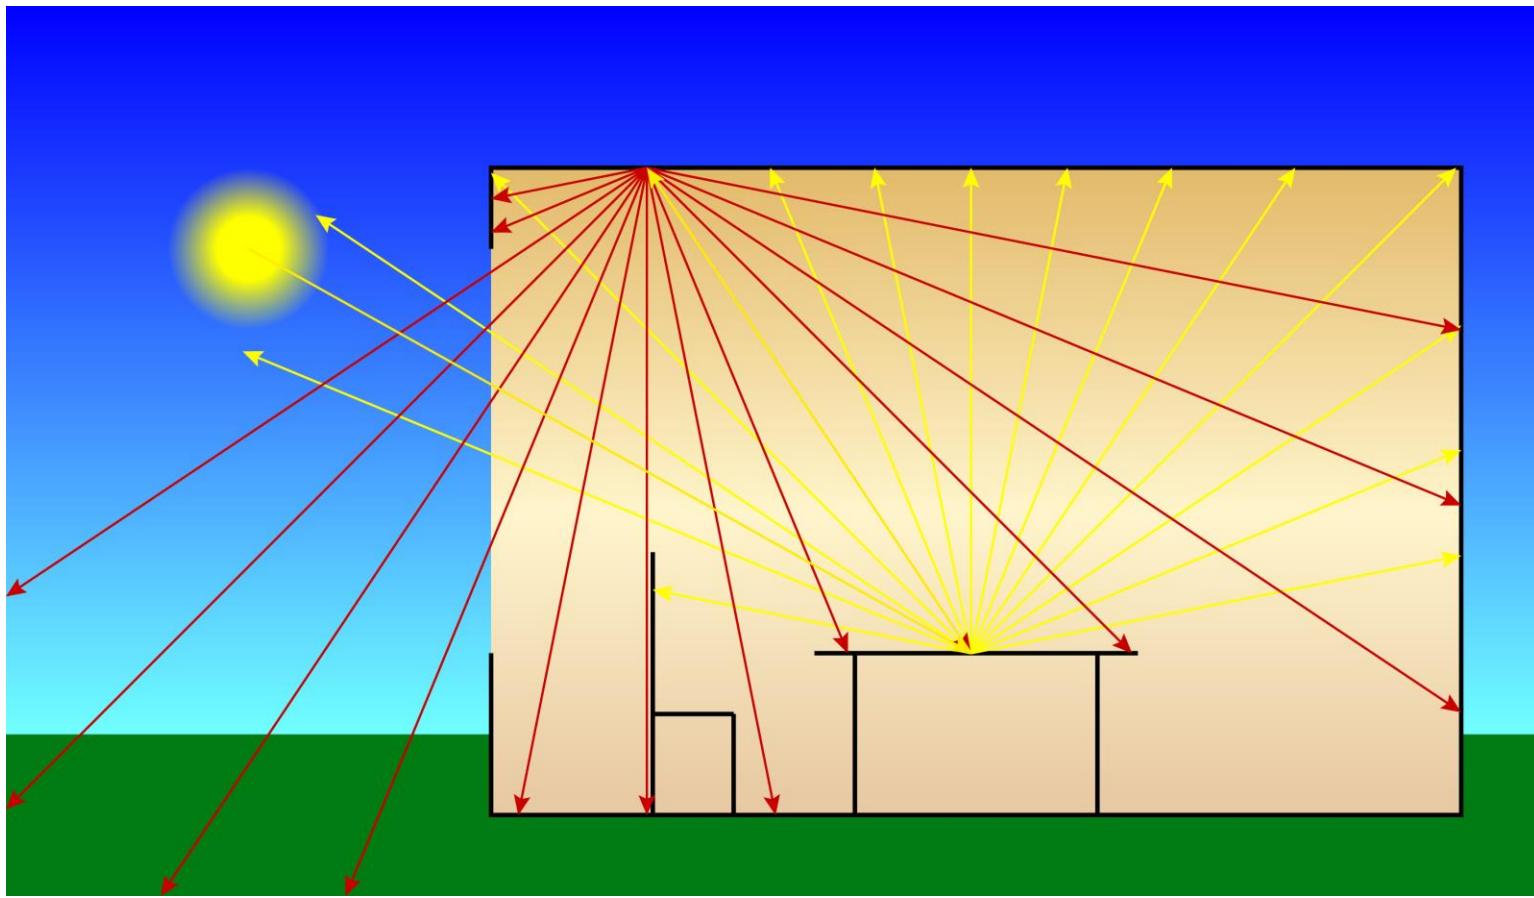


... depends on the material properties ...



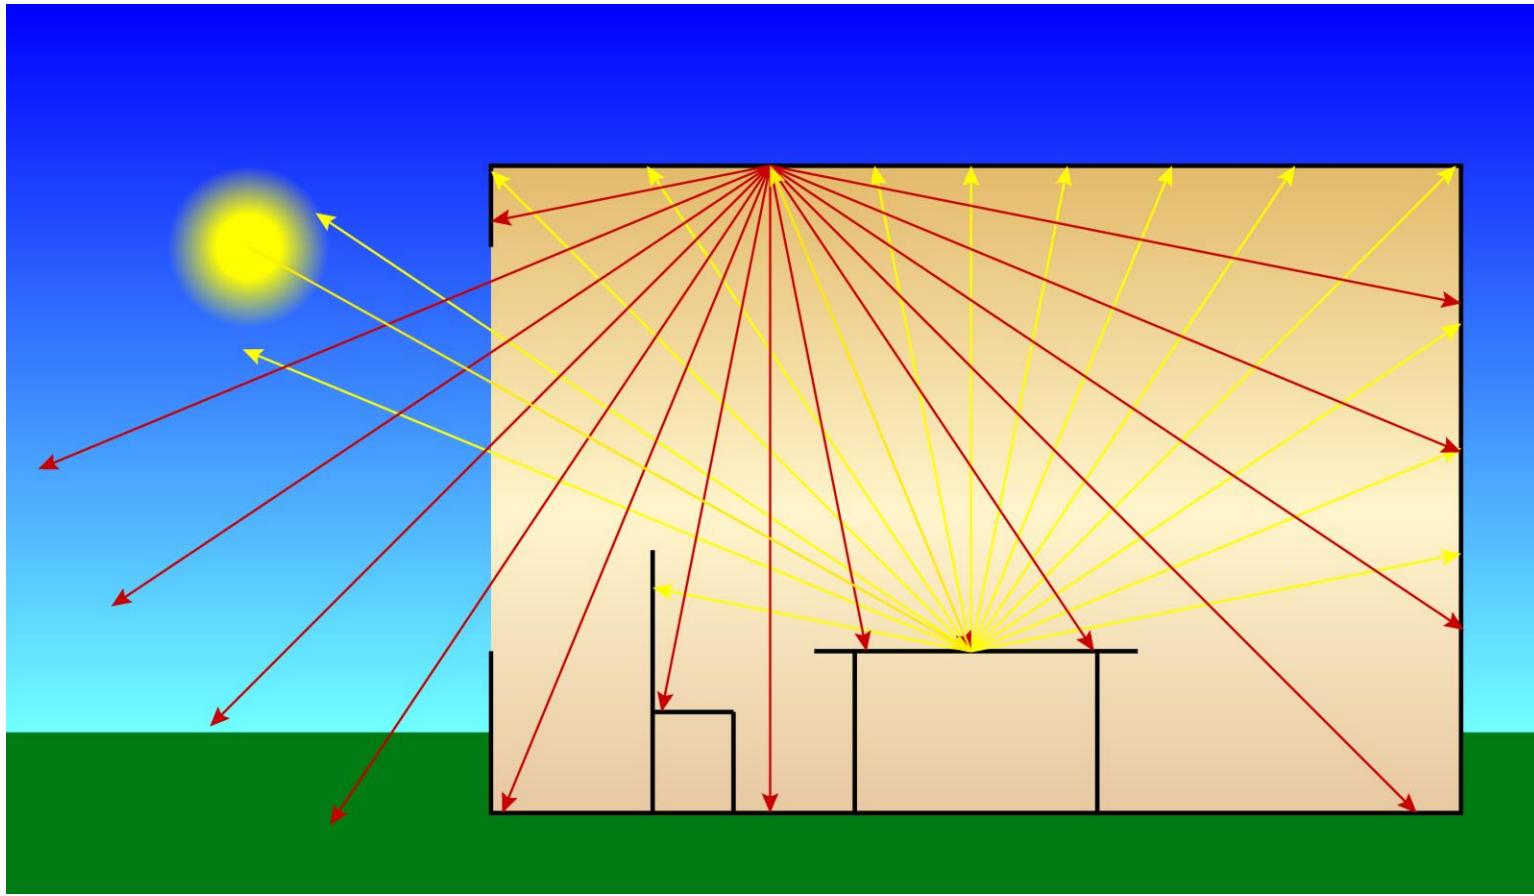


... of the surfaces in the scene.



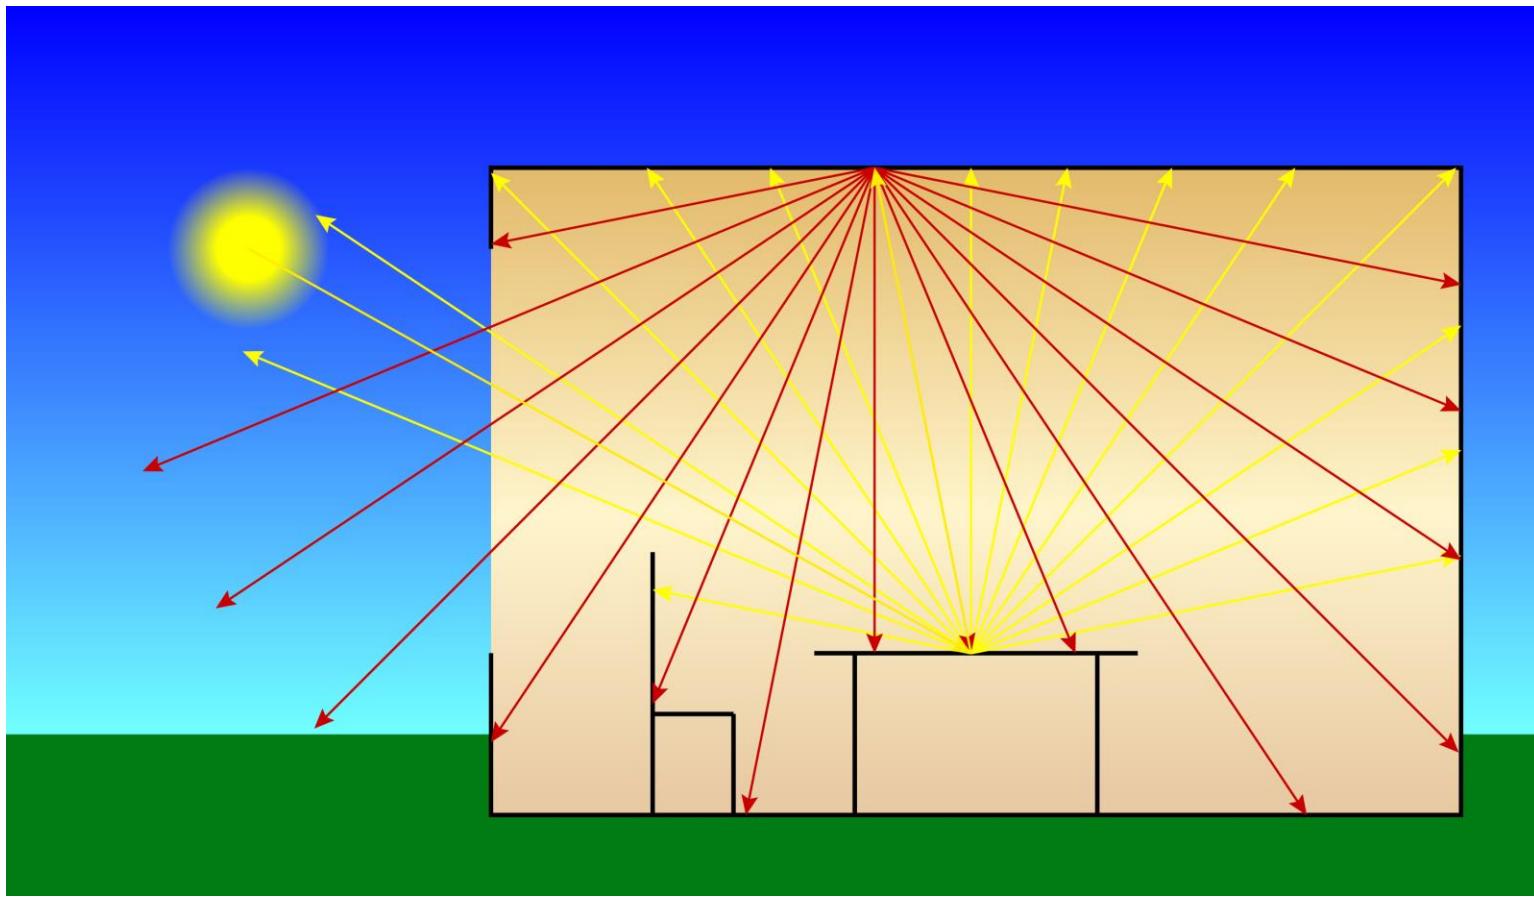


# Materials are transparent ...



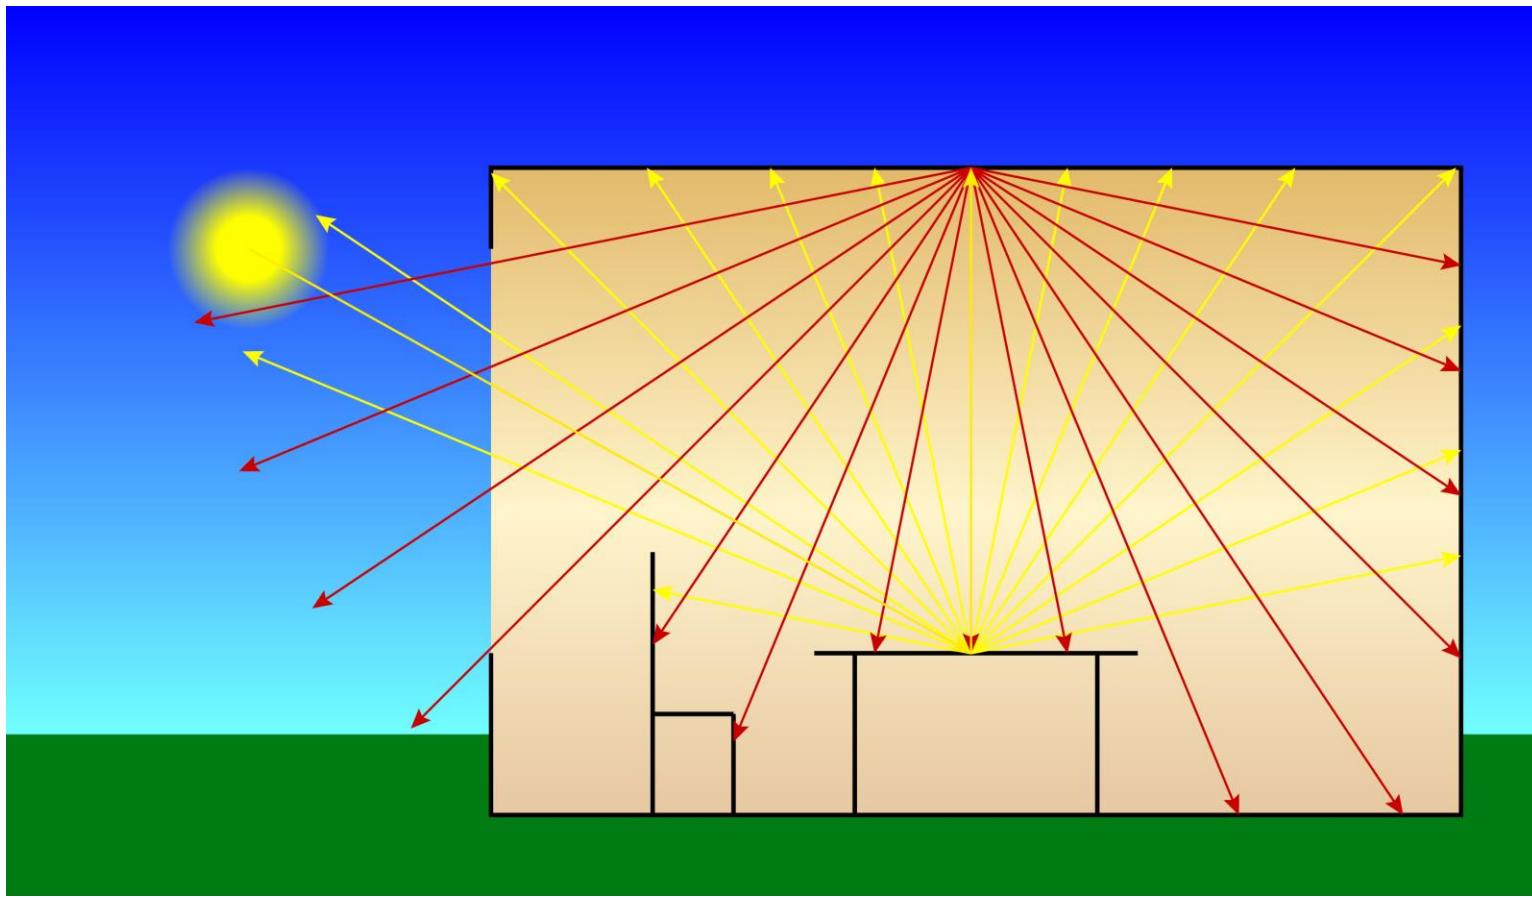


... to some wavelengths ...



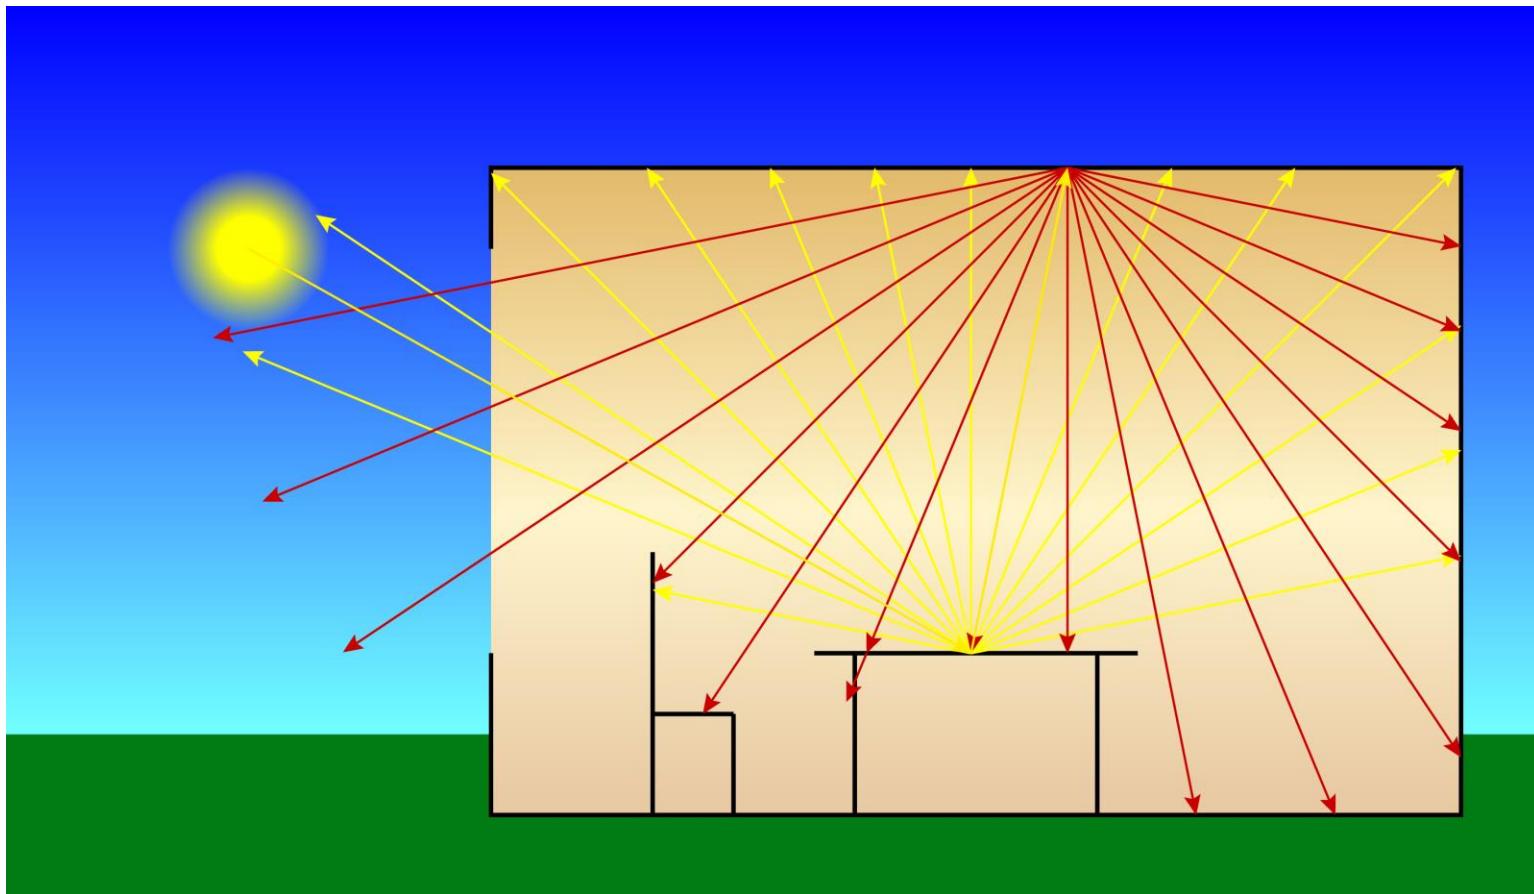


... opaque to others ...



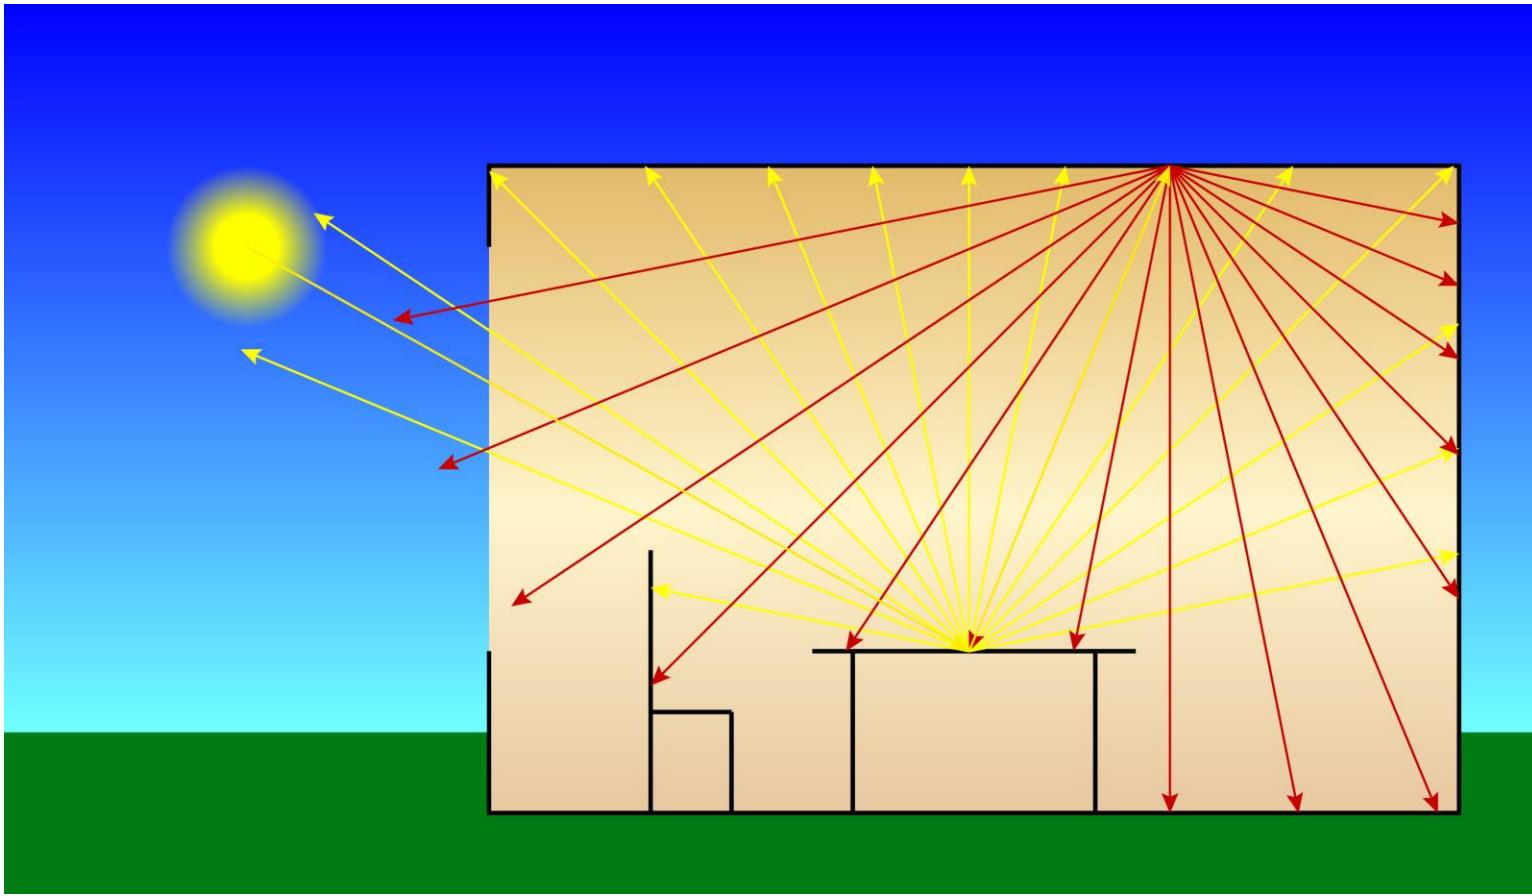


... and reflect yet other wavelengths.



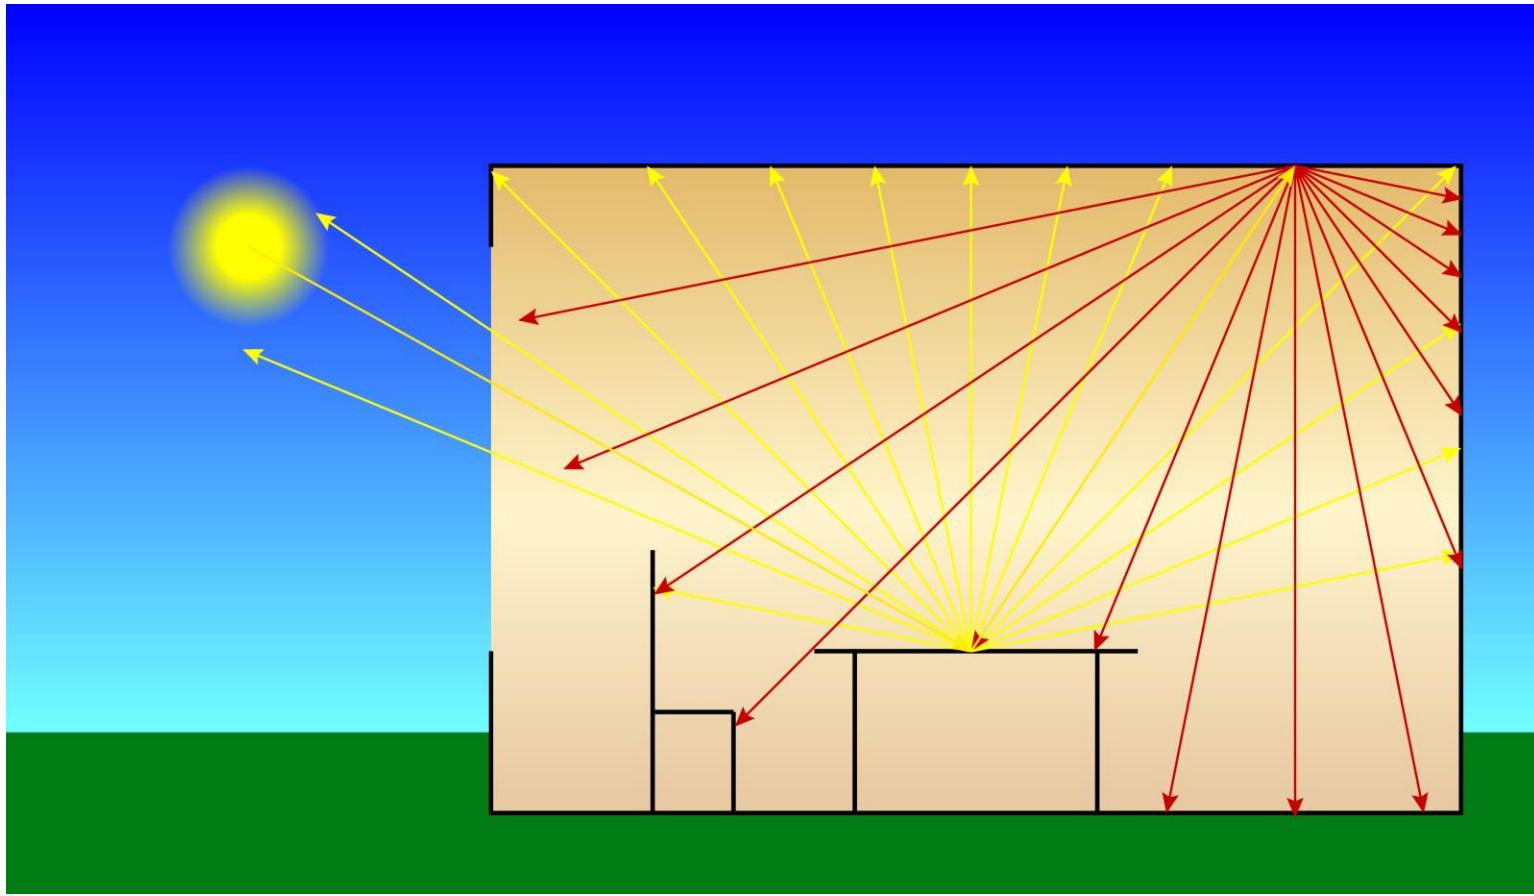


# Smooth surfaces may be mirror-like, ...



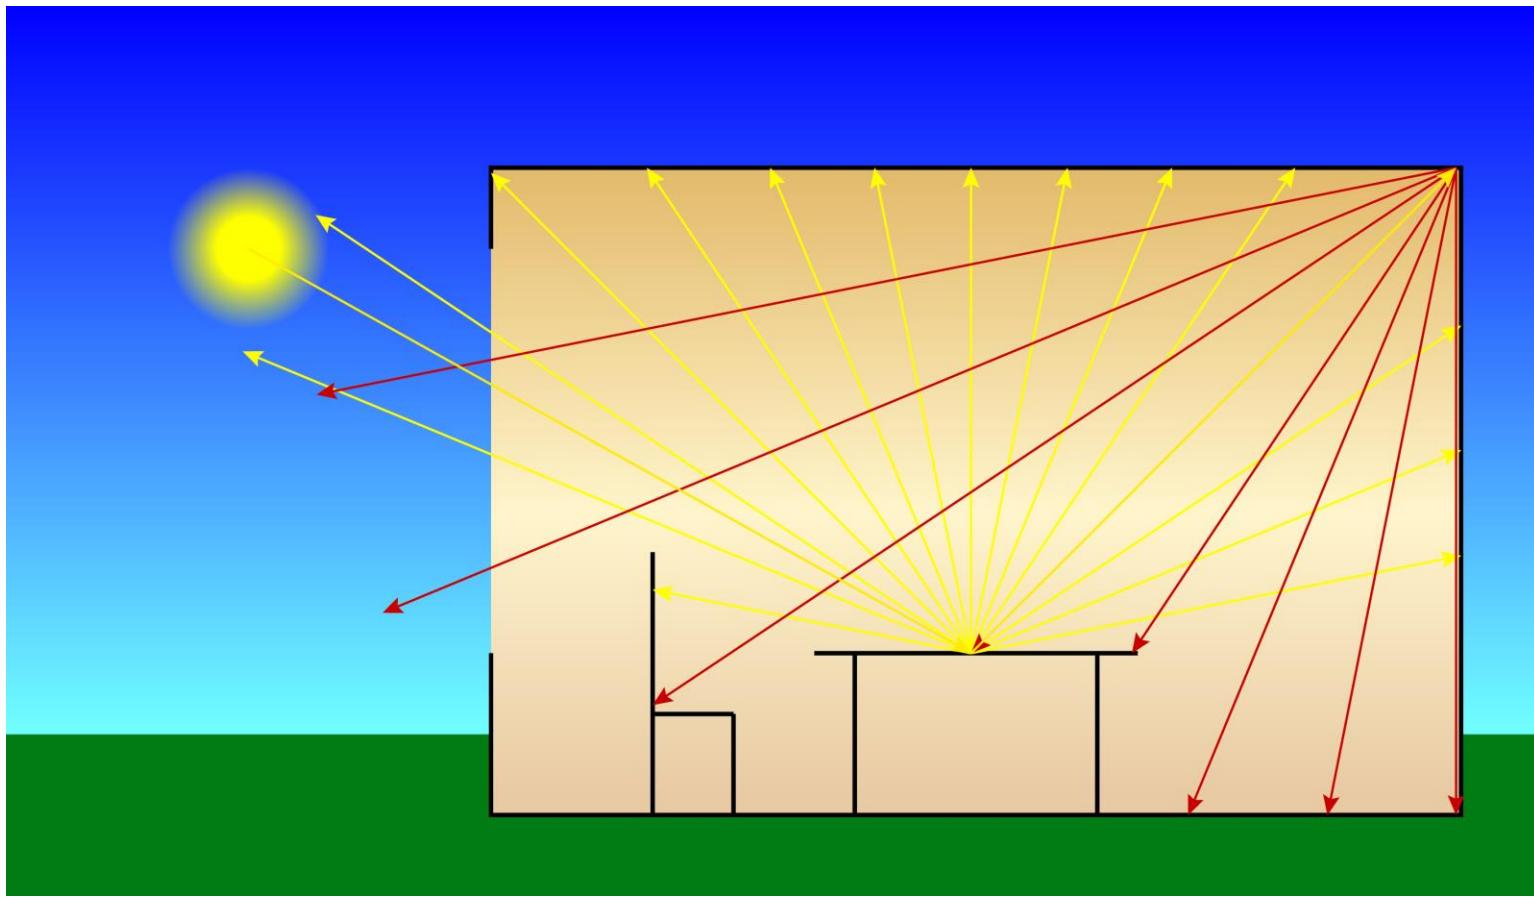


...rough surfaces, matte or diffuse.



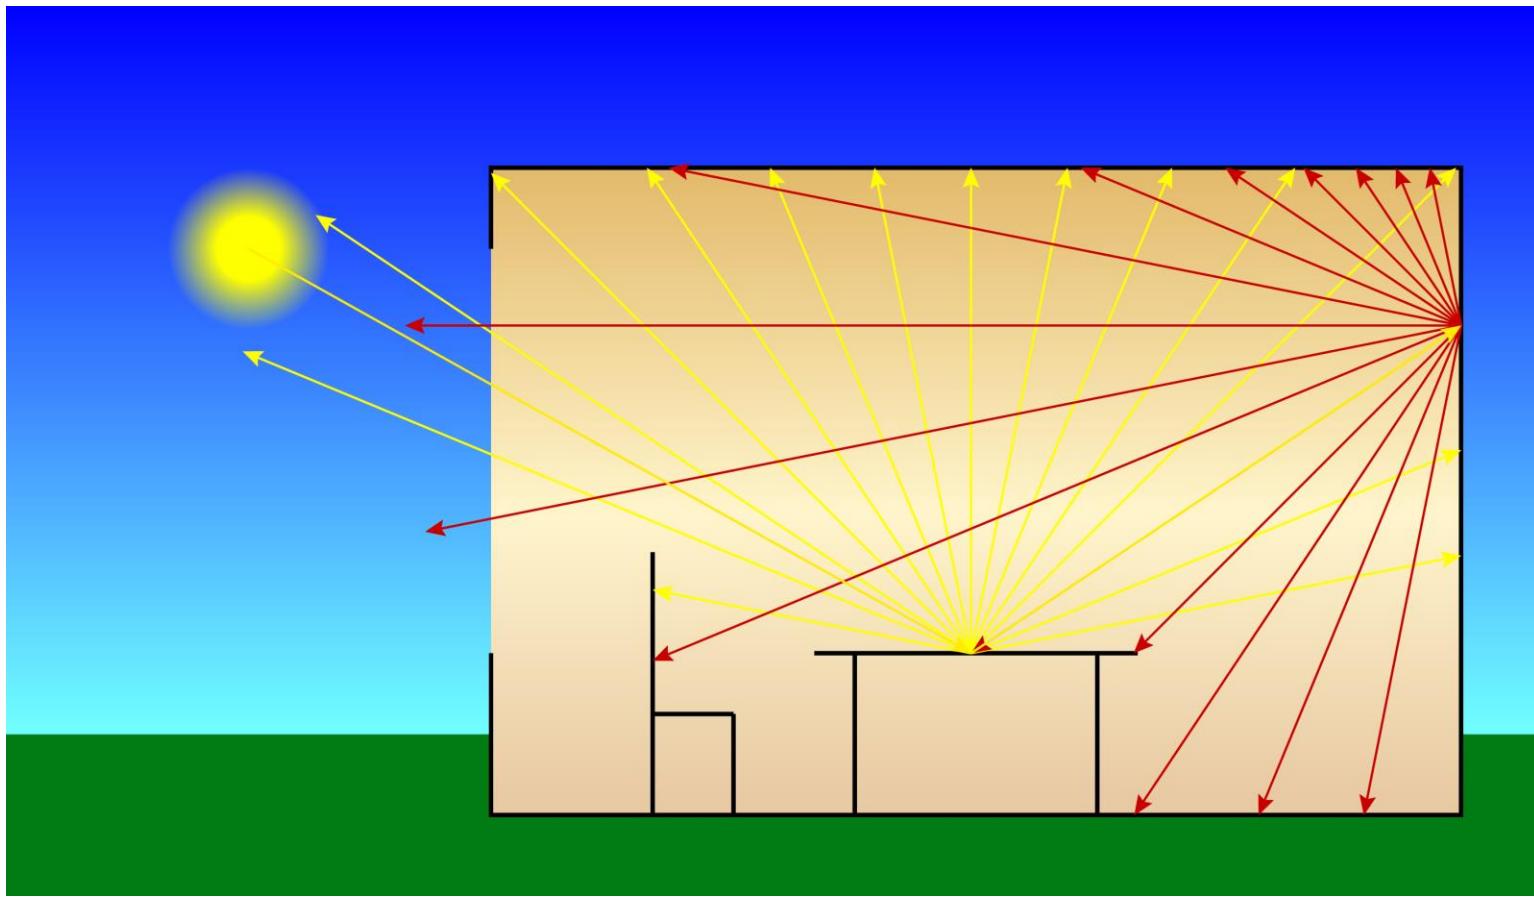


# Anisotropic surfaces scatter more in ...



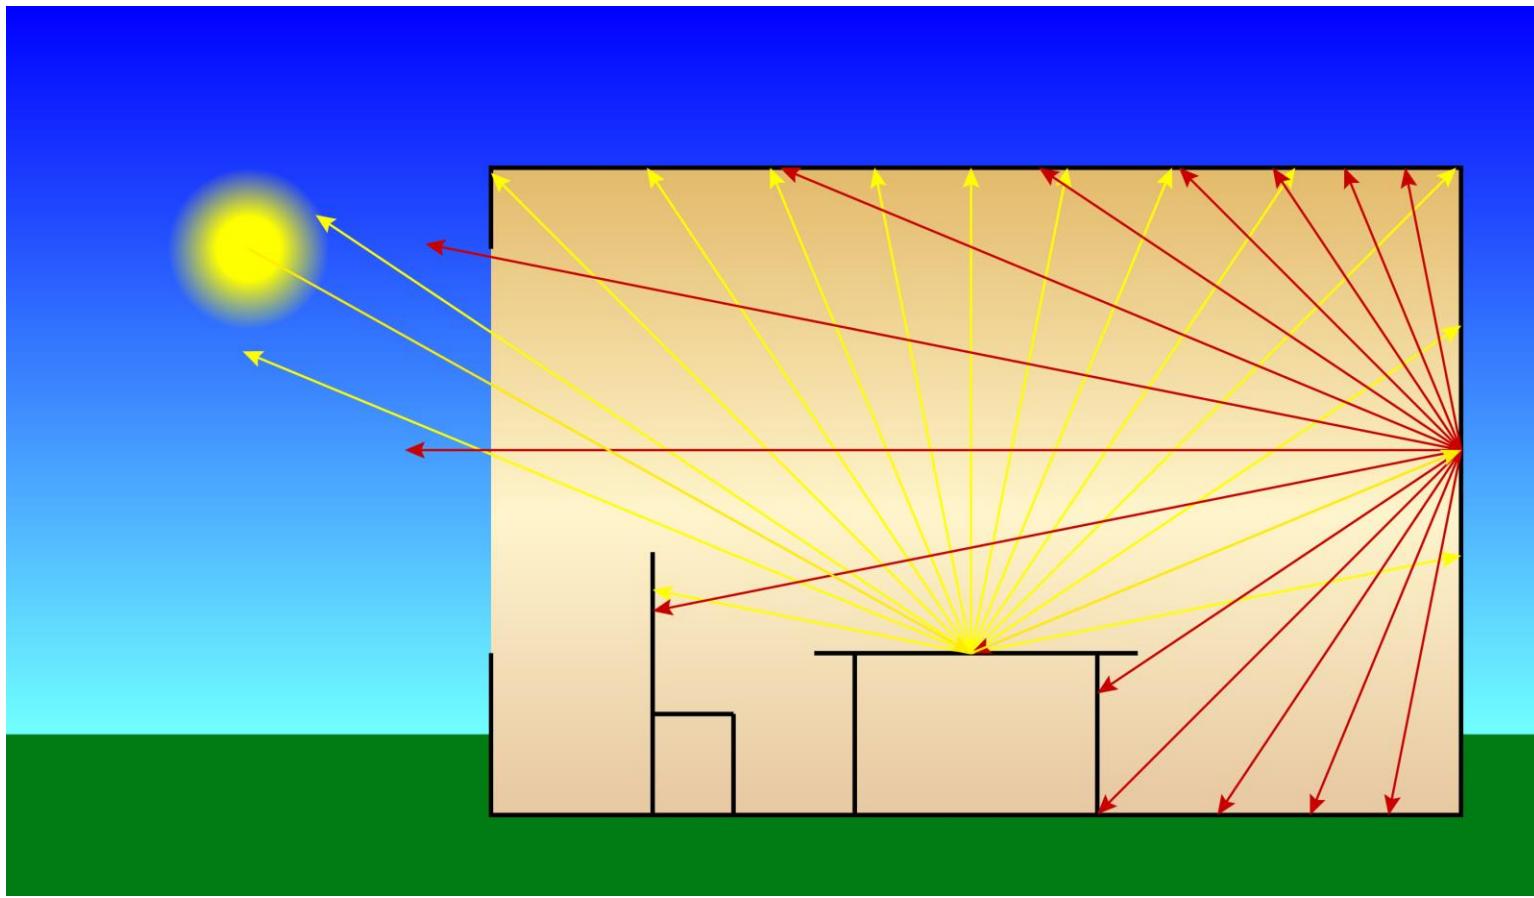


... in one direction than others.



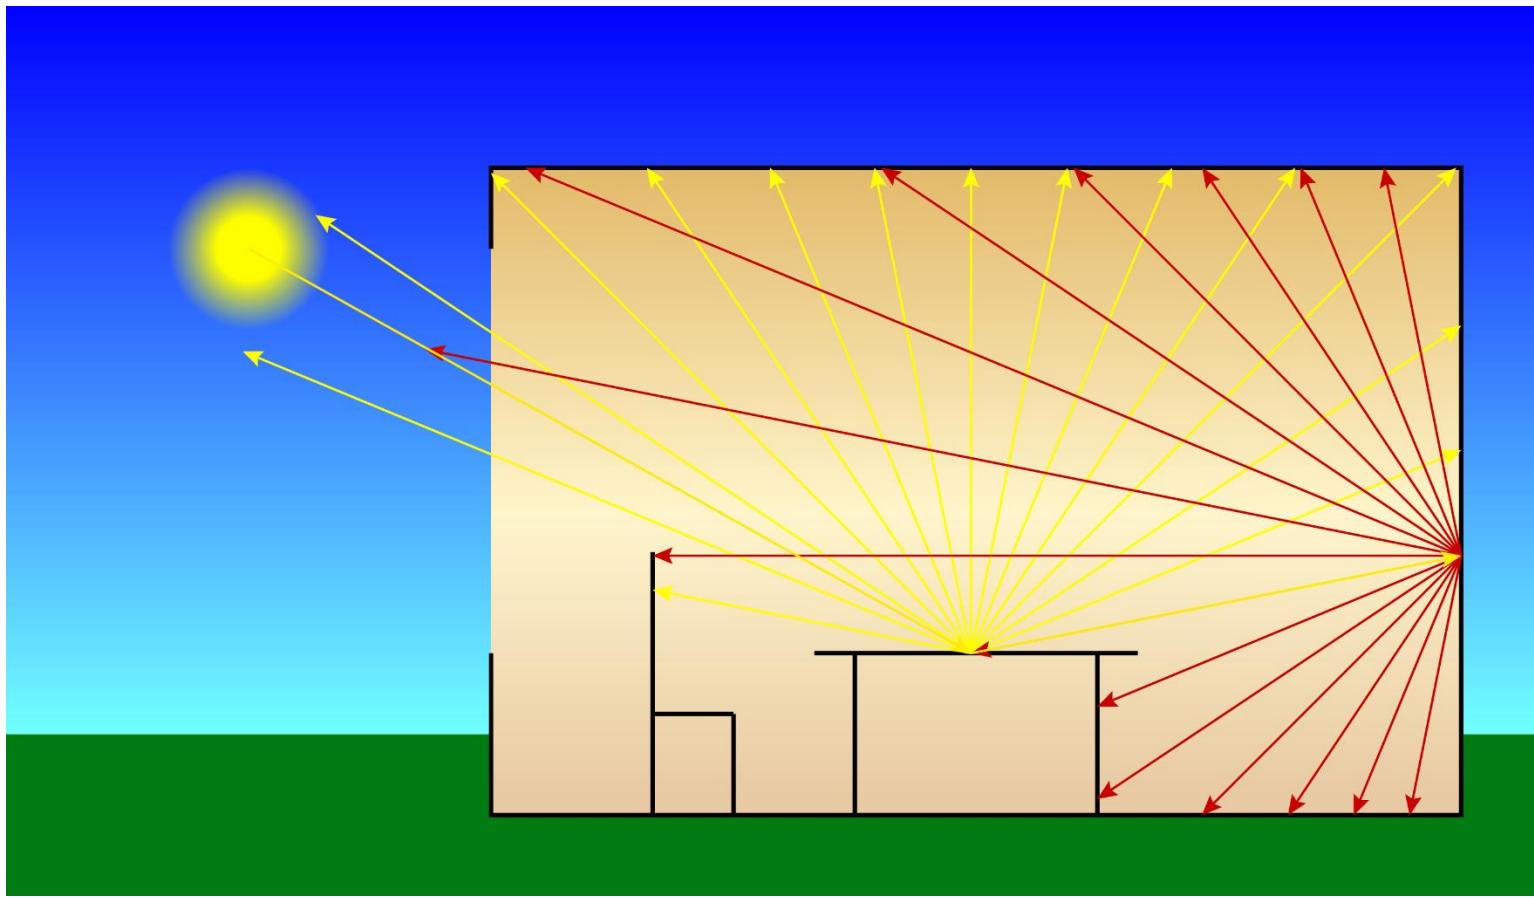


# The color of a surface affects those ...



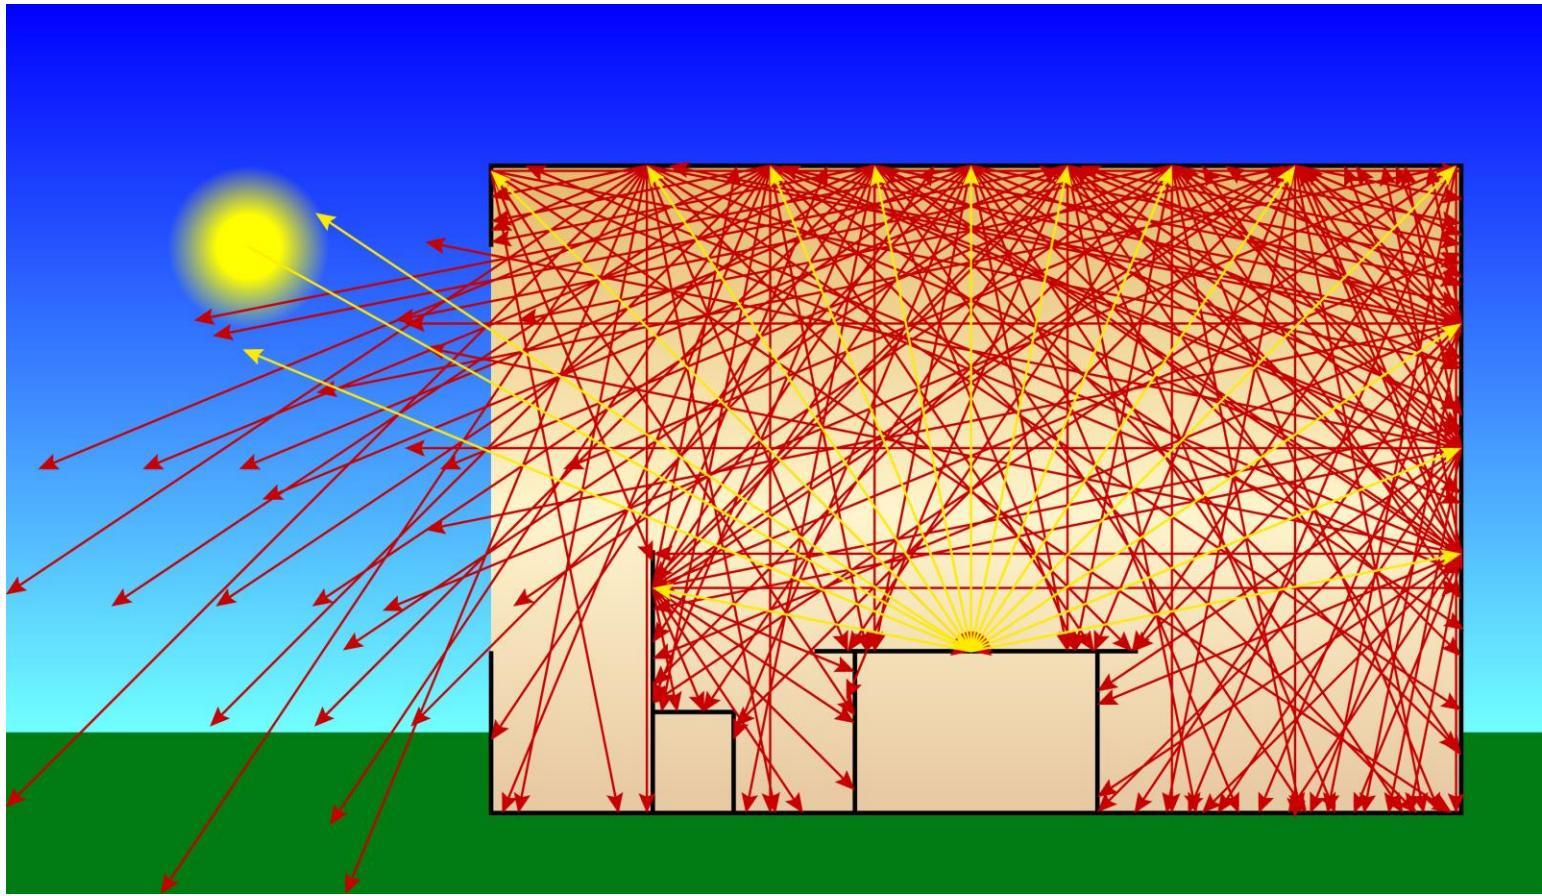


... of nearby surfaces. The result ...



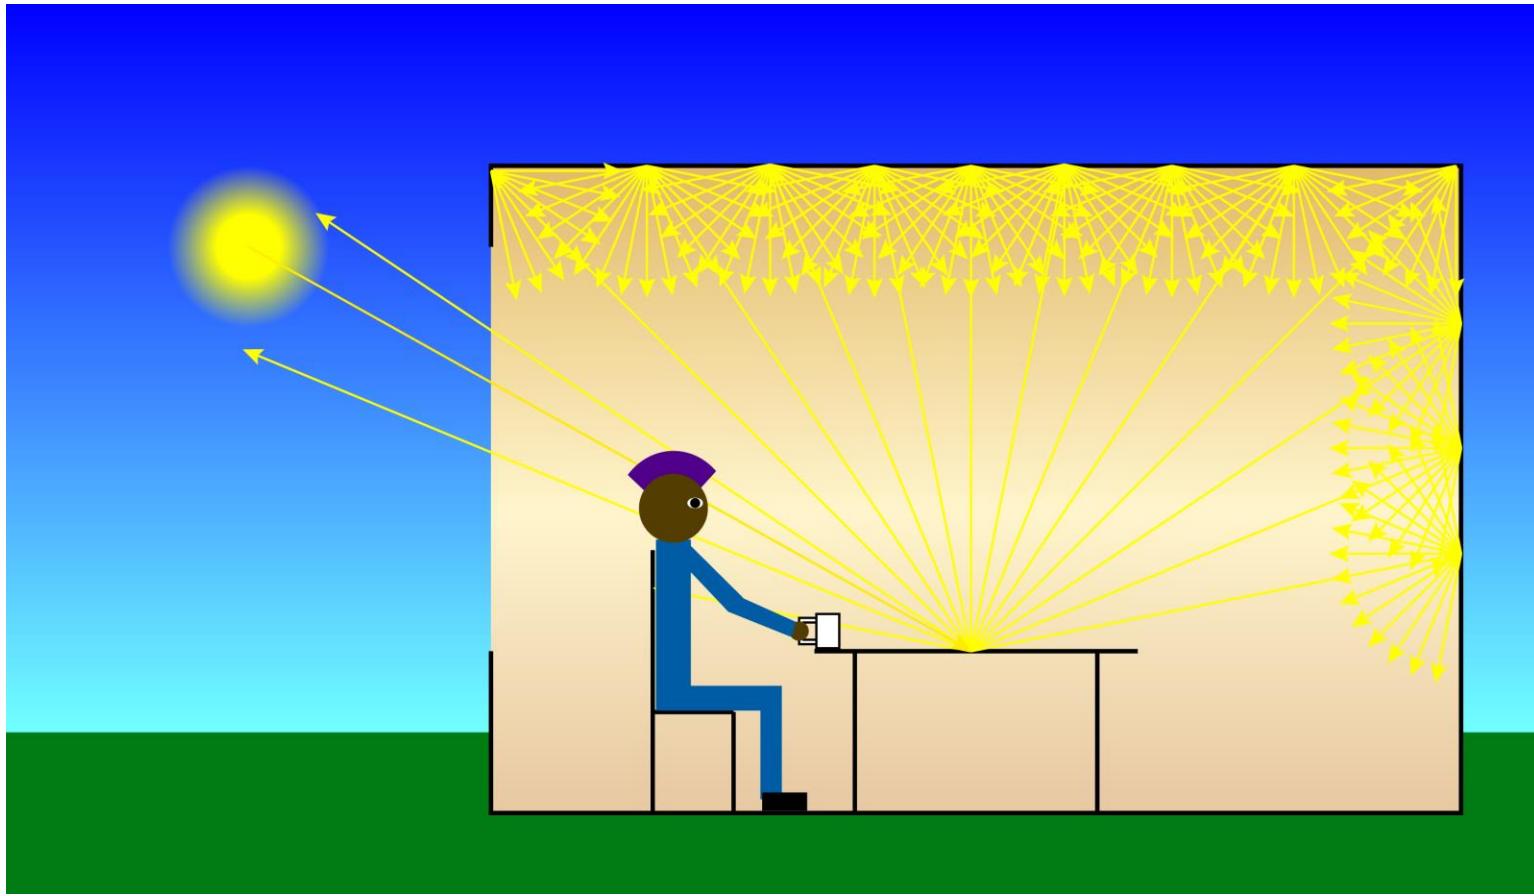


... is highly complex.



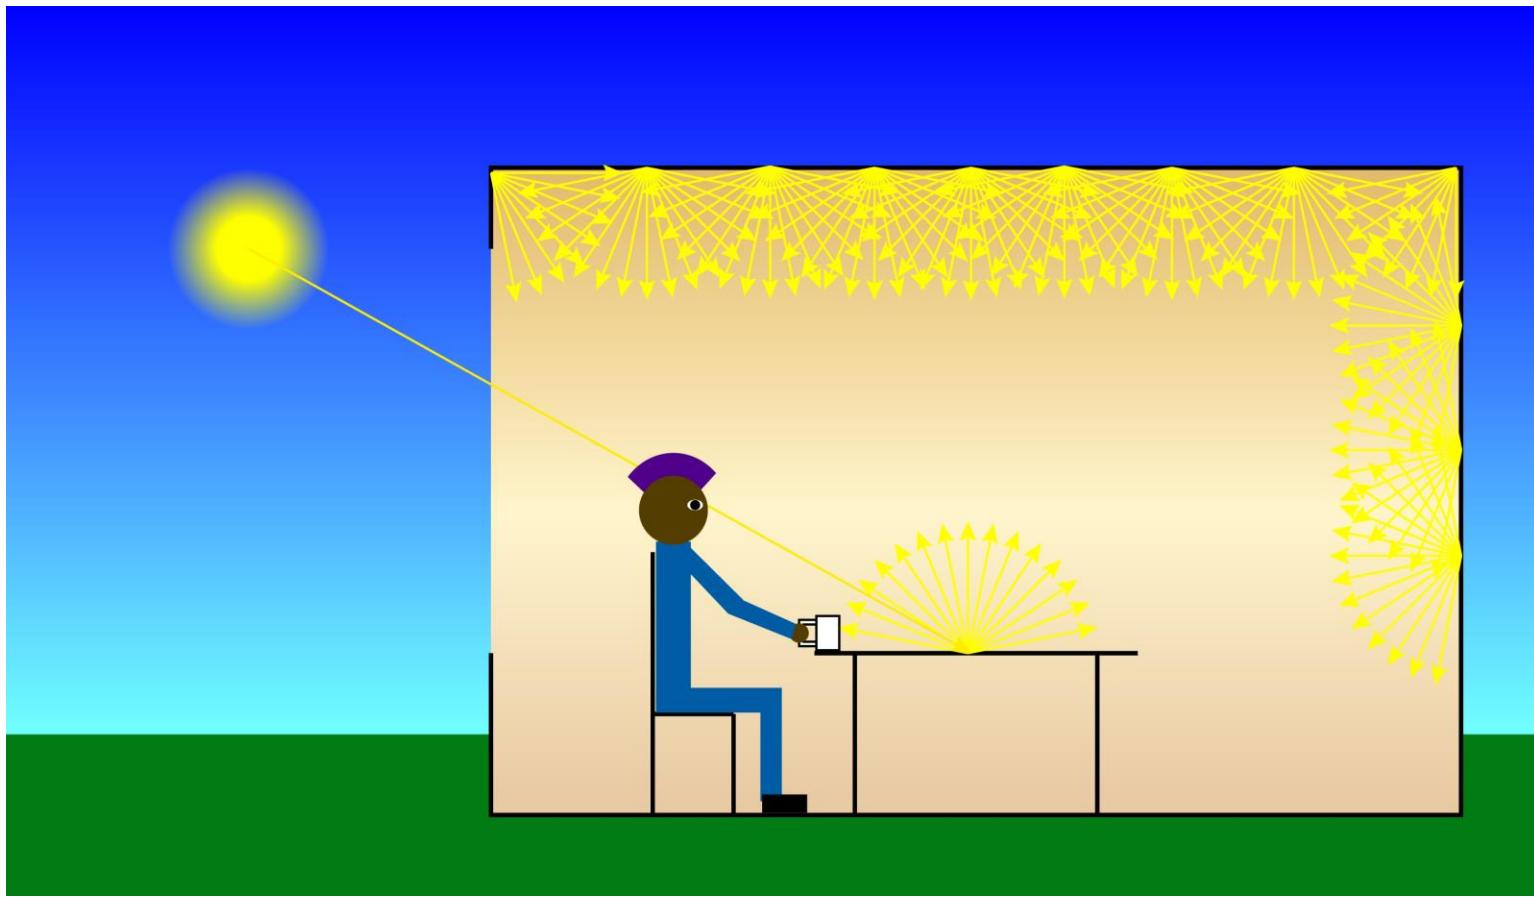


# An observer, awash in all the light, ...



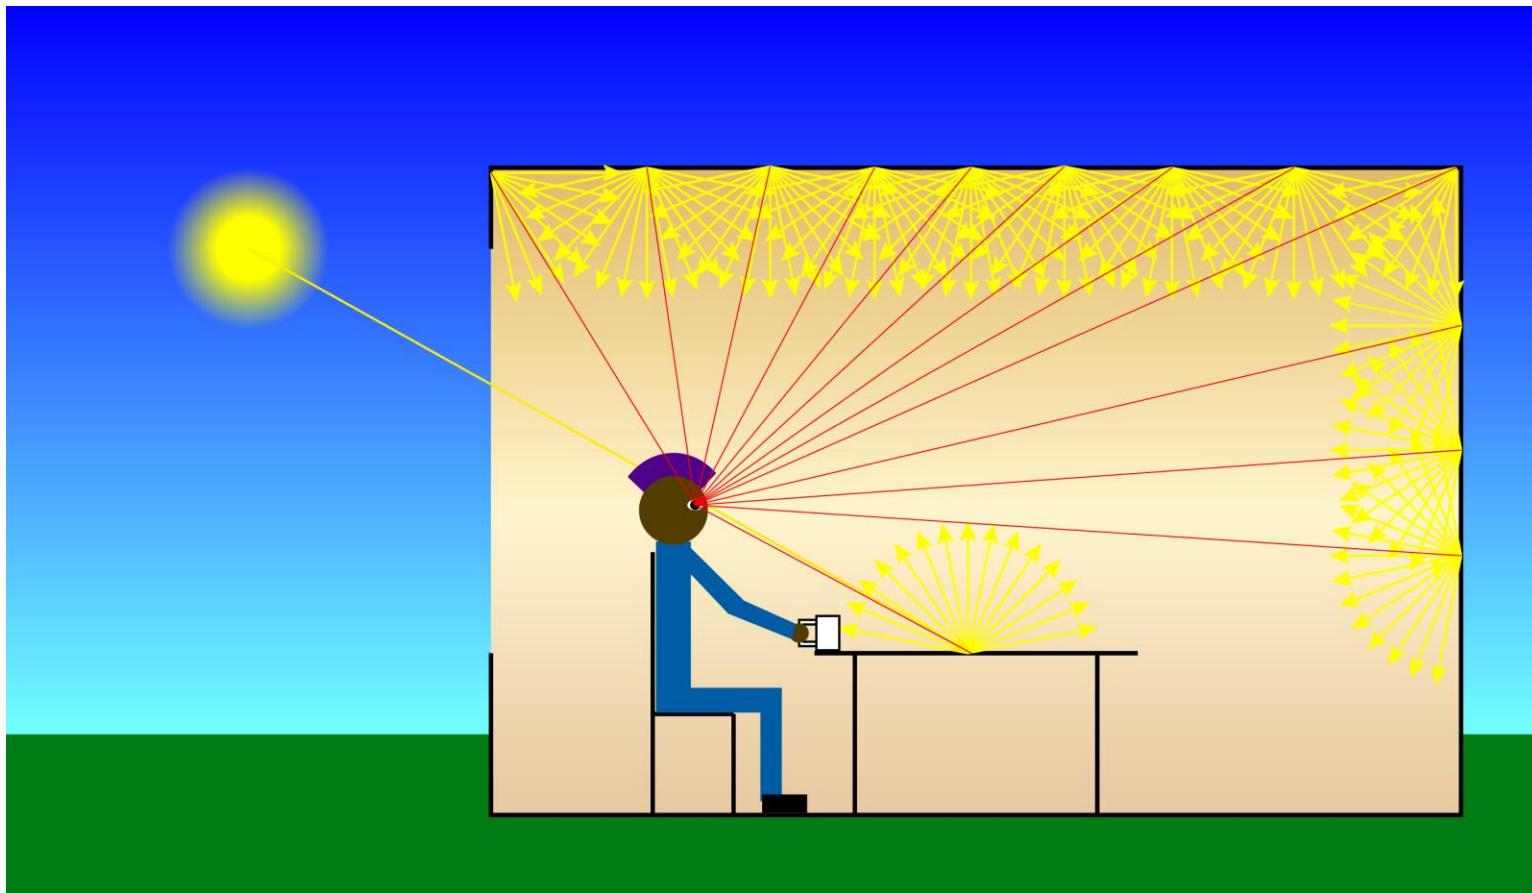


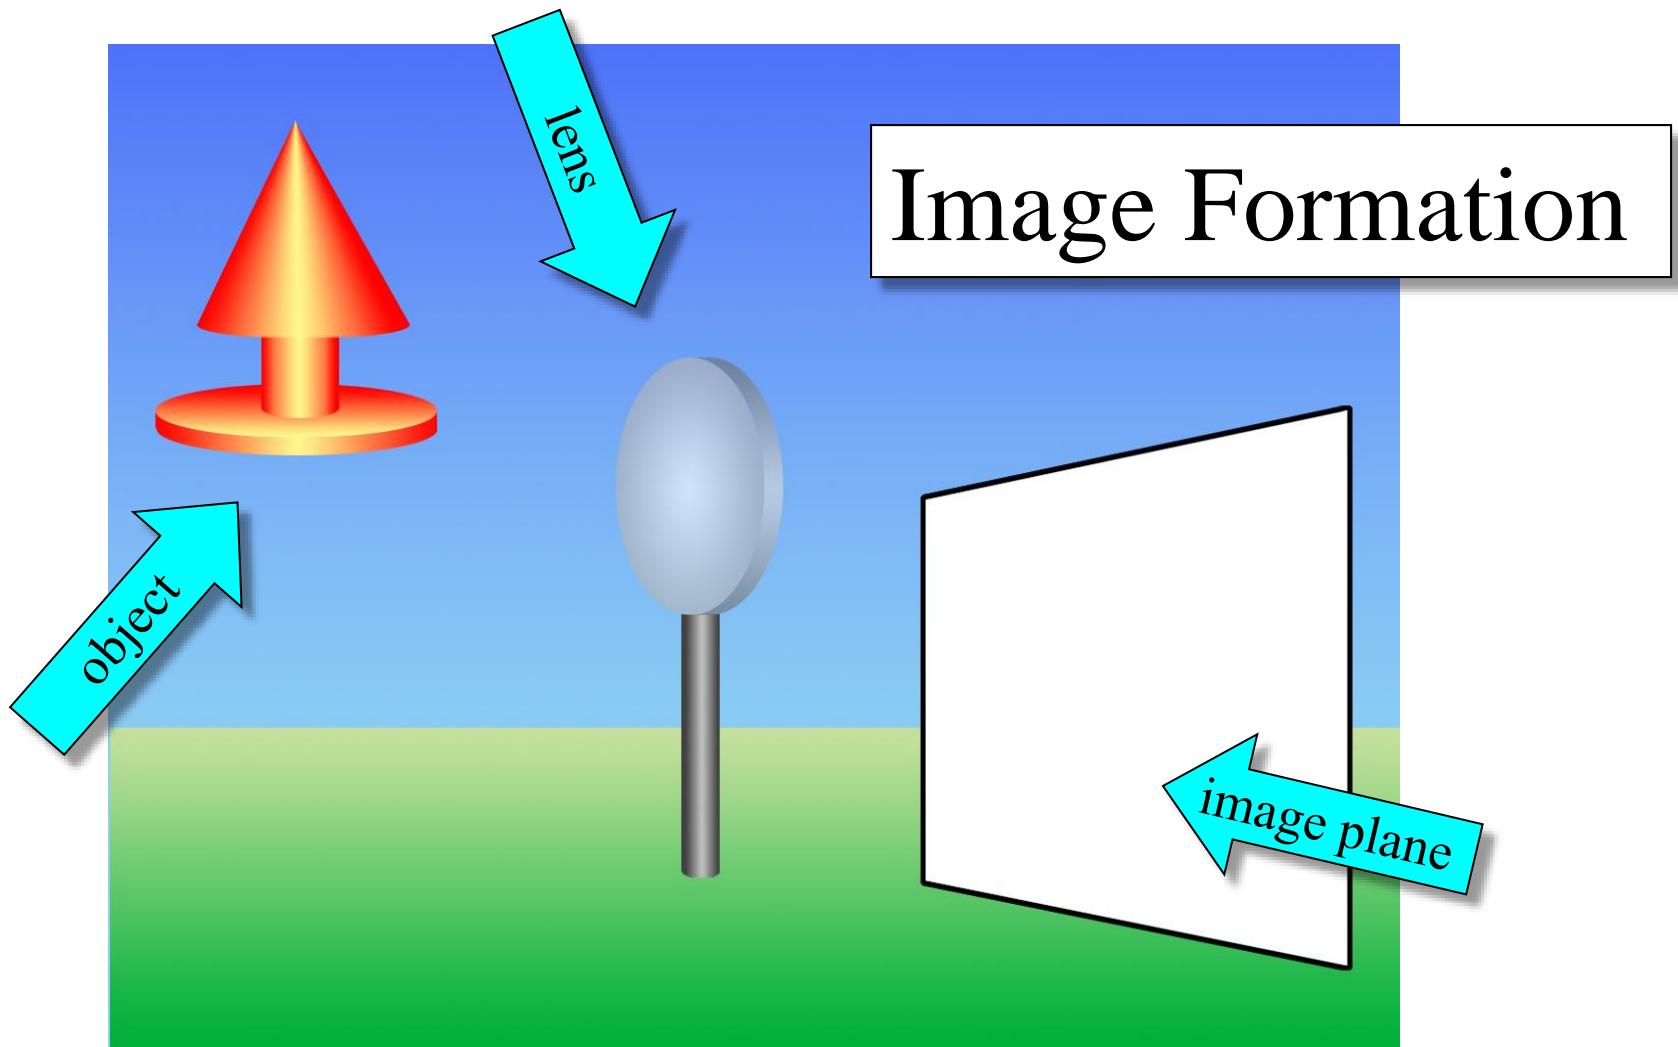
... sees the result of both illumination ...

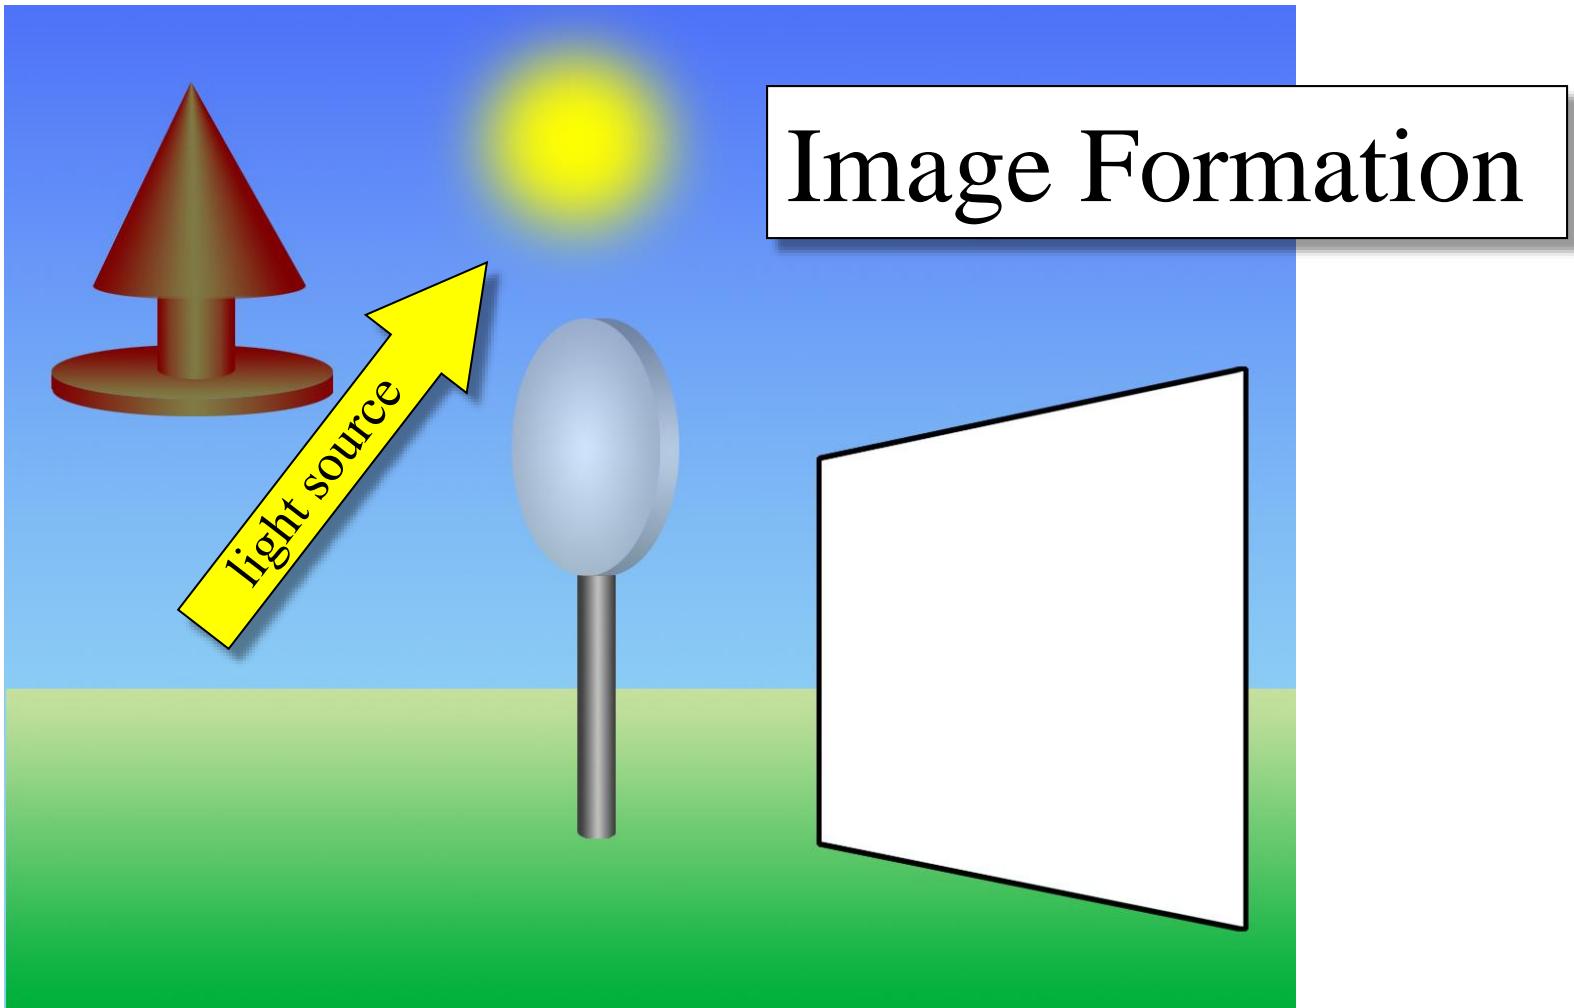


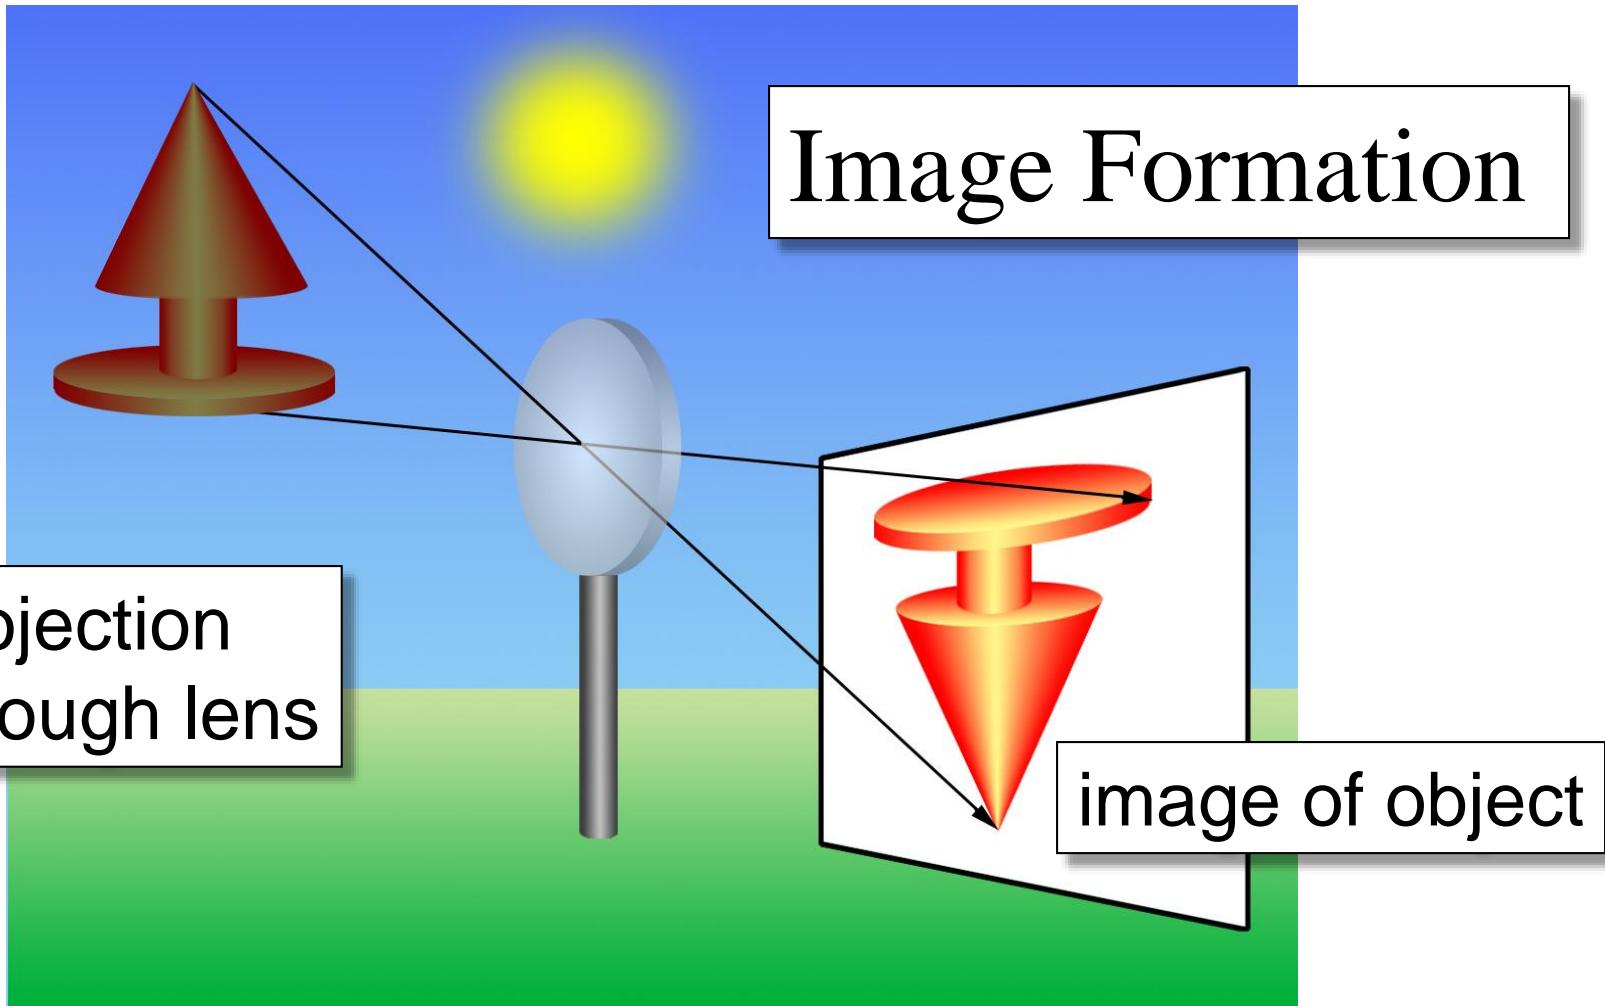


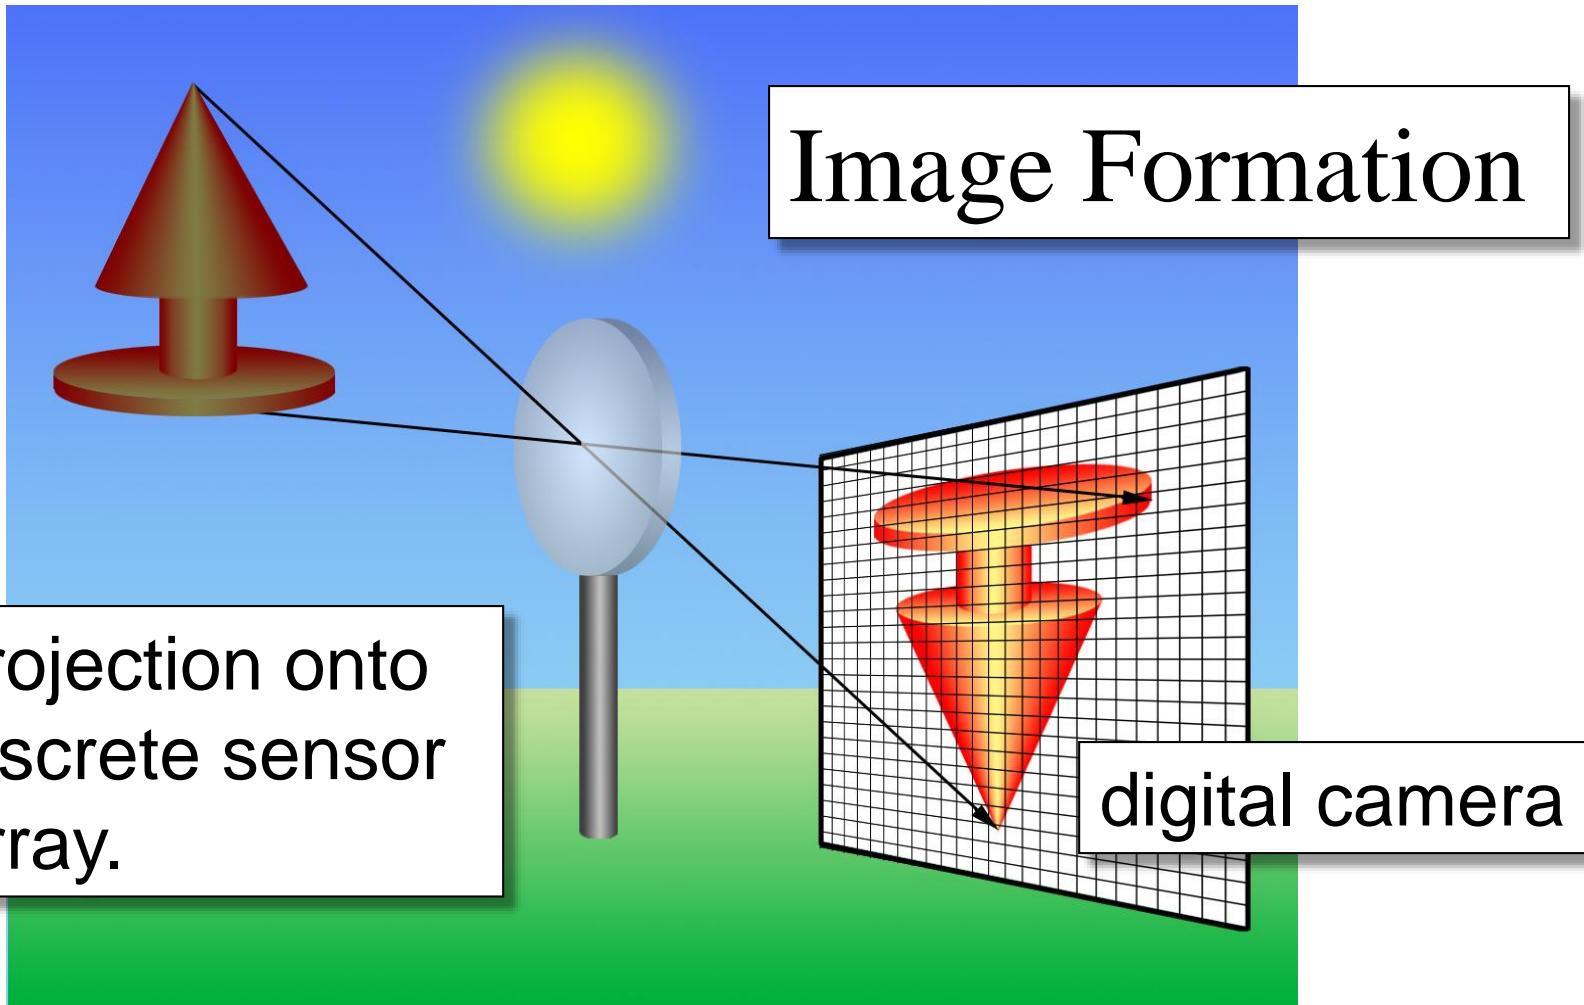
... and surface interreflection.

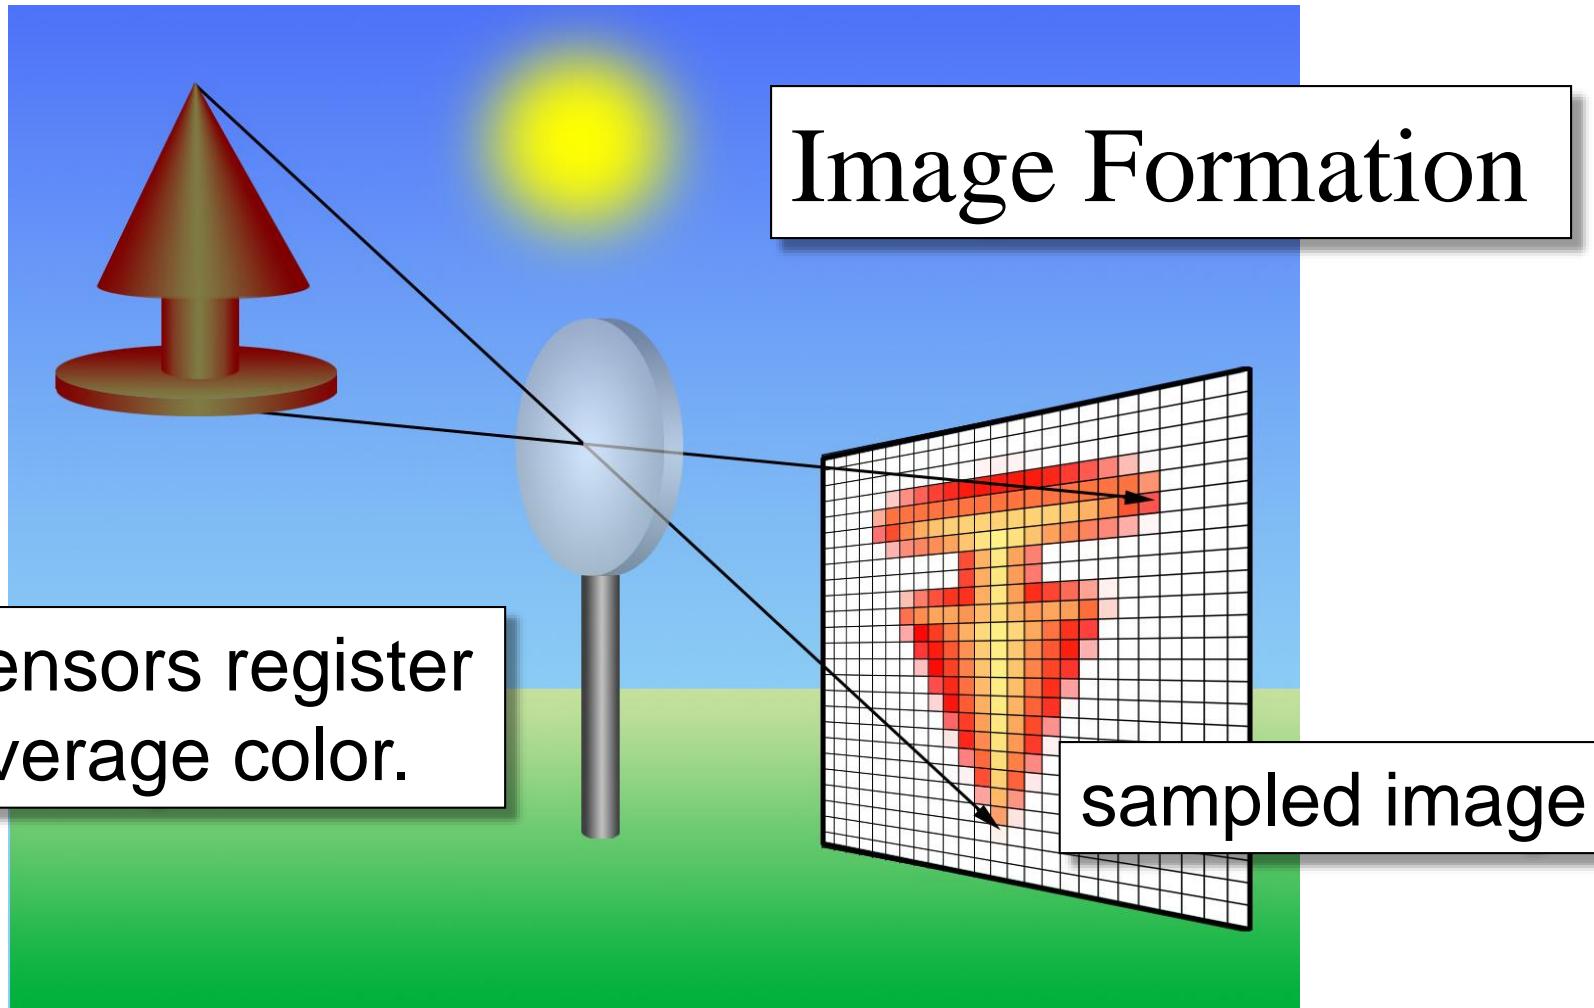


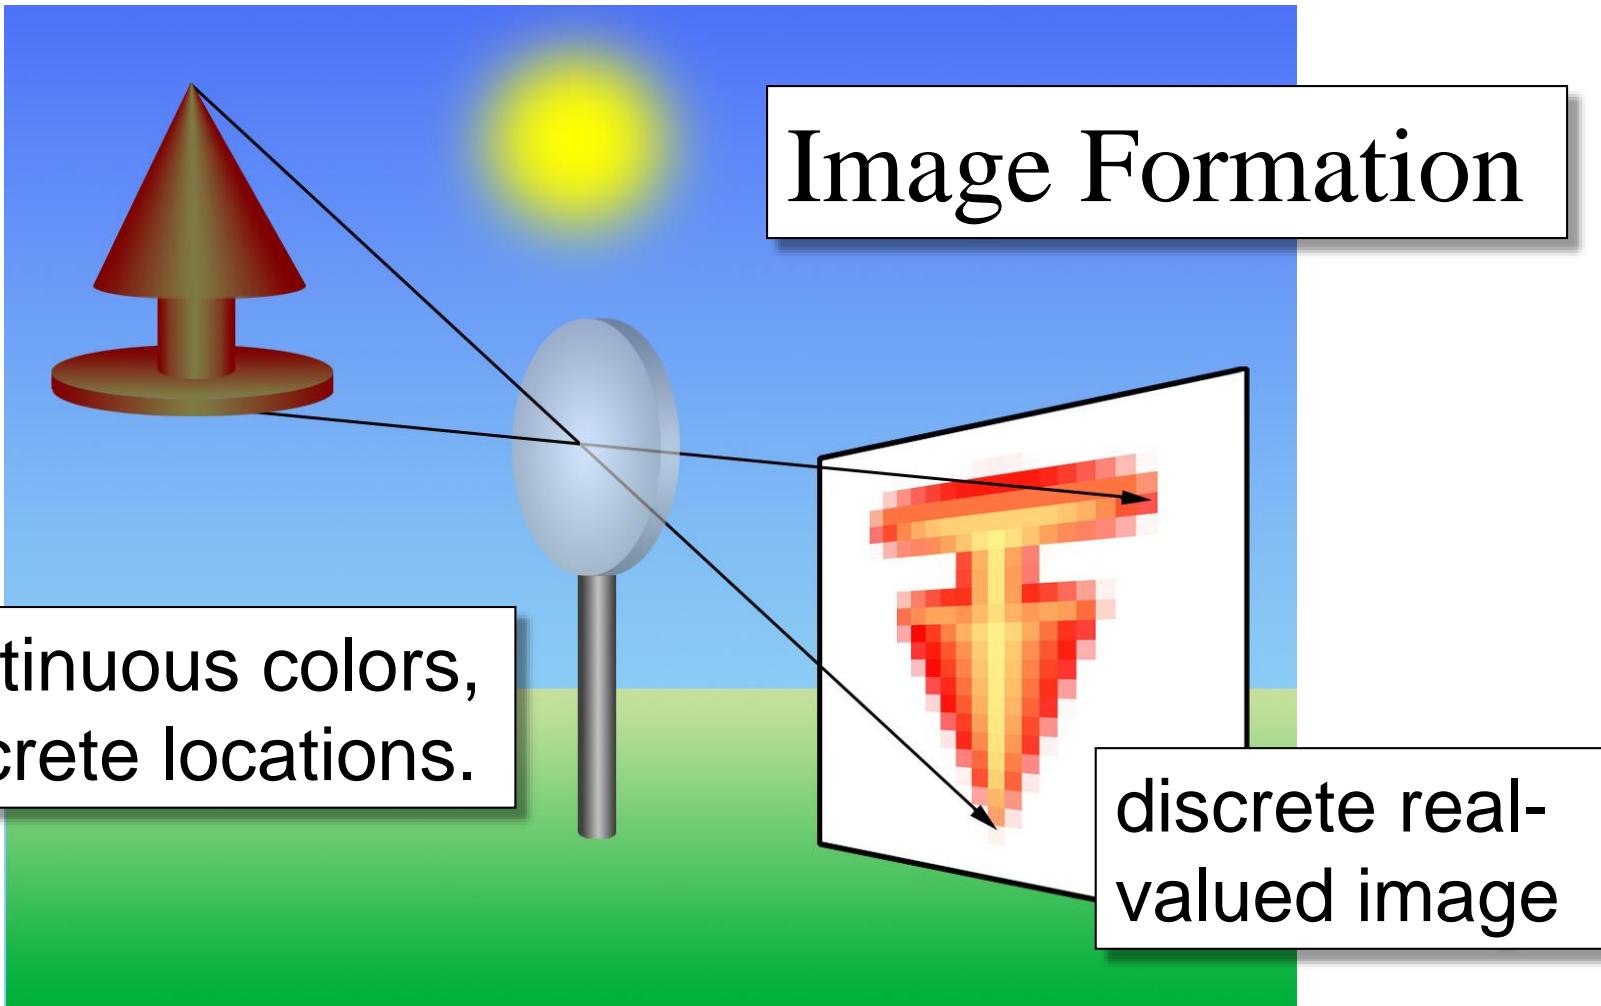






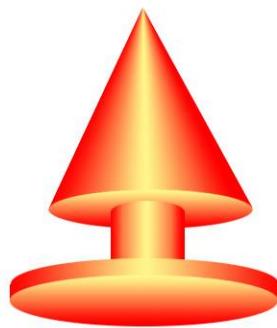




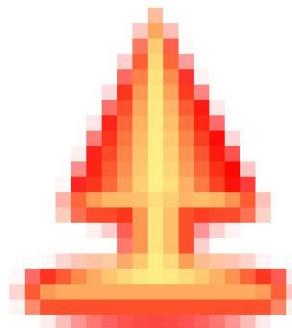




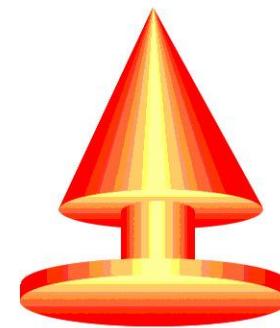
# Sampling and Quantization



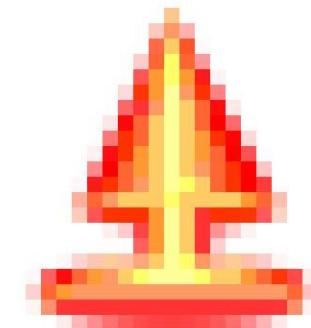
real image



sampled



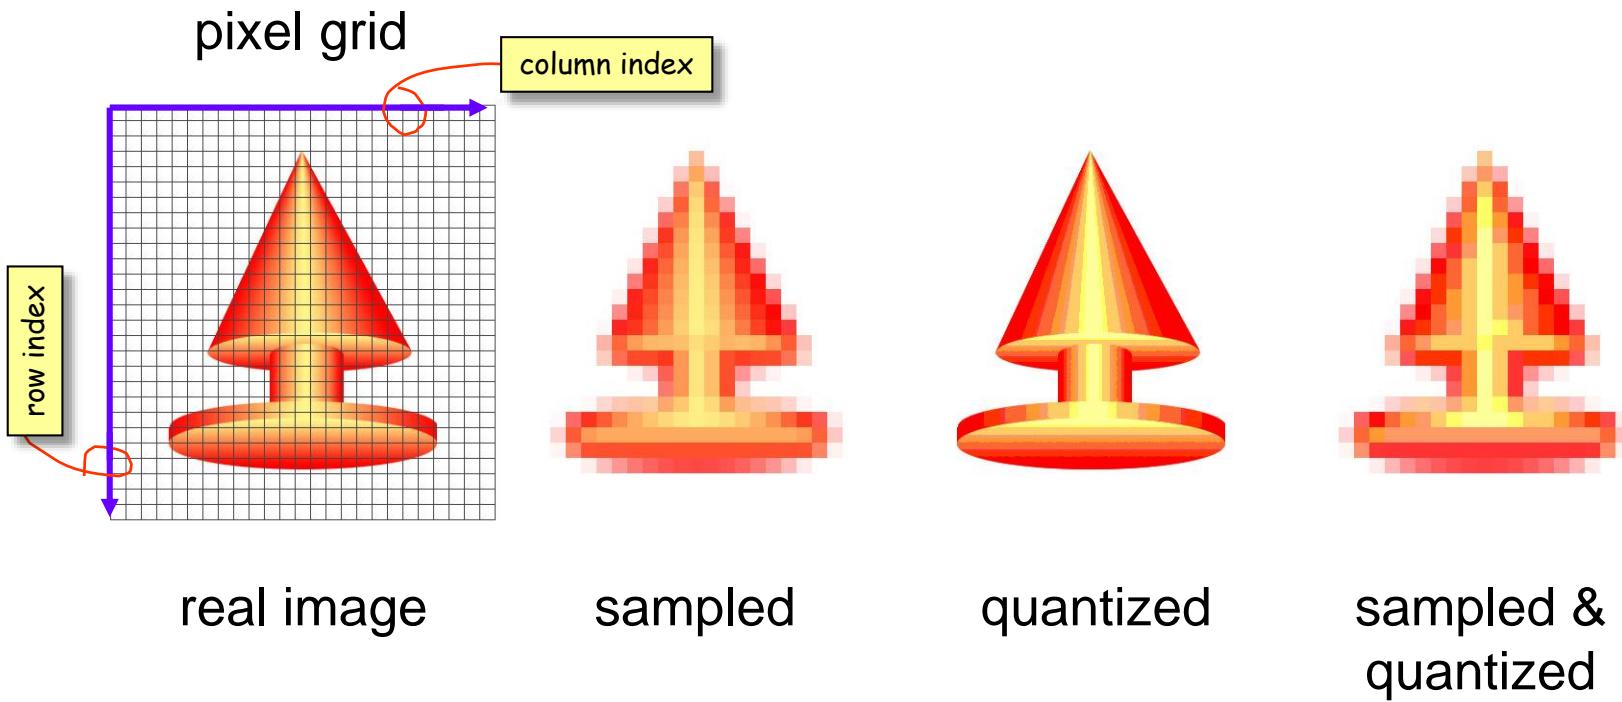
quantized



sampled &  
quantized



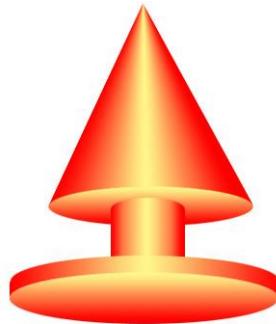
# Sampling and Quantization





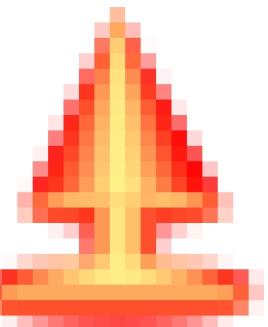
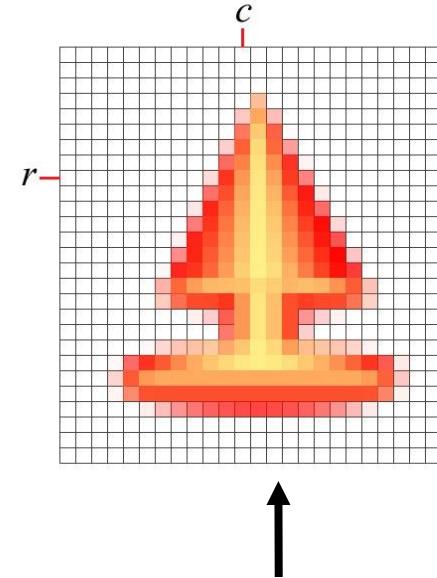
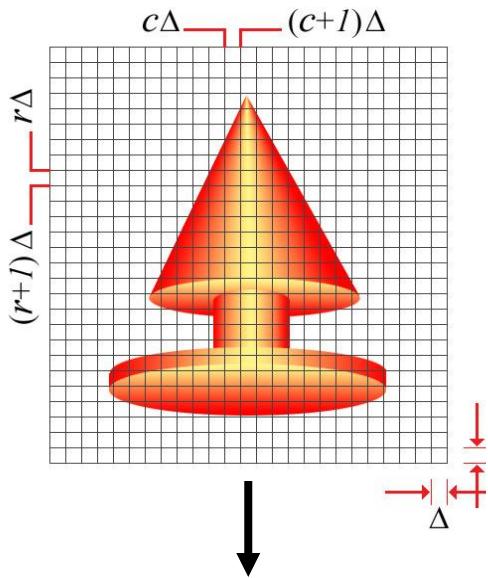
# Sampling

Take the average  
within each square.



$$\mathbf{I}_C(\rho, \chi)$$

continuous image



$$\mathbf{I}_S(r, c)$$

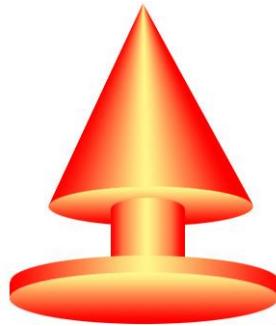
sampled image

$$\mathbf{I}_S(r, c) \frac{1}{\Delta^2} \int_{r\Delta}^{(r+1)\Delta} \int_{c\Delta}^{(c+1)\Delta} \mathbf{I}_C(\rho, \chi) \delta\rho \delta\chi$$



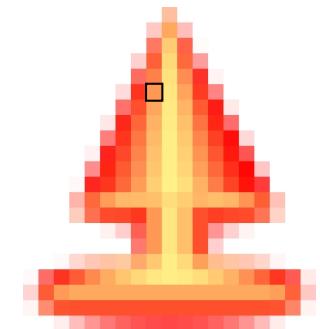
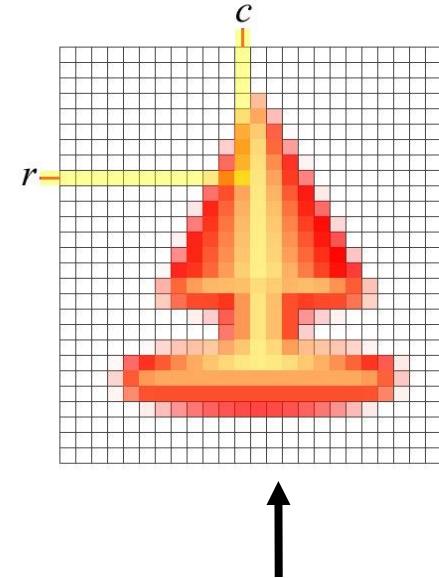
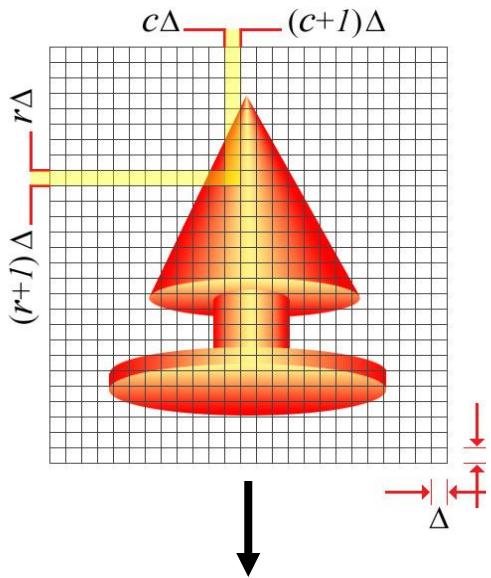
# Sampling

Take the average  
within each square.



$$\mathbf{I}_C(\rho, \chi)$$

continuous image



$$\mathbf{I}_S(r, c)$$

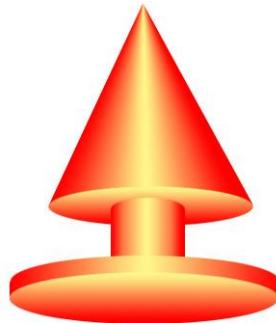
sampled image

$$\mathbf{I}_S(r, c) \frac{1}{\Delta^2} \int_{r\Delta}^{(r+1)\Delta} \int_{c\Delta}^{(c+1)\Delta} \mathbf{I}_C(\rho, \chi) \delta\rho \delta\chi$$



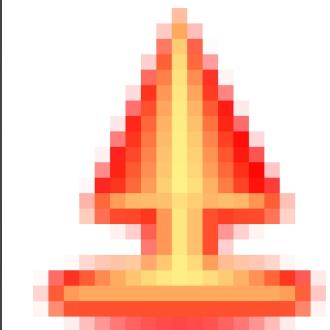
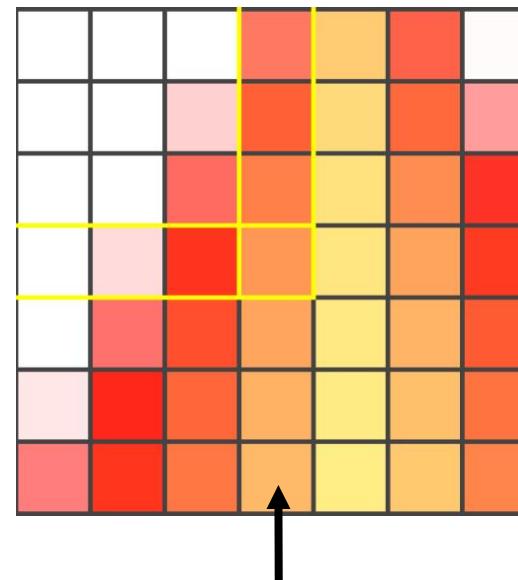
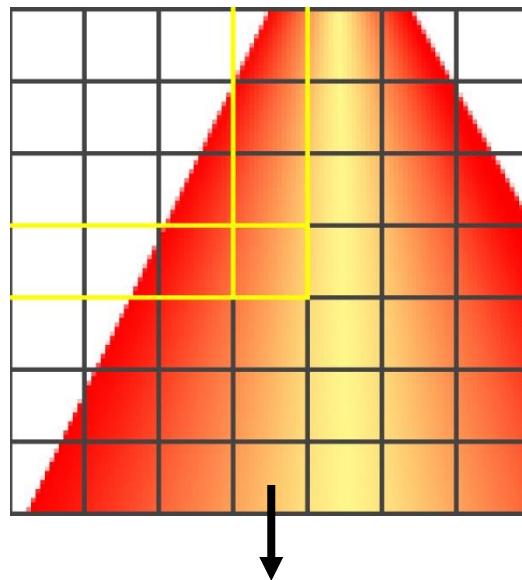
# Sampling

Take the average  
within each square.



$$\mathbf{I}_C(\rho, \chi)$$

continuous image



$$\mathbf{I}_S(r, c)$$

sampled image

$$\mathbf{I}_S(r, c) \frac{1}{\Delta^2} \int_{r\Delta}^{(r+1)\Delta} \int_{c\Delta}^{(c+1)\Delta} \mathbf{I}_C(\rho, \chi) \delta\rho \delta\chi$$



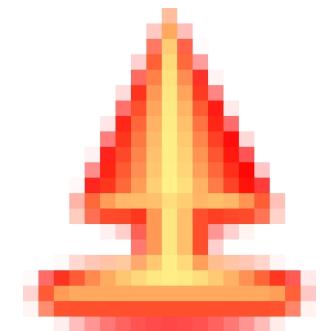
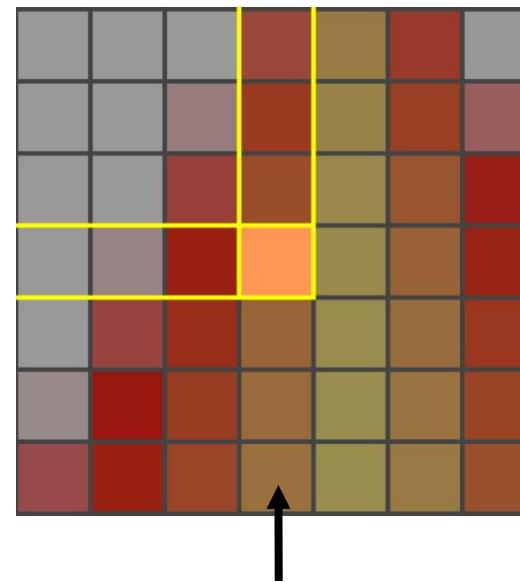
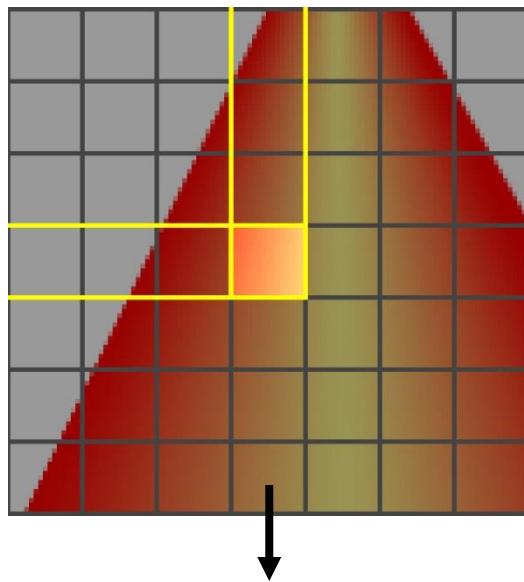
# Sampling

Take the average  
within each square.



$$\mathbf{I}_C(\rho, \chi)$$

continuous image



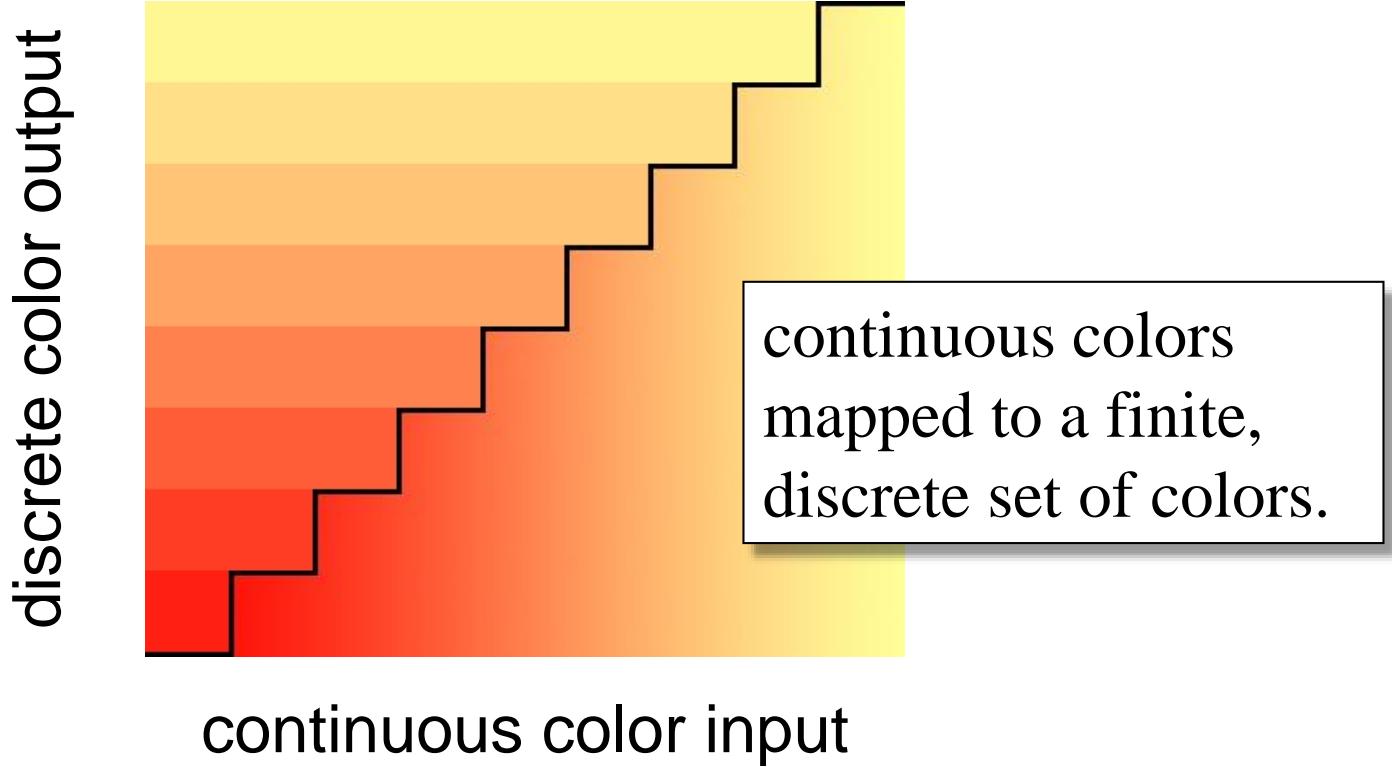
$$\mathbf{I}_S(r, c)$$

sampled image

$$\mathbf{I}_S(r, c) \frac{1}{\Delta^2} \int_{r\Delta}^{(r+1)\Delta} \int_{c\Delta}^{(c+1)\Delta} \mathbf{I}_C(\rho, \chi) \delta\rho \delta\chi$$



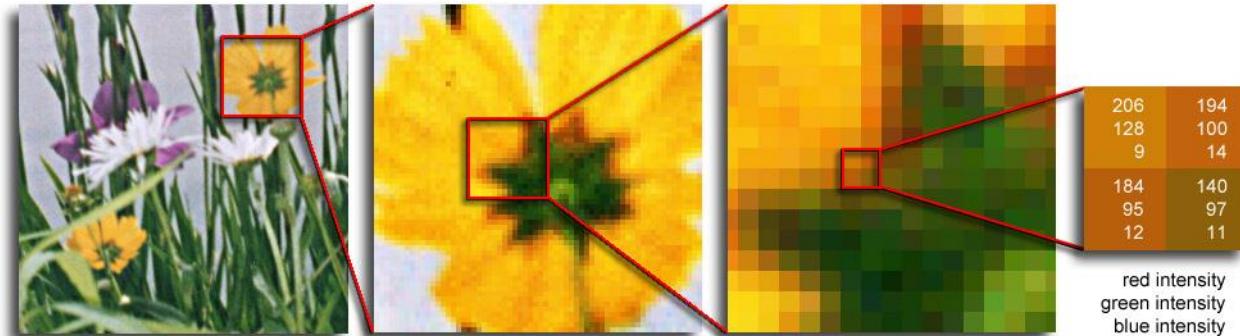
# Digital Image Formation: Quantization



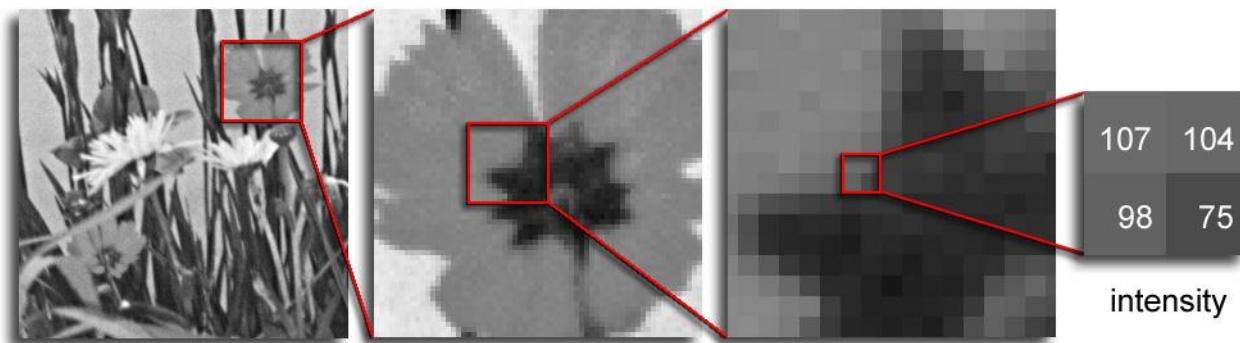


# Digital Image

a grid of squares, each of which contains a single color



each square is called a pixel (for *picture element*)





# Pixels

- A digital image,  $\mathbf{I}$ , is a mapping from a 2D grid of uniformly spaced discrete points,  $\{\mathbf{p} = (r,c)\}$ , into a set of positive integer values,  $\{\mathbf{I}(\mathbf{p})\}$ , or a set of vector values, e.g.,  $\{[\mathbf{R} \; \mathbf{G} \; \mathbf{B}]^T(\mathbf{p})\}$ .
- At each column location in each row of  $\mathbf{I}$  there is a value.
- The pair (  $\mathbf{p}$ ,  $\mathbf{I}(\mathbf{p})$  ) is called a “pixel” (for *picture element*).

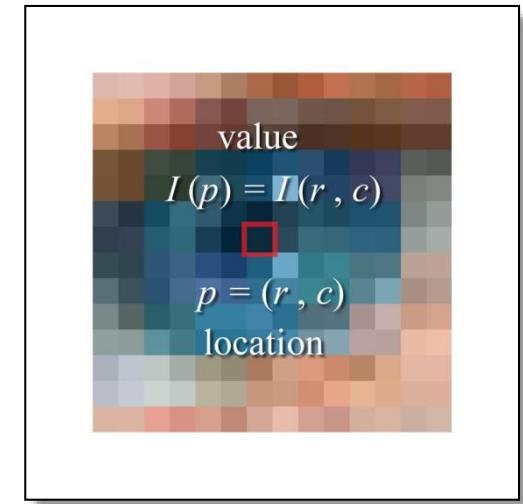
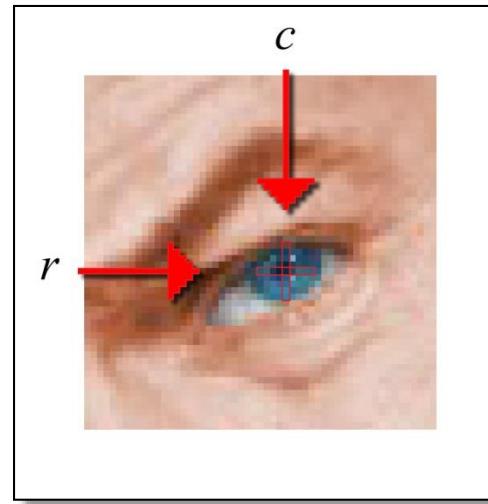
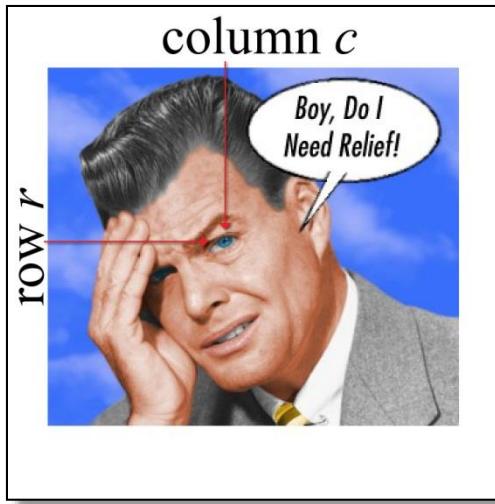


# Pixels

- $\mathbf{p} = (r,c)$  is the pixel location indexed by row,  $r$ , and column,  $c$ .
- $\mathbf{I}(\mathbf{p}) = \mathbf{I}(r,c)$  is the value of the pixel at location  $\mathbf{p}$ .
- If  $\mathbf{I}(\mathbf{p})$  is a single number then  $\mathbf{I}$  is monochrome.
- If  $\mathbf{I}(\mathbf{p})$  is a vector (ordered list of numbers) then  $\mathbf{I}$  has multiple bands (*e.g.*, a color image).



# Pixels



Pixel Location:  $\mathbf{p} = (r, c)$

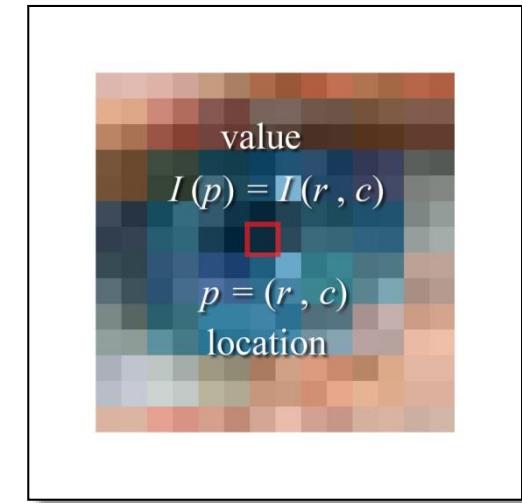
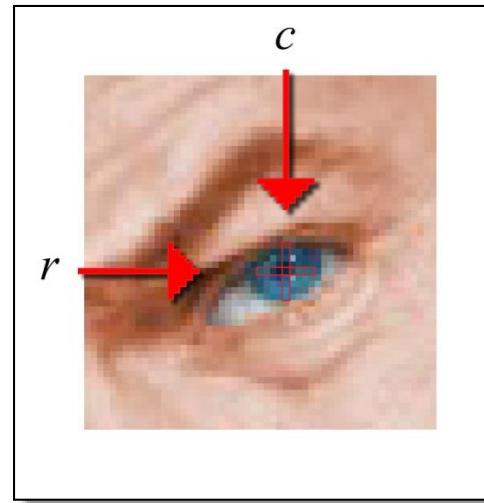
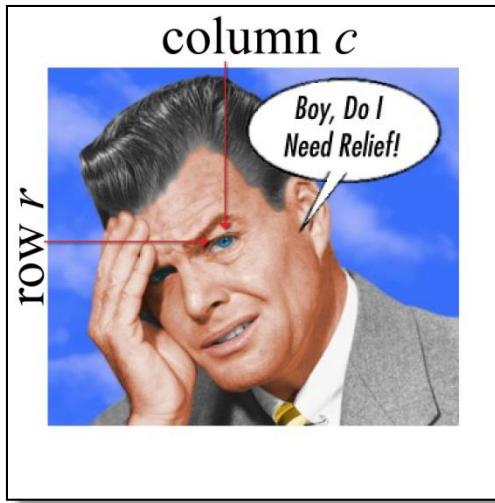
Pixel Value:  $I(\mathbf{p}) = I(r, c)$

Pixel : [  $\mathbf{p}$ ,  $I(\mathbf{p})$  ]



# Pixels

Pixel : [ p, I(p) ]



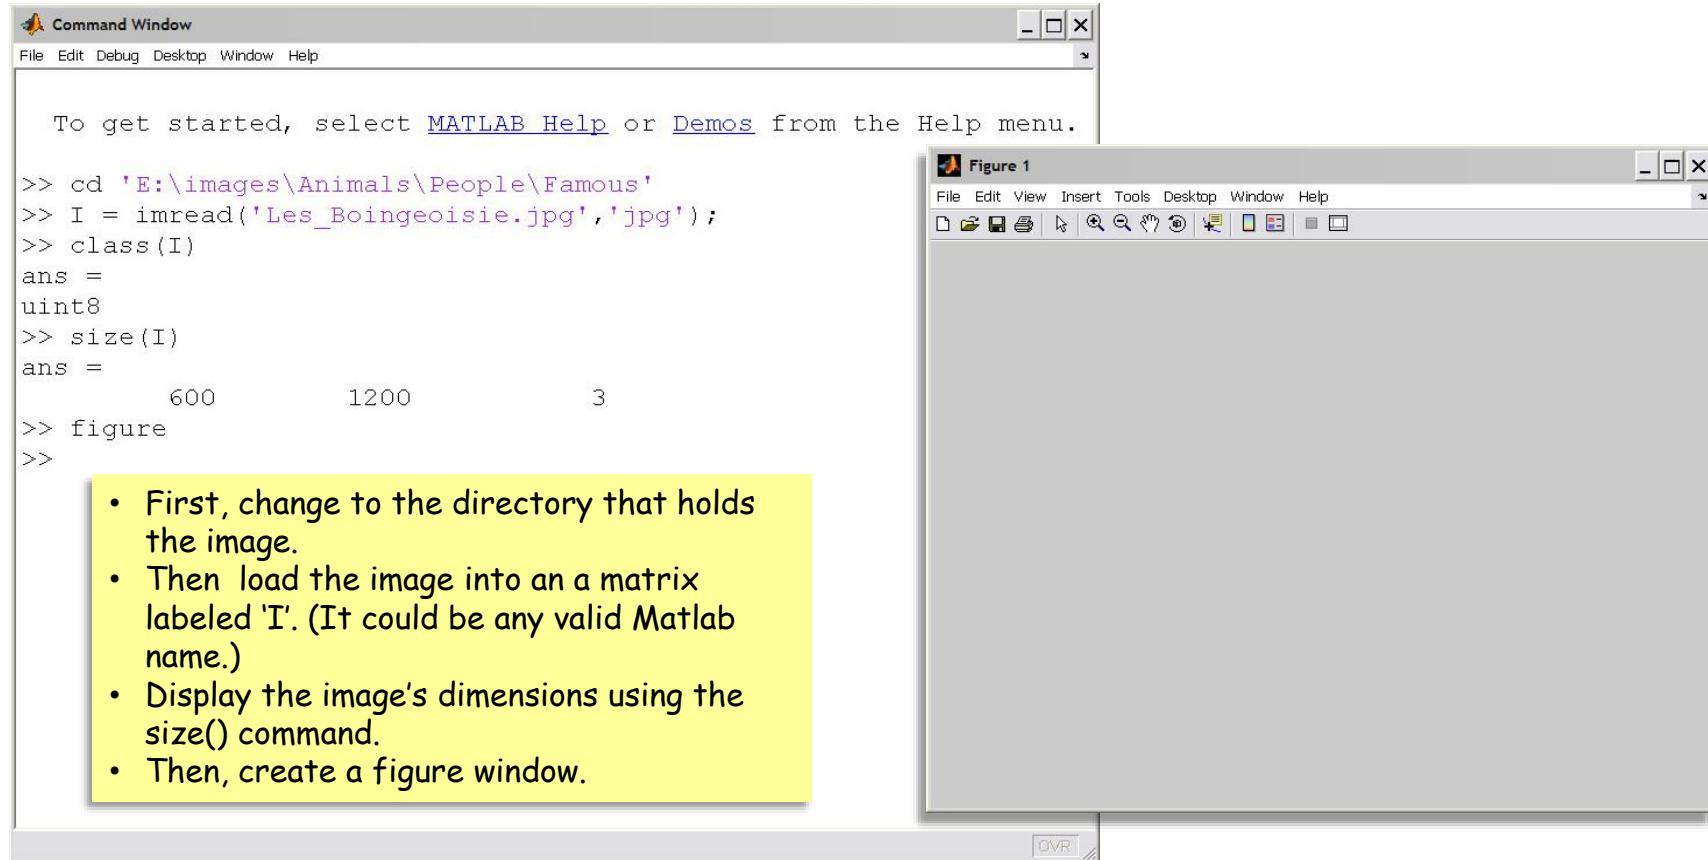
$$\begin{aligned} \mathbf{p} &= (r, c) \\ &= (\text{row \#}, \text{col \#}) \\ &= (272, 277) \end{aligned}$$

$$\mathbf{I}(\mathbf{p}) = \begin{bmatrix} \text{red} \\ \text{green} \\ \text{blue} \end{bmatrix} = \begin{bmatrix} 12 \\ 43 \\ 61 \end{bmatrix}$$



The next few slides demonstrate the results of the commands shown in the command window. Each successive slide starts at the end of the previous one.

## Read a Truecolor Image into Matlab



To get started, select [MATLAB Help](#) or [Demos](#) from the Help menu.

```
>> cd 'E:\images\Animals\People\Famous'  
>> I = imread('Les_Boingeoisie.jpg','jpg');  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
       600        1200         3  
>> figure  
>>
```

- First, change to the directory that holds the image.
- Then load the image into an a matrix labeled 'I'. (It could be any valid Matlab name.)
- Display the image's dimensions using the `size()` command.
- Then, create a figure window.



# Read a Truecolor Image into Matlab

```
Command Window
File Edit Debug Desktop Window Help
To get started, select MATLAB Help or
>> cd 'E:\images\Animals\People\Famous'
>> I = imread('Les_Boingeoisie.jpg', 'jp
>> class(I)
ans =
uint8
>> size(I)
ans =
       600      1200      3
>> figure
>>
```

On the `size()` command:

- If  $X$  is a scalar  $\text{size}(X)$  returns  $1 \ 1$
- If  $X$  is a 1-D row vector w/ 5 entries, it returns  $1 \ 5$
- If  $X$  is a 1-D column vec w/ 5 entries, it returns  $5 \ 1$
- If  $X$  is a 4-row by 5-column matrix, it returns  $4 \ 5$
- If  $X$  is an  $R$ -row by  $C$ -column, 1-band image, it returns  $R \ C$
- If  $X$  is an  $R$ -row by  $C$ -column, 3-band image, it returns  $R \ C \ 3$

E.g. Above,  $I$  is a 600-row by 1200-column by 3-band image.

On the `class()` command:

- If  $X$  is a matrix of 64-bit floating-point numbers (Matlab's default data class) `class(X)` returns 'double' (without the quotes).
- If  $X$  is a matrix of 8-bit unsigned integers (like most common images) `class(X)` returns 'uint8'.
- Matlab has other data classes, too. Search for 'Fundamental MATLAB Classes' in the help window for a description of them.

Related to the `size()` command:  
`length(X)`

- returns the size of the largest dimension of  $X$ .
- `length(X(:))` returns the number of elements in  $X$ .



# Read a Truecolor Image into Matlab



To get started, select [MATLAB Help](#) or [Demos](#) from the Help menu.

```
>> cd 'E:\images\Animals\People\Famous'  
>> I = imread('Les_Boingeoisie.jpg','jpg');  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
       600      1200         3  
>> figure  
>> image(I)  
>> title('Les Boingeoisie: The Boing-Boing Bloggers')  
>> xlabel('Photo: Bart Nagel, 2006, www.bartnagel.com')  
>>
```

- Display the image in the figure window using `image()`.
- Place a title at the top.
- Put a label under the x-axis using `xlabel()`.
- `ylabel()` puts a label to the left of the y-axis, oriented 90° from horizontal

Matlab scales the image to fit the figure window so both the images size and aspect ratio are usually incorrect.





# Read a

To get started:

```
>> cd 'E:\image'
>> I = imread('LesBoingBoingers.jpg')
>> class(I)
ans =
uint8
>> size(I)
ans =
      600
      600
>> figure
>> image(I)
>> title('Les Boing-Boingers')
>> xlabel('Photo by Bart Nagel')
>> truesize
>>
```

truesize makes the figure window fit the image – unless the image is too large for the display in which case it is resized to fit but with the correct aspect ratio.

Figure 1

File Edit View Insert Tools Desktop Window Help

Command Window

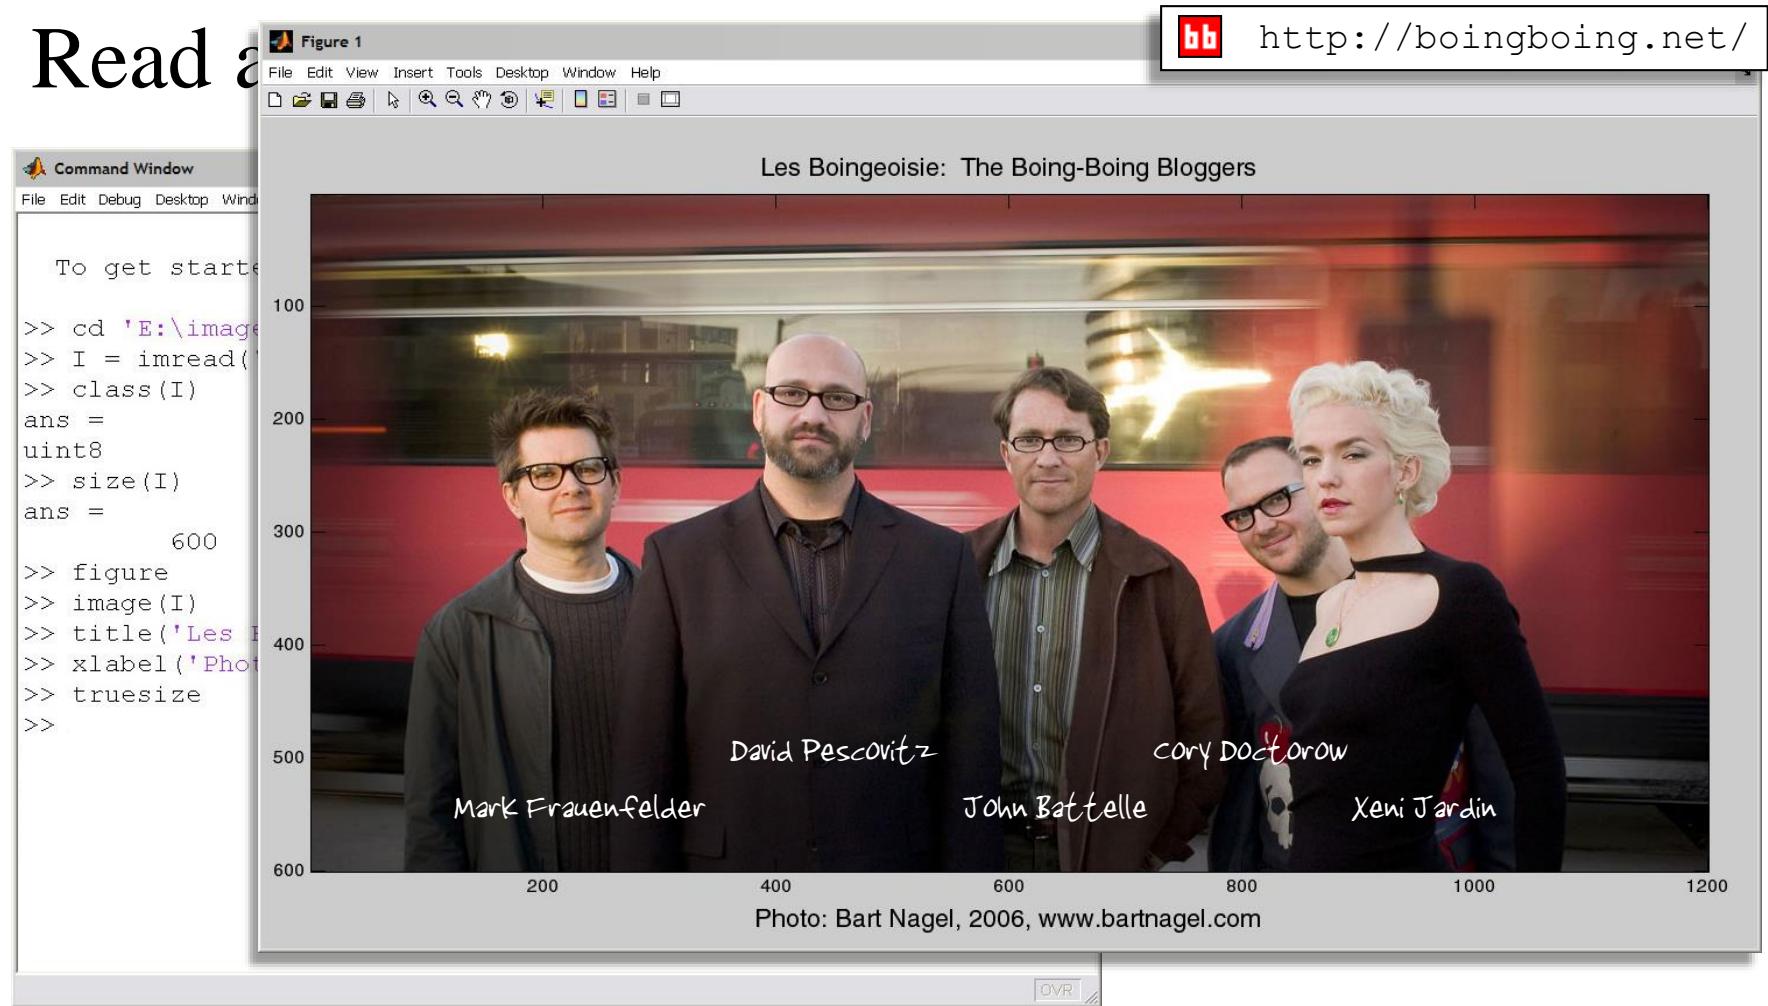
File Edit Debug Desktop Window

Les Boing-Boingers: The Boing-Boing Bloggers

Photo: Bart Nagel, 2006, [www.bartnagel.com](http://www.bartnagel.com)

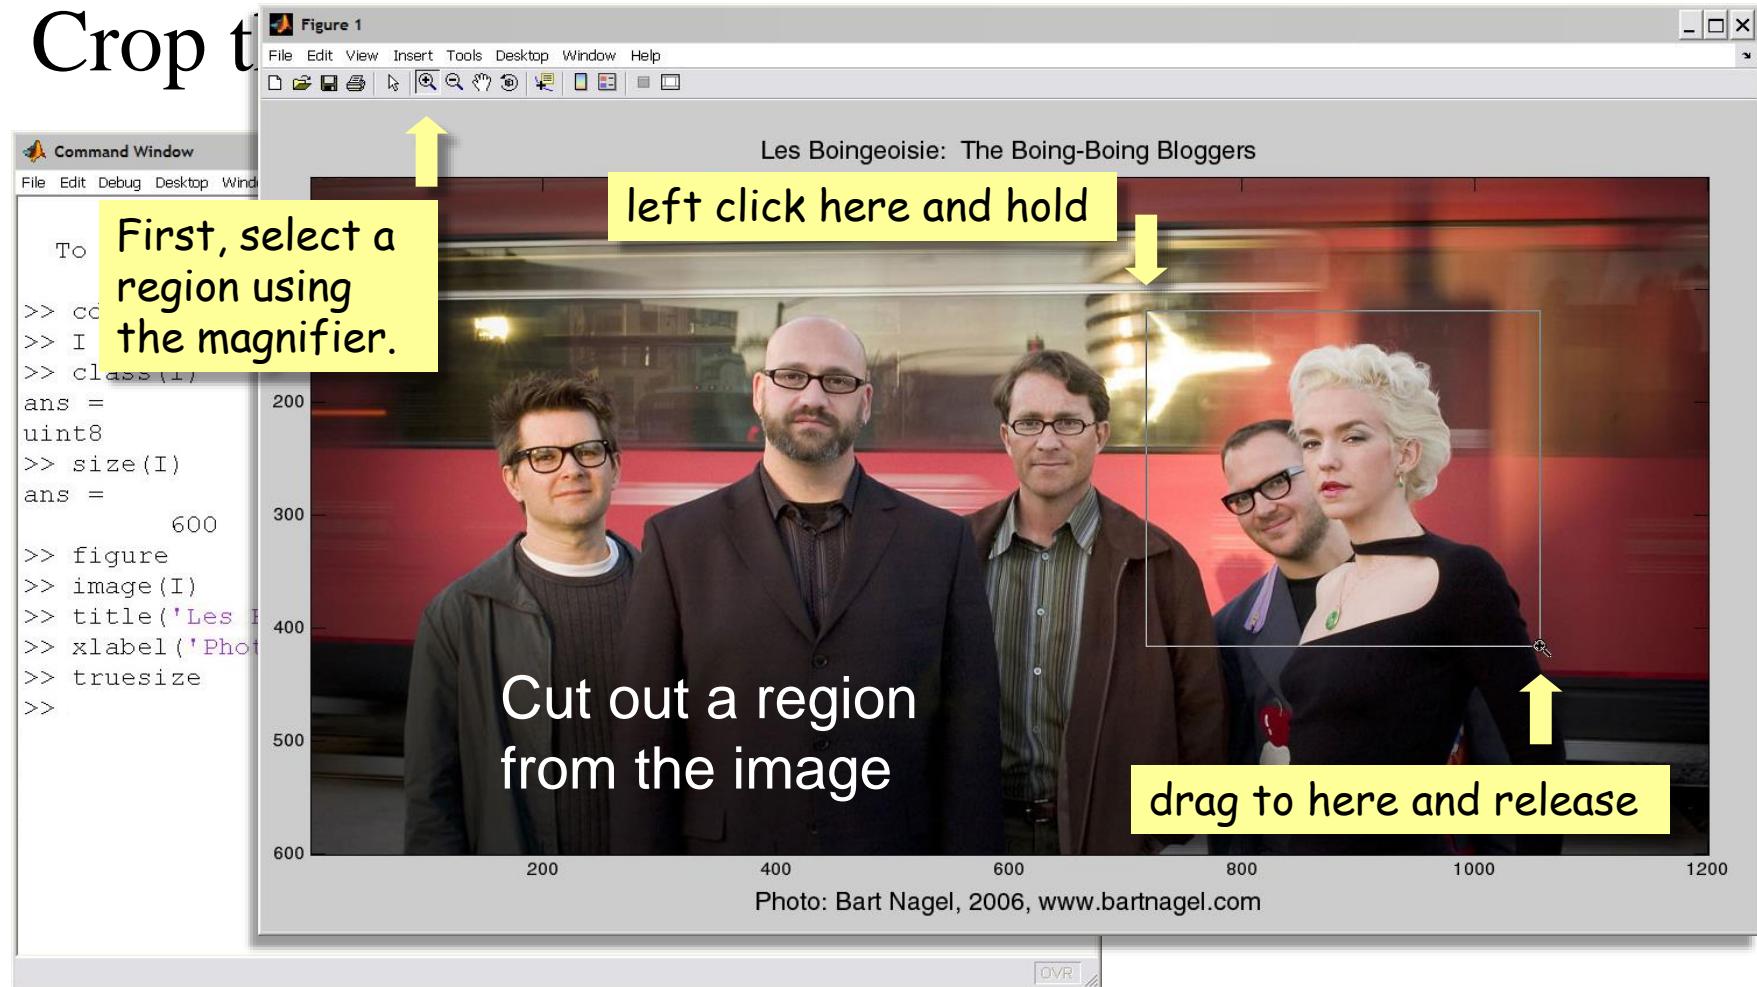


# Read a





# Crop t





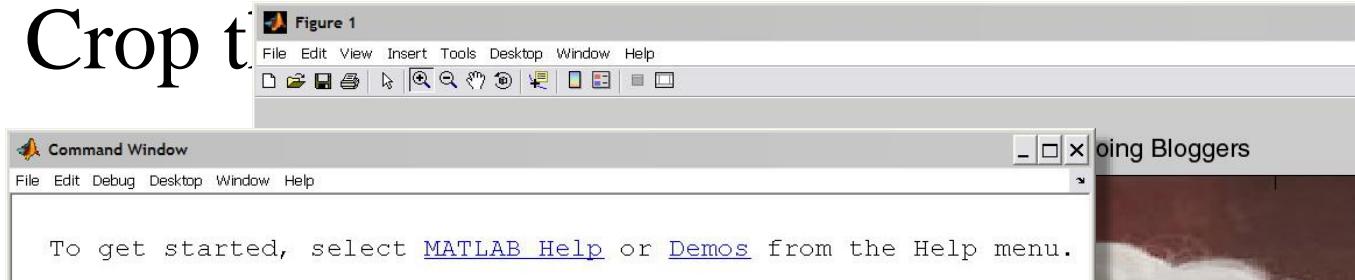
# Crop t

From this close-up  
we can estimate  
the coordinates of  
the region:





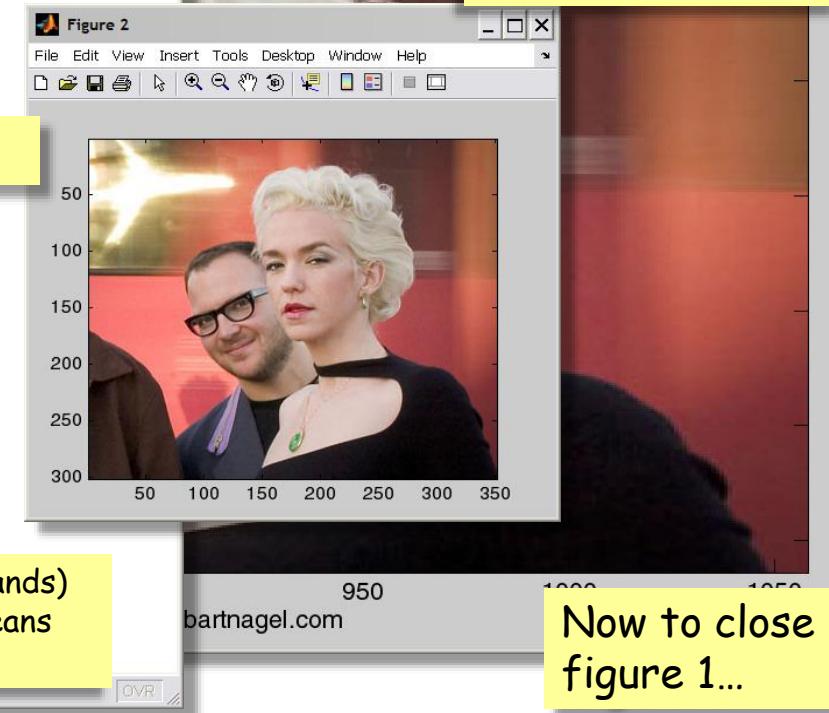
## Crop t



```
>> cd 'E:\images\Animals\People\Famous'  
>> I = imread('Les_Boingeoisie.jpg','jpg');  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
       600      1200         3  
>> figure  
>> image(I)  
>> title('Les Boingeoisie: The Boing-Boing Bloggers')  
>> xlabel('Photo: Bart Nagel, 2006, www.bartnagel.com')  
>> truesize  
>> J = I(125:425,700:1050,:);  
>> figure  
>> image(J)  
>> truesize  
>>
```

$J = I(\text{first row:last row}, \text{first column:last column}, \text{all bands})$   
':' alone means 'all bands' here. In the 1<sup>st</sup> position, ':' means  
'all rows'; in the 2<sup>nd</sup> it means 'all columns'.

Here it is:



$I(:)$  means all I's elements in a vertical list: 1<sup>st</sup> col followed by 2<sup>nd</sup>, etc. until the 1<sup>st</sup> band is catenated.  
Same procedure is applied to remaining bands for one long list.

Now to close figure 1...



# Crop t

Figure 1

File Edit View Insert Tools Desktop Window Help

Command Window

To get started

```
>> cd 'E:\image'
>> I = imread('Les Boingeoisie.jpg')
>> class(I)
ans =
uint8
>> size(I)
ans =
600 600
>> figure
>> image(I)
>> title('Les Boingeoisie: The Boing-Boing Bloggers')
>> xlabel('Photo by Bart Nagel')
>> truesize
>> J = I(125:415, 750:950);
>> figure
>> image(J)
>> truesize
>> figure(1)
```

Les Boingeoisie: The Boing-Boing Bloggers

Photo: Bart Nagel, 2006, www.bartnagel.com

DVR

... Bring figure 1 to the front using the figure() command,

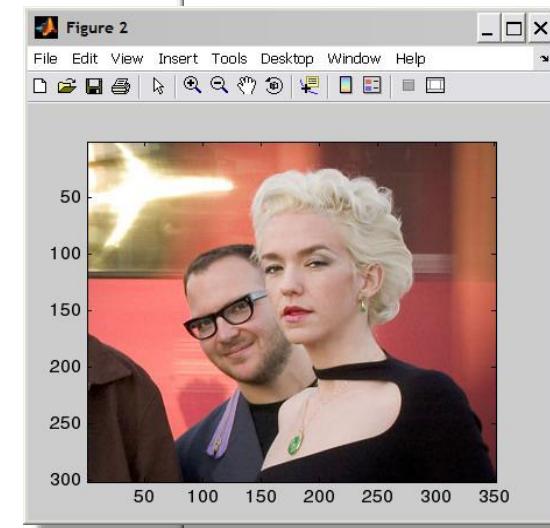


# Crop the Image

Command Window

```
To get started, select MATLAB Help or Demos from the Help menu.
```

```
>> cd 'E:\images\Animals\People\Famous'  
>> I = imread('Les_Boingeoisie.jpg','jpg');  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
       600      1200         3  
>> figure  
>> image(I)  
>> title('Les Boingeoisie: The Boing-Boing Bloggers')  
>> xlabel('Photo: Bart Nagel, 2006, www.bartnagel.com')  
>> truesize  
>> J = I(125:425,700:1050,:);  
>> figure  
>> image(J)  
>> truesize  
>> figure(1)  
>> close  
>>
```



then type 'close'  
at the prompt.



# Read a Colormapped Image into Matlab

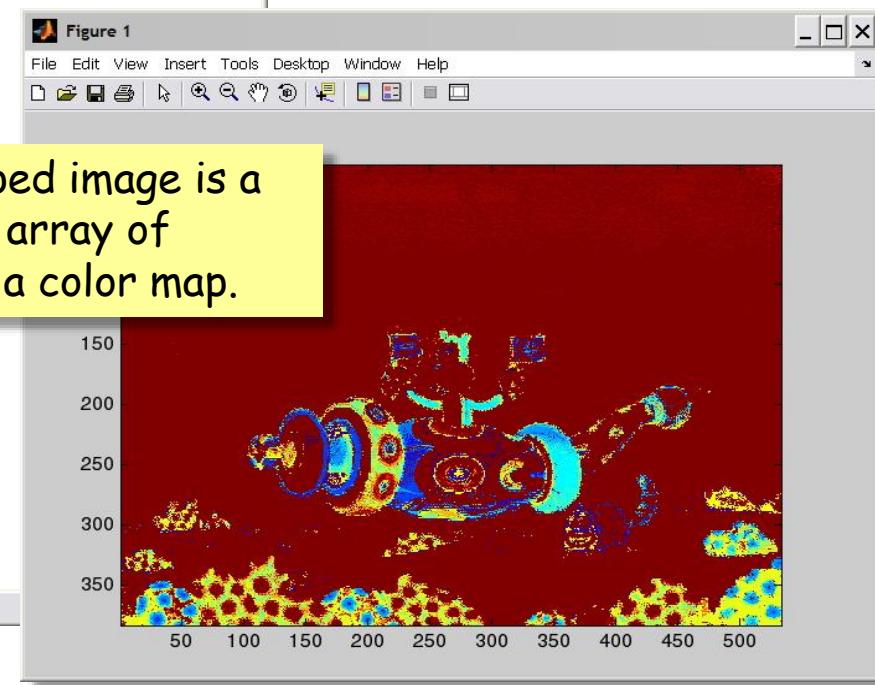
To get started, select [MATLAB Help](#) or [Demos](#) from the Help menu.

```
>> cd 'D:\classes\EECE253\Fall 2006\graphics\matlab intro'  
>> [I,cmap] = imread('Jim Woodring - PlusMinus.gif','gif');  
>> figure  
>> image(I)  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
383 533  
>> |
```

Matlab assumes that a 1-band, 8-bit / pixel image is colormapped. Initial display uses a default colormap which is highly likely to be incorrect for any given image.

A colormap is an  $n \times 3$  table of color values. 8-bits per pixel implies  $n = 256$  (at most) colors out of  $256^3 = 16777216$

A colormapped image is a rectangular array of indices into a color map.





# Read a Colormapped Image into Matlab

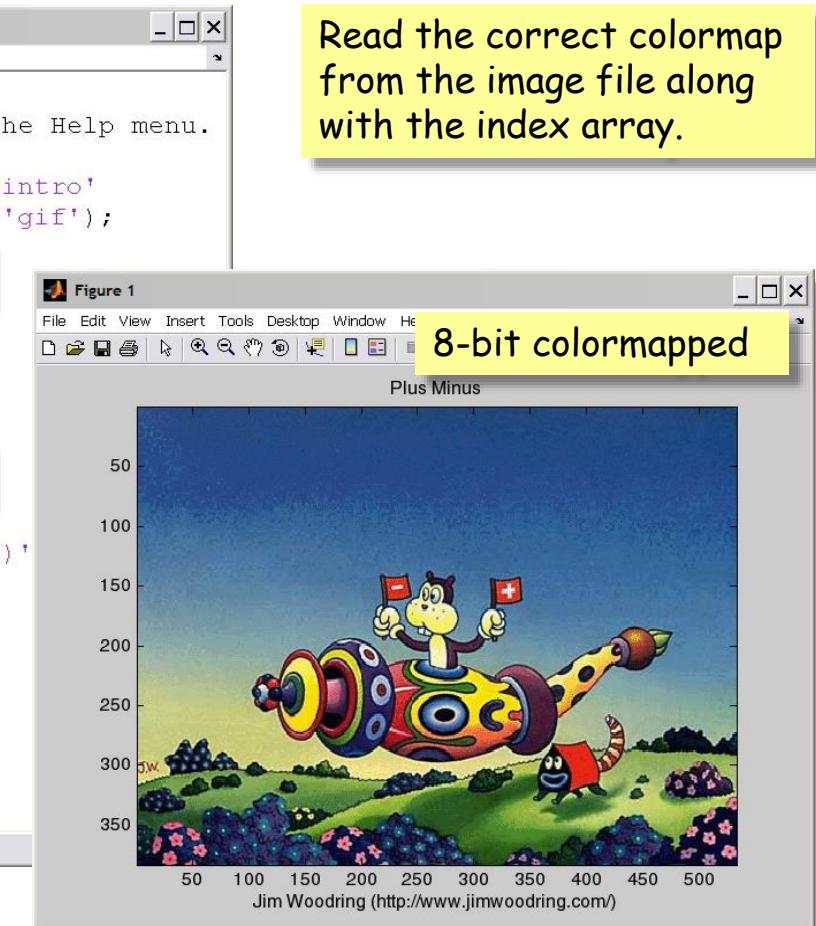
To get started, select [MATLAB Help](#) or [Demos](#) from the Help menu.

```
>> cd 'D:\classes\EECE253\Fall 2006\graphics\matlab intro'  
>> [I,cmap] = imread('Jim Woodring - PlusMinus.gif','gif');  
>> figure  
>> image(I)  
>> class(I)  
ans =  
uint8  
>> size(I)  
ans =  
    383   533  
>> colormap(cmap)  
>> title('Plus Minus');  
>> xlabel('Jim Woodring (http://www.jimwoodring.com/)');  
>> truesize  
>>
```

I = index array

cmap = colormap

Display the index array using `image()`. Then load the colormap using `colormap()`.





# Colormapped vs. Truecolor in Matlab

To get started, select [MATLAB Help](#) or [Demos](#) from the Help menu.

```
>> cd 'D:\classes\EECE253\Fall 2006\graphics\matlab intro'
>> [I,cmap] = imread('Jim Woodring - PlusMinus.gif','gif');
>> figure
>> image(I)
>> class(I)
ans =
uint8
>> size(I)
ans =
    383    533
>> colormap(cmap)
>> title('Plus Minus');
>> xlabel('Jim Woodring (http://www.jimwoodring.com)');
>> truesize
>> T = imread('Jim Woodring - PlusMinus.jpg','jpg')
>> figure
>> image(T)
>> truesize
>> |
```

**Figure 2**

24-bit truecolor

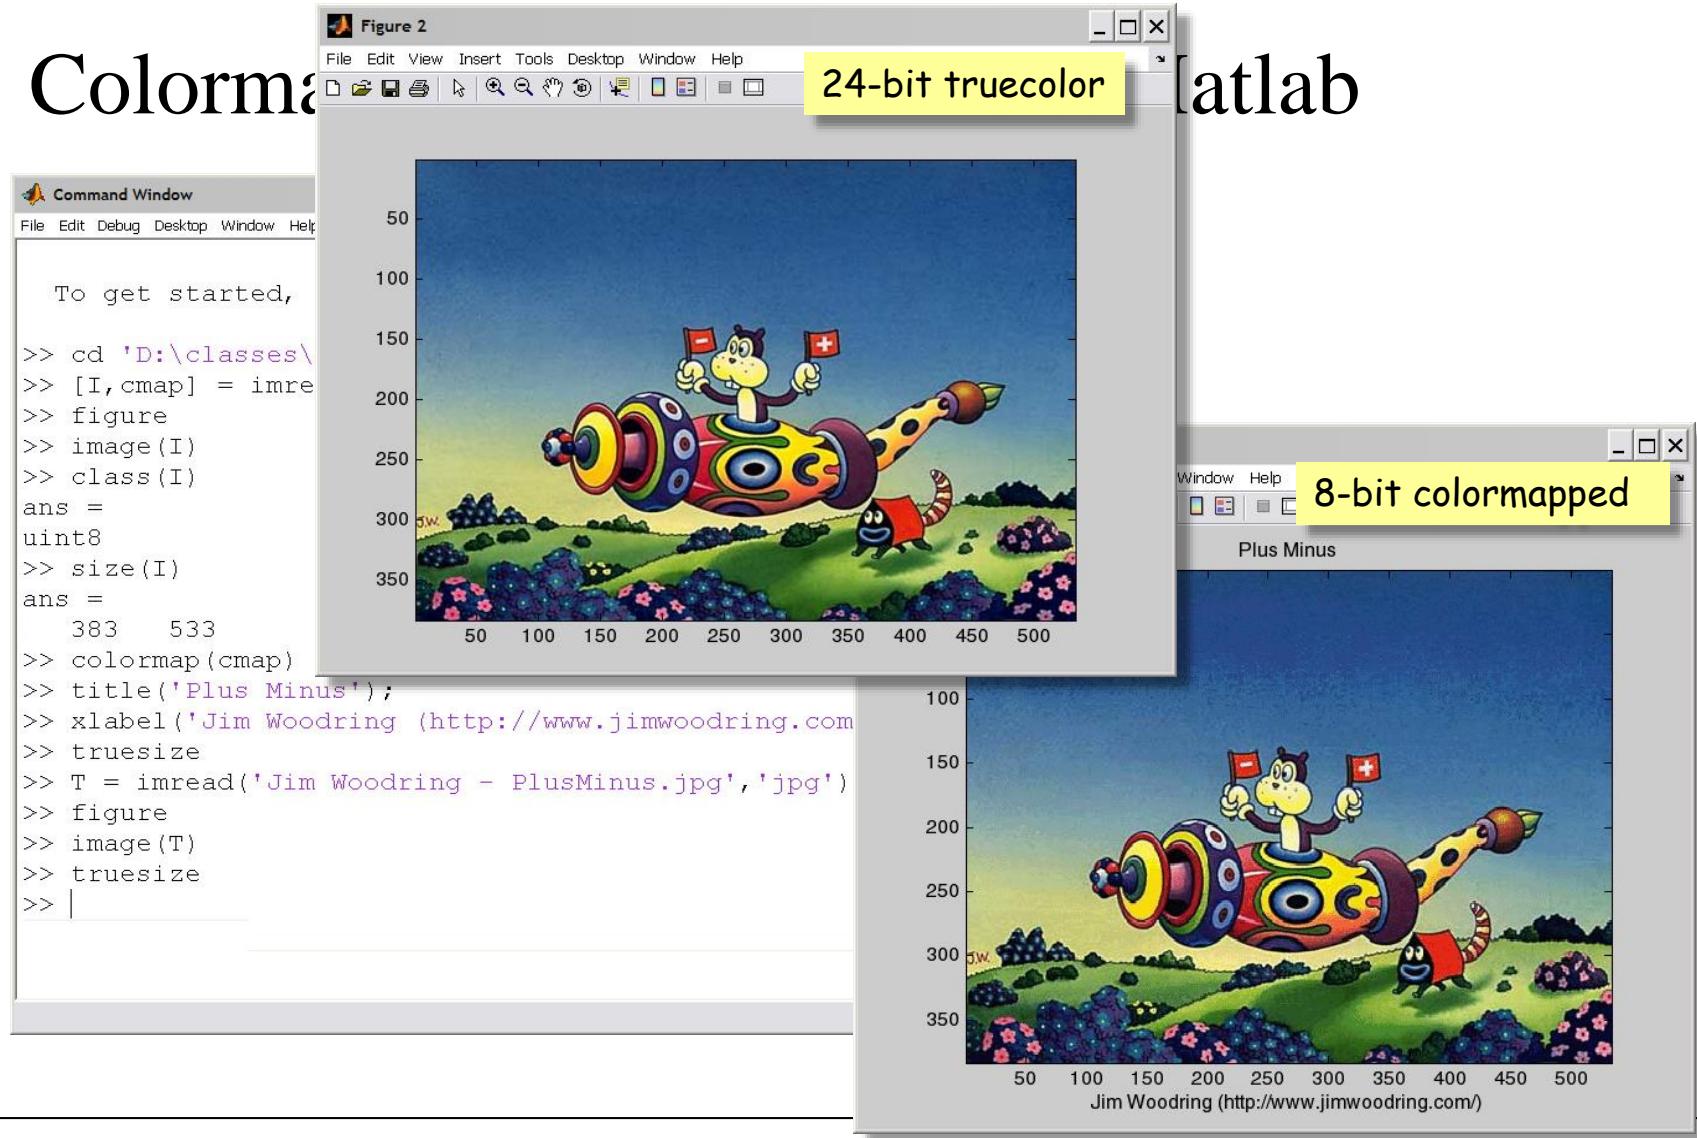
http://www.jimwoodring.com)'."/>

50 100 150 200 250 300 350 400 450 500

50 100 150 200 250 300 350 400 450 500

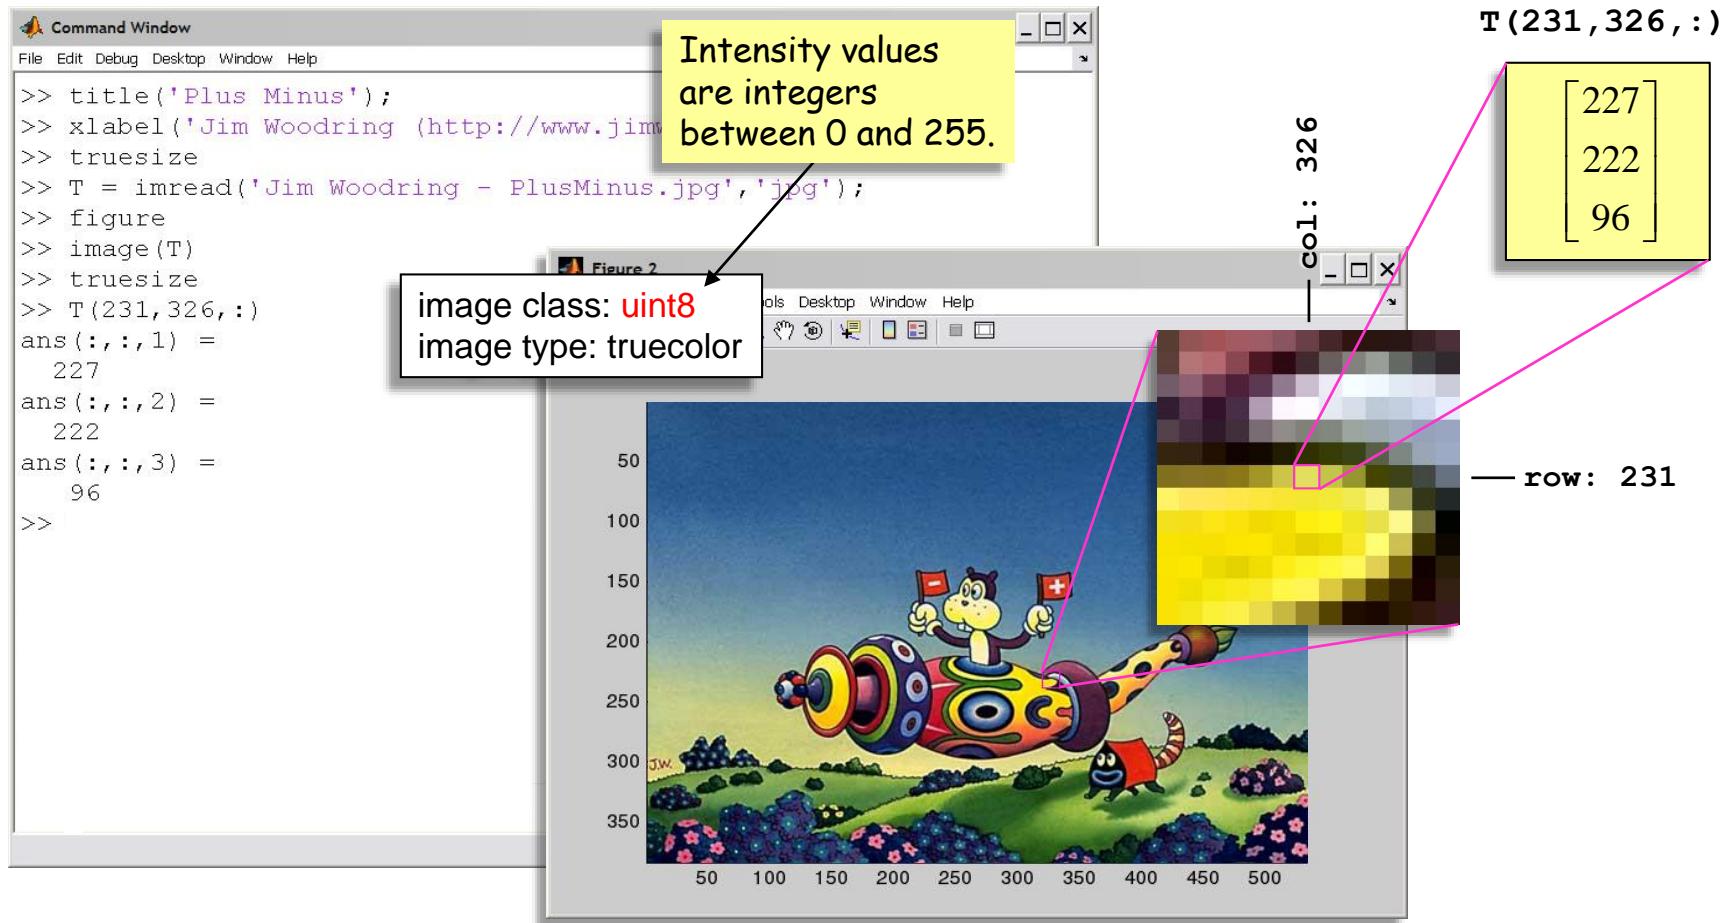


# Colormapping



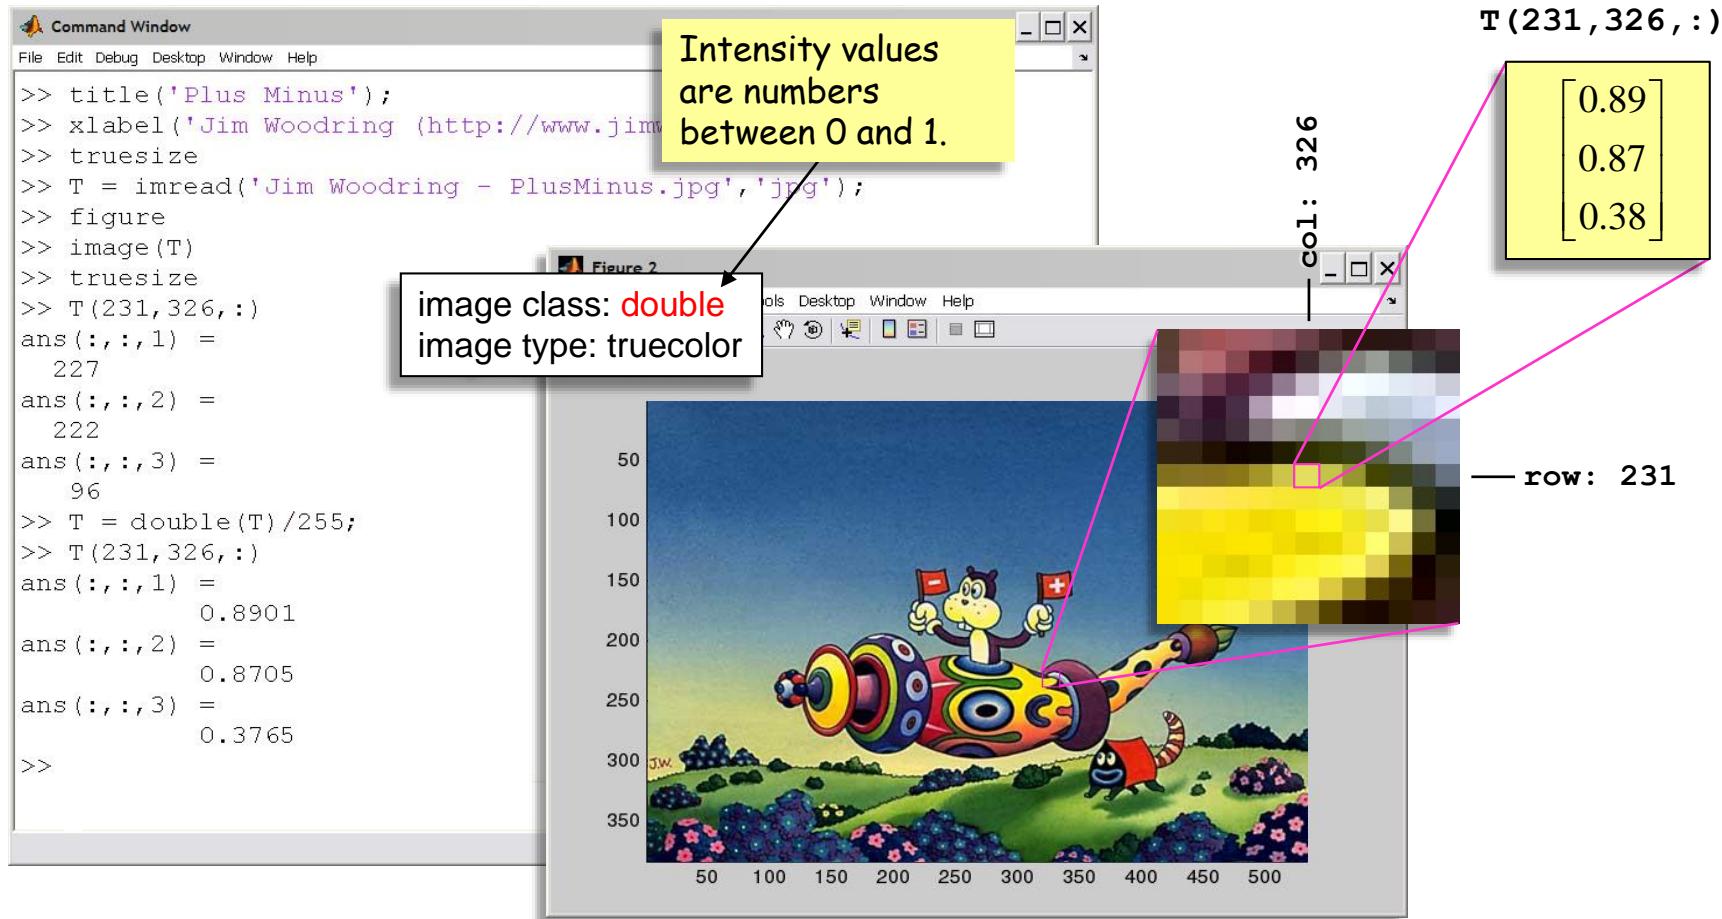


# Colormapped vs. Truecolor in Matlab





# Colormapped vs. Truecolor in Matlab





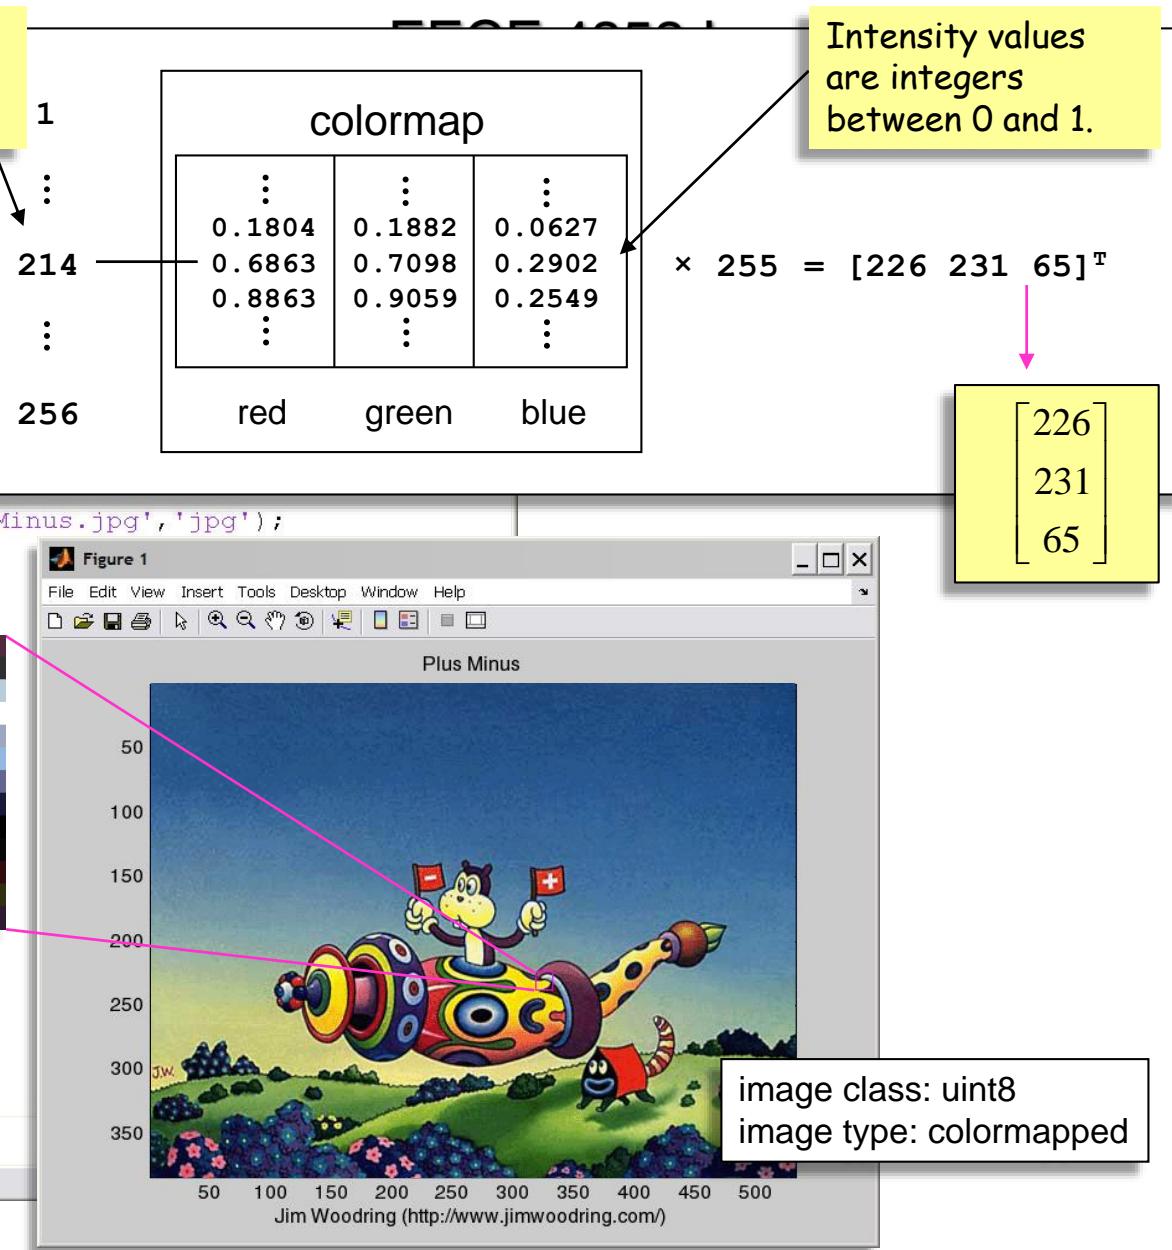
# Color

Number at pixel location is an index into a colormap.

```
>> title('Plus Minus')
>> xlabel('Jim Woodring')
>> truesize
>> T = imread('Jim Woodring - PlusMinus.jpg', 'jpg');
>> figure
>> image(T)
>> truesize
>> T(231,326,:)
ans(:,:,1) =
227
ans(:,:,2) =
222
row: 231
ans(:,:,3) =
96
>> I(231,326,:)
ans =
214
>> cmap(214,:)
ans =
0.8863    0.9059    0.2549
>> round(255*cmap(214,:))
ans =
226    231     65
>>
```

Intensity values are integers between 0 and 1.

226  
231  
65





## Example truecolor and colormapped images



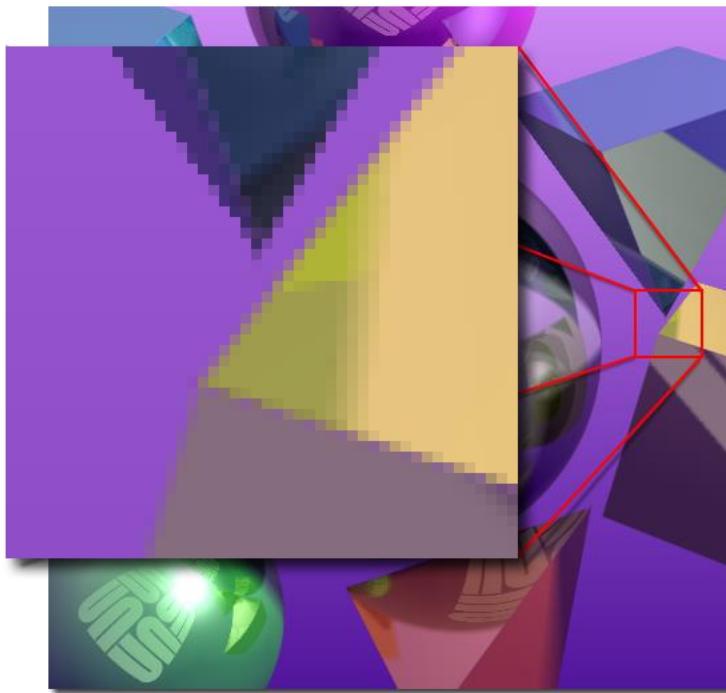
24-bit truecolor



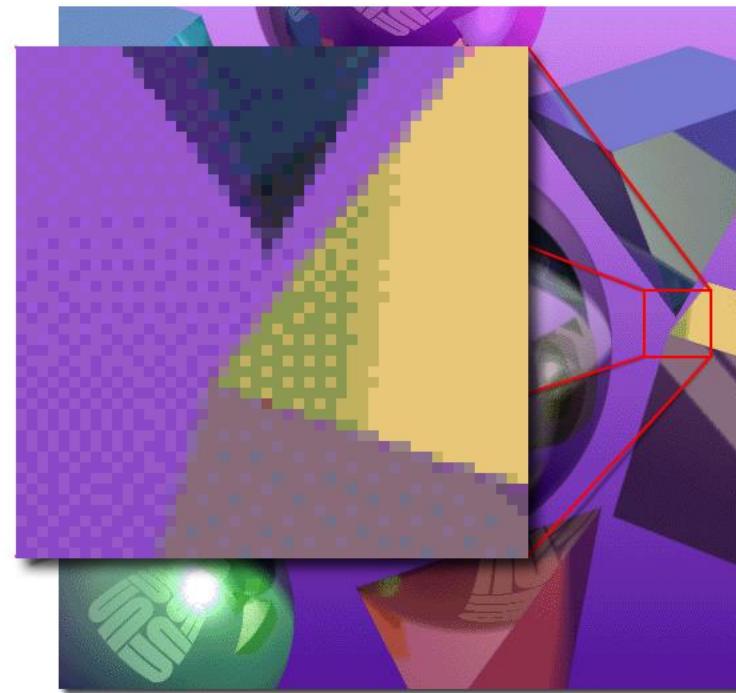
8-bit colormapped to 24 bits



# Example truecolor and colormapped images



24-bit truecolor



8-bit colormapped to 24 bits



# Colormapped Image, Indices, & Color Map

```
>> [M, CMAP] = imread('button_mapped.bmp', 'bmp');
```

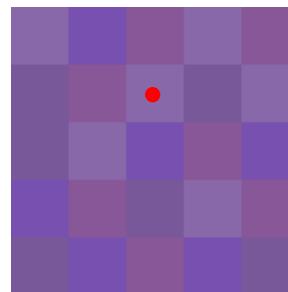
Indices contained in M(254:258,254:258)

:	:	:	:	:	:	
...	111	121	48	111	48	...
...	110	48	111	110	111	...
...	110	111	121	48	121	...
...	121	48	110	111	48	...
...	110	121	48	121	110	...
:	:	:	:	:	:	



actual values in CMAP(109:113,:)

:	R :	G :	B :
109	0.6588	0.4706	0.8471
110	0.2196	0.1569	0.2824
111	0.4706	0.3451	0.5961
112	0.5333	0.4078	0.6588
113	0.2824	0.2196	0.3451
:	:	:	:



255\*CMAP(109:113,:)

:	R:	G:	B:
109	168	120	216
110	56	40	72
111	120	88	152
112	136	104	168
113	72	56	88
:	:	:	:

Last  
3  
cols.  
only



# How to convert a colormapped image to true color

M is a  $512 \times 512 \times 1$ , 8-bit image.  
It has 262,144 pixels.  
Each pixel has a value between 0 & 255.

cmap is the colormap that is stored in 'button\_mapped.bmp' along with image. cmap is a  $256 \times 3$  type-double matrix, each row of which lists a color in terms of its R, G, & B intensities, which are given as fractions between 0 and 1.

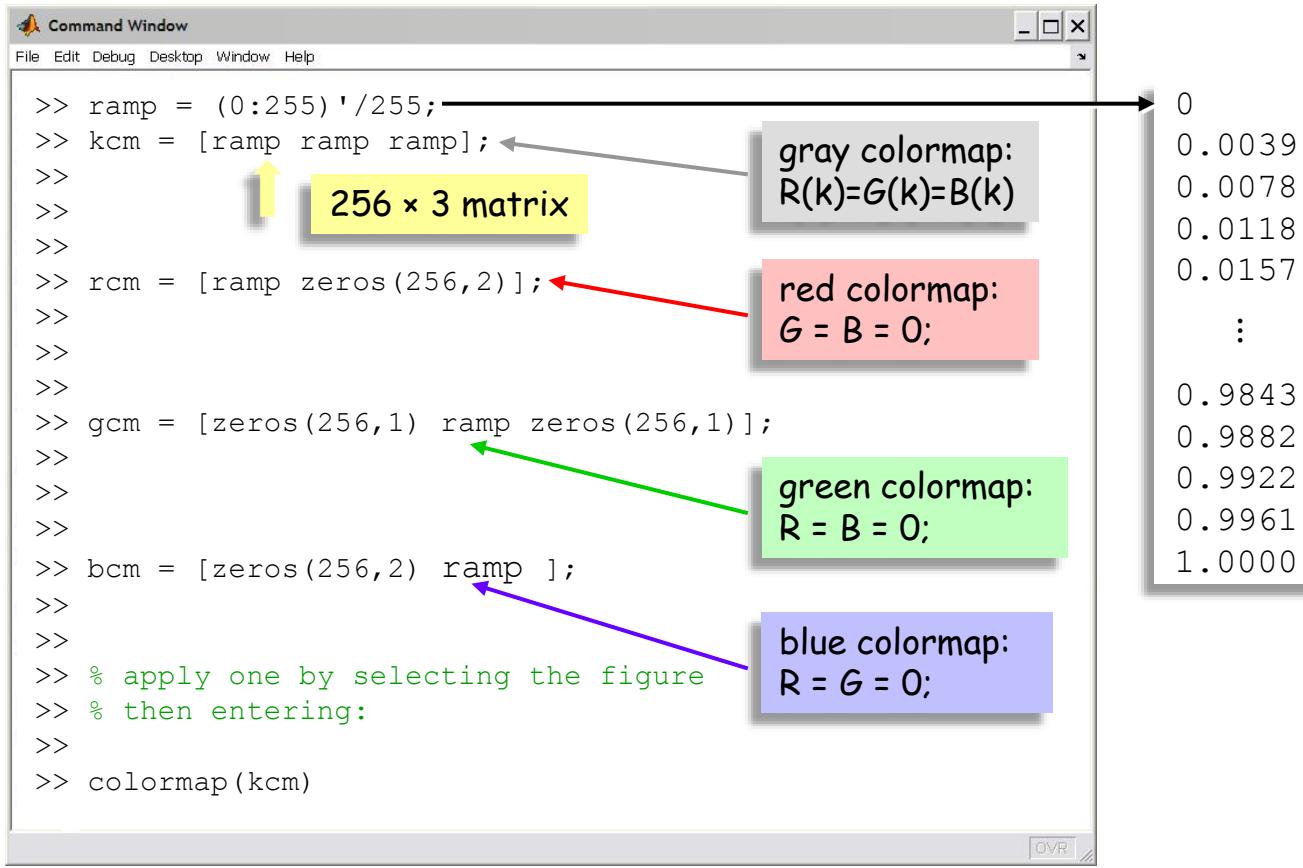
```
>> [M, cmap] = imread('button_mapped.bmp', 'bmp');  
>> T = uint8(reshape(cmap(M+1, :), [size(M) 3]) * 255);
```

The  $262,144 \times 3$  matrix of intensity values is reshaped into a  $512 \times 512 \times 3$  image of type double. The values are scaled to lie between 0 & 255 then converted to type uint8.

By concatenating M's columns, Matlab rearranges M into a  $262,144 \times 1$  list. Each number in the list (if it has 1 added to it) refers to a row of the colormap. Then,  $\text{cmap}(M+1, :)$  produces a  $262,144 \times 3$  matrix of intensity values of type double between 0 & 1.



# How to Make Colormaps



The image shows a MATLAB Command Window with the following code:

```
>> ramp = (0:255)'/255;
>> kcm = [ramp ramp ramp];
>>
>>
>> rcm = [ramp zeros(256,2)];
>>
>>
>> gcm = [zeros(256,1) ramp zeros(256,1)];
>>
>>
>> bcm = [zeros(256,2) ramp];
>>
>>
>> % apply one by selecting the figure
>> % then entering:
>>
>> colormap(kcm)
```

Annotations explain the code:

- A yellow arrow points to the first line of code: `ramp = (0:255)'/255;` with the text "256 x 3 matrix".
- A red arrow points to the line `rcm = [ramp zeros(256,2)];` with the text "red colormap:  $G = B = 0;$ ".
- A green arrow points to the line `gcm = [zeros(256,1) ramp zeros(256,1)];` with the text "green colormap:  $R = B = 0;$ ".
- A purple arrow points to the line `bcm = [zeros(256,2) ramp];` with the text "blue colormap:  $R = G = 0;$ ".
- An arrow points from the text "gray colormap:  $R(k)=G(k)=B(k)$ " to the right side of the window, where a vertical list of values is shown.

The right side of the window displays a vertical list of values:

0
0.0039
0.0078
0.0118
0.0157
:
0.9843
0.9882
0.9922
0.9961
1.0000

This code, `0:255`, generates a 1 row by 256 element vector of class double that contains numbers 0 through 255 inclusive.

This, `(0:255)'`, has the same contents and class but is a 256 row by 1 column vector. The apostrophe (' ) is the matrix transpose operator.



R, G, & B bands of a  
truecolor image displayed  
with grayscale colormaps

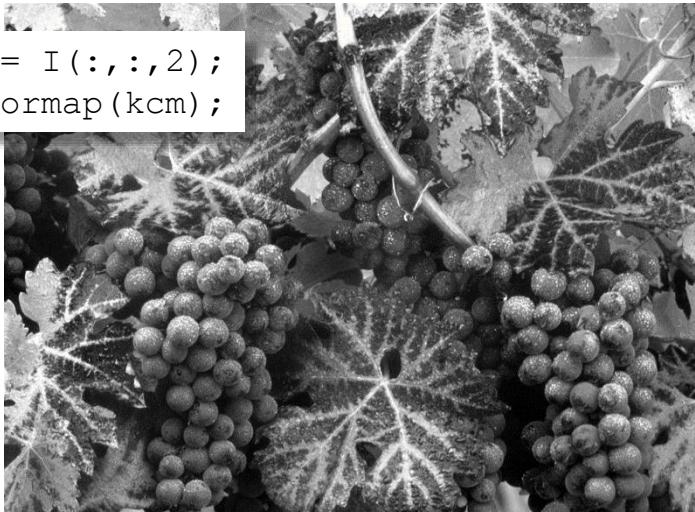
# EECE 4353 Image Processing

Vanderbilt University School of Engineering

```
>> I = imread('blue_grapes_sm.jpg', 'jpg');
```



```
>> Gn = I(:,:,2);  
>> colormap(kcm);
```



```
>> Rd = I(:,:,1);  
>> colormap(kcm);
```



```
>> Bl = I(:,:,3);  
>> colormap(kcm);
```





R, G, & B bands of a  
truecolor image displayed  
with grayscale colormaps

# EECE 4353 Image Processing

Vanderbilt University School of Engineering

>> I =



>> Gn =  
>> col



:, 1);  
kcm);



:, 3);  
kcm);





R, G, & B bands of a  
truecolor image displayed  
with tinted colormaps

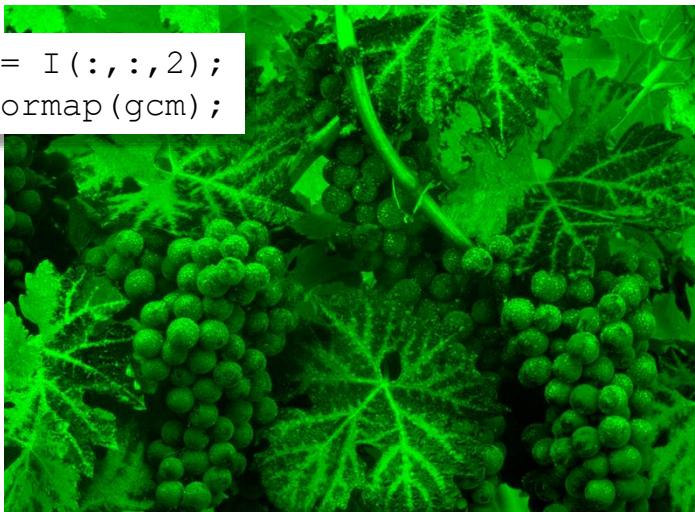
# EECE 4353 Image Processing

Vanderbilt University School of Engineering

```
>> I = imread('blue_grapes_sm.jpg', 'jpg');
```



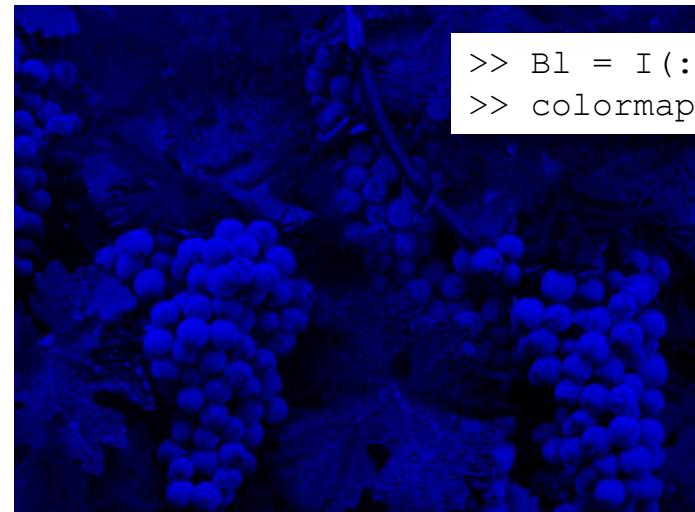
```
>> Gn = I(:,:,2);  
>> colormap(gcm);
```



```
>> Rd = I(:,:,1);  
>> colormap(rcm);
```



```
>> Bl = I(:,:,3);  
>> colormap(bcm);
```





R, G, & B bands of a  
truecolor image displayed  
with tinted colormaps

# EECE 4353 Image Processing

Vanderbilt University School of Engineering

>> I =



>> Gn =  
>> col



: , 1);  
rcm);

: , 3);  
bcm);





R, G, & B bands of a  
truecolor image displayed  
with grayscale colormaps

# EECE 4353 Image Processing

Vanderbilt University School of Engineering

```
>> I =
```



```
>> Gn =  
>> col
```



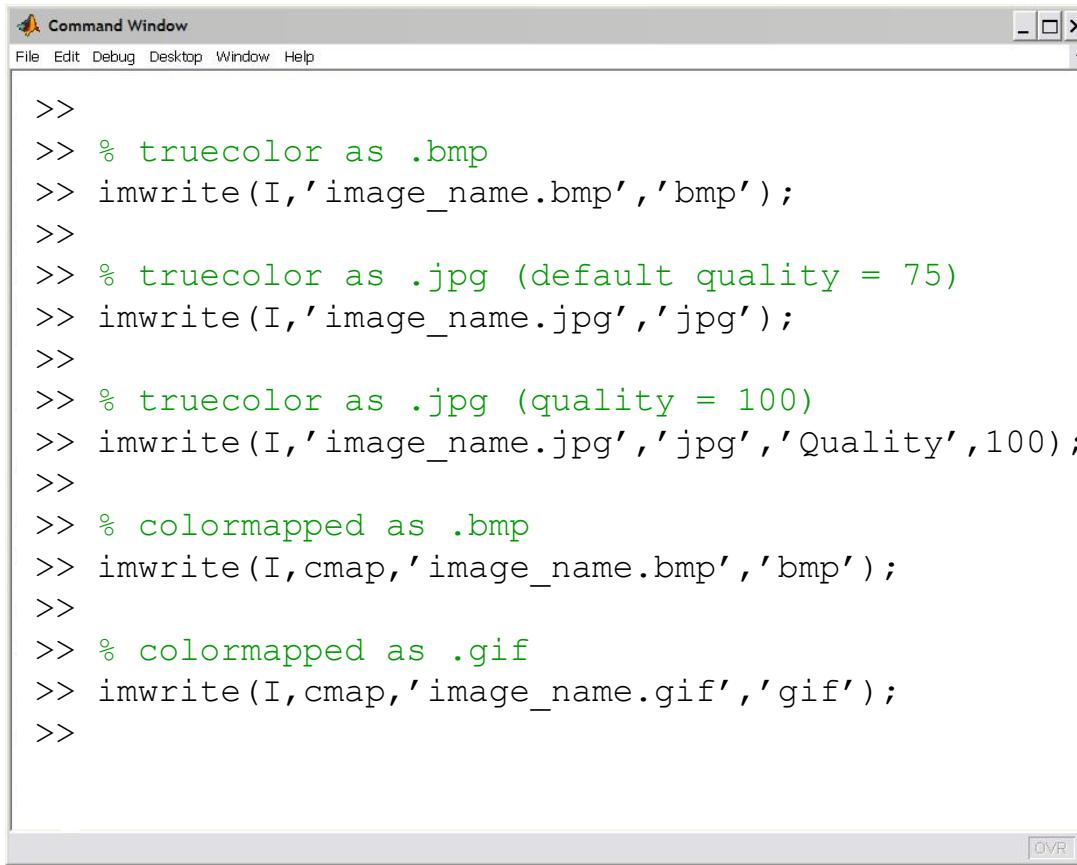
```
:, 1);  
kcm);
```



```
:, 3);  
kcm);
```



# Saving Images as Files



A screenshot of a MATLAB Command Window. The title bar says "Command Window". The menu bar includes "File", "Edit", "Debug", "Desktop", "Window", and "Help". The window contains the following MATLAB code:

```
>>
>> % truecolor as .bmp
>> imwrite(I,'image_name.bmp','bmp');
>>
>> % truecolor as .jpg (default quality = 75)
>> imwrite(I,'image_name.jpg','jpg');
>>
>> % truecolor as .jpg (quality = 100)
>> imwrite(I,'image_name.jpg','jpg','Quality',100);
>>
>> % colormapped as .bmp
>> imwrite(I,cmap,'image_name.bmp','bmp');
>>
>> % colormapped as .gif
>> imwrite(I,cmap,'image_name.gif','gif');
>>
```

Assuming that  
'I' contains the image of  
the correct class,  
that  
'cmap' is a colormap,  
and that  
'image\_name' is the  
file-name that you  
want.

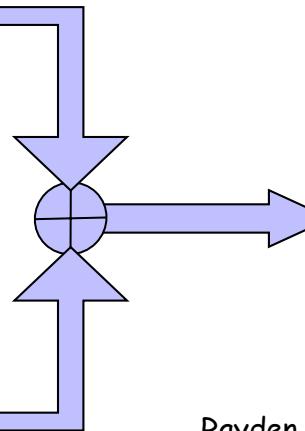


Jim Woodring - Bumperillo



Mark Rayden - The Ecstasy of Cecelia

## Double Exposure: Adding Two Images



Rayden Woodring - The Ecstasy of Bumperillo (?)



# Double Exposure: Adding Two Images

```
>> cd 'D:\Classes\EECE253\Fall 2006\Graphics\matlab intro'  
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');  
>> figure  
>> image(JW)  
>> truesize  
>> title('Bumperillo')  
>> xlabel('Jim Woodring')  
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');  
>> figure  
>> image(MR)  
>> truesize  
>> title('The Ecstasy of Cecelia')  
>> xlabel('Mark Ryden')  
>> [RMR,CMR,DMR] = size(MR);  
>> [RJW,CJW,DJW] = size(JW);  
>> rb = round((RJW-RMR)/2);  
>> cb = round((CJW-CMR)/2);  
>> JWplusMR = uint8((double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))+double(MR))/2);  
>> figure  
>> image(JWplusMR)  
>> truesize  
>> title('The Ecstasy of Bumperillo')  
>> xlabel('Jim Woodring + Mark Ryden')
```

Example  
Matlab Code



# Double Exposure: Adding Two Images

```
>> cd 'D:\Classes\EECE253\Fall 2006\Graphics\matlab intro'  
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');  
>> figure  
>> image(JW)  
>> truesize  
>> title('Bumperillo')  
>> xlabel('Jim Woodring')  
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');  
>> figure  
>> image(MR)  
>> truesize  
>> title('The Ecstasy of Cecelia')  
>> xlabel('Mark Ryden')  
>> [RMR,CMR,DMR] = size(MR);  
>> [RJW,CJW,DJW] = size(JW);  
>> rb = round((RJW-RMR)/2);  
>> cb = round((CJW-CMR)/2);  
>> JWplusMR = uint8((double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))+double(MR))/2);  
>> figure  
>> image(JWplusMR)  
>> truesize  
>> title('The Ecstasy of Bumperillo')  
>> xlabel('Jim Woodring + Mark Ryden')
```

Example  
Matlab Code

Cut a section out of the middle of the larger image the same size as the smaller image.



# Double Exposure: Adding Two Images

```
>> cd 'D:\Classes\EECE253\Fall 2006\Graphics\matlab intro'  
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');  
>> figure  
>> image(JW)  
>> truesize  
>> title('Bumperillo')  
>> xlabel('Jim Woodring')  
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');  
>> figure  
>> image(MR)  
>> truesize  
>> title('The Ecstasy of Cecelia')  
>> xlabel('Mark Ryden')  
>> [RMR,CMR,DMR] = size(MR);  
>> [RJW,CJW,DJW] = size(JW);  
>> rb = round((RJW-RMR)/2);  
>> cb = round((CJW-CMR)/2);  
>> JWplusMR = uint8((double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))+double(MR))/2);  
>> figure  
>> image(JWplusMR)  
>> truesize  
>> title('The Ecstasy of Bumperillo')  
>> xlabel('Jim Woodring + Mark Ryden')
```

Example  
Matlab Code

Note that the images are averaged,  
pixelwise.



# Double Exposure: Adding Two Images

```
>> cd 'D:\Classes\EECE253\Fall 2006\Graphics\matlab intro'  
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');  
>> figure  
>> image(JW)  
>> truesize  
>> title('Bumperillo')  
>> xlabel('Jim Woodring')  
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');  
>> figure  
>> image(MR)  
>> truesize  
>> title('The Ecstasy of Cecelia')  
>> xlabel('Mark Ryden')  
>> [RMR,CMR,DMR] = size(MR);  
>> [RJW,CJW,DJW] = size(JW);  
>> rb = round((RJW-RMR)/2);  
>> cb = round((CJW-CMR)/2);  
>> JWplusMR = uint8(double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))-double(MR))/2;  
>> figure  
>> image(JWplusMR)  
>> truesize  
>> title('The Ecstasy of Bumperillo')  
>> xlabel('Jim Woodring + Mark Ryden')
```

Example  
Matlab Code

Note the data class  
conversions.

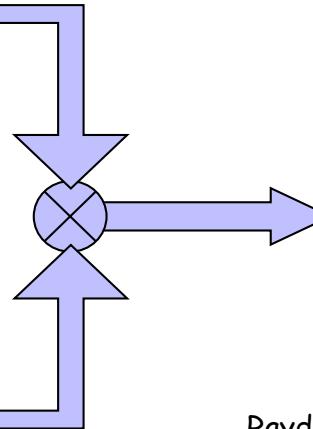


Jim Woodring - Bumperillo

## Intensity Masking: Multiplying Two Images



Mark Rayden - The Ecstasy of Cecelia



Rayden Woodring - Bumperillo Ecstasy (?)



# Intensity Masking: Multiplying Two Images

```
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');
>> [RMR,CMR,DMR] = size(MR);
>> [RJW,CJW,DJW] = size(JW);
>> rb = round((RJW-RMR)/2);
>> cb = round((CJW-CMR)/2);
>> JWplusMR = uint8((double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))+double(MR))/2);
>> figure
>> image(JWplusMR)
>> truesize
>> title('The Extacy of Bumperillo')
>> xlabel('Jim Woodring + Mark Ryden')
>> JWtimesMR = double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:)).*double(MR);
>> M = max(JWtimesMR(:));
>> m = min(JWtimesMR(:));
>> JWtimesMR = uint8(255*(double(JWtimesMR)-m)/(M-m));
>> figure
>> image(JWtimesMR)
>> truesize
>> title('EcstasyBumperillo')
```

Example  
Matlab Code



# Intensity Masking: Multiplying Two Images

```
>> JW = imread('Jim Woodring - Bumperillo.jpg','jpg');
>> MR = imread('Mark Ryden - The Ecstasy of Cecelia.jpg','jpg');
>> [RMR,CMR,DMR] = size(MR);
>> [RJW,CJW,DJW] = size(JW);
>> rb = round((RJW-RMR)/2);
>> cb = round((CJW-CMR)/2);
>> JWplusMR = uint8((double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:))+double(MR))/2);
>> figure
>> image(JWplusMR)
>> truesize
>> title('The Extacy of Bumperillo')
>> xlabel('Jim Woodring + Mark Ryden')
>> JWtimesMR = double(JW(rb:(rb+RMR-1),cb:(cb+CMR-1),:)).*double(MR);
>> M = max(JWtimesMR(:));
>> m = min(JWtimesMR(:));
>> JWtimesMR = uint8(255*(double(JWtimesMR)-m)/(M-m));
>> figure
>> image(JWtimesMR)
>> truesize
>> title('EcstasyBumperillo')
```

Example  
Matlab Code

Note that the images are multiplied, pixelwise.

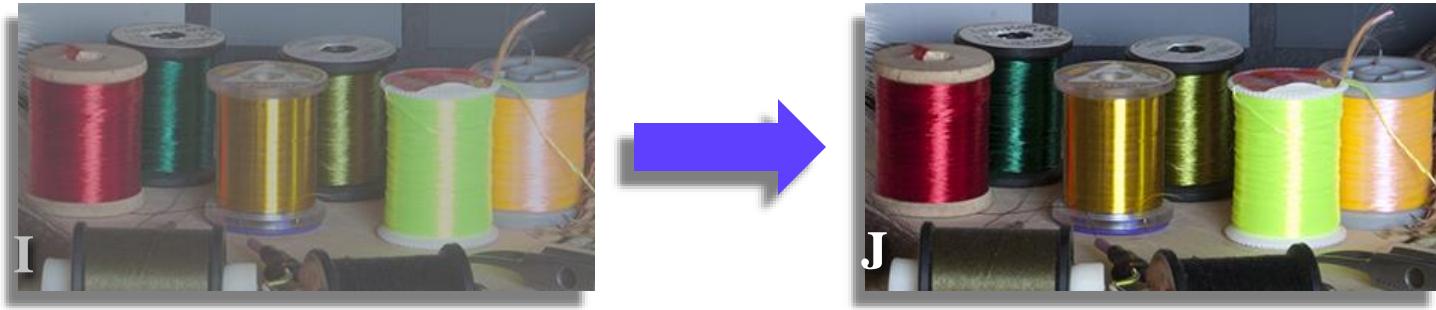
Note how the image intensities are scaled back into the range 0-255.



# A note on image intensity scaling

In the code on the previous slide we scaled the result of the image multiplication so that it would have the maximum possible dynamic range (0-255). The formula for such linear scaling is

$$m = \min(\mathbf{I}), \quad M = \max(\mathbf{I}), \quad \text{and} \quad J = 255(\mathbf{I} - m)/(M - m).$$



In Matlab, if the images are of type `uint8`, this requires class conversions:

```
J = uint8(255*double(I-m)/double(M-m));
```

Without the `double` casts, the arithmetic is performed with 8 bits of precision, which yields incorrect results.

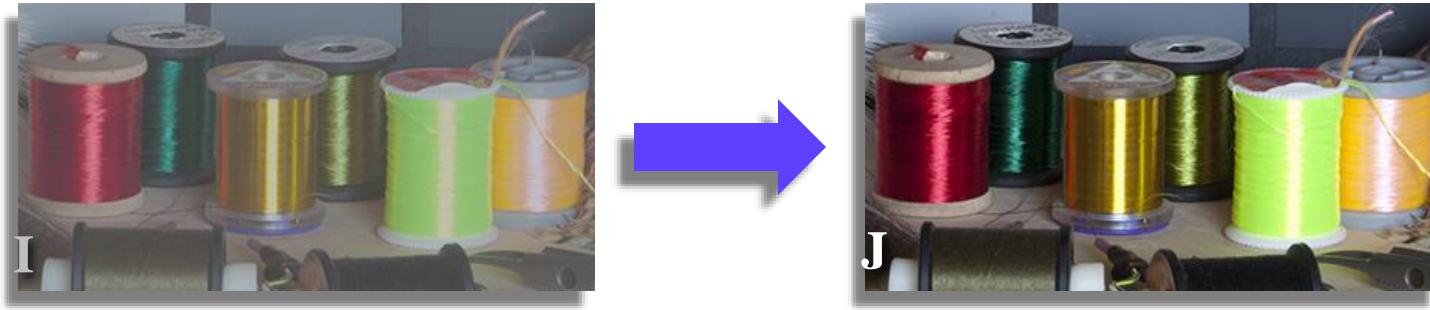


# A note on image intensity scaling

In the code on the previous slide we scaled the result of the image multiplication so that it would have the range  $[0, 255]$ . In Matlab, the min and max are computed like this:

$$m = \min(I(:)); M = \max(I(:));$$

$$m = \min(\mathbf{I}), M = \max(\mathbf{I}), \text{ and } J = 255(\mathbf{I} - m)/(M - m).$$



In Matlab, if the images are of type `uint8`, this requires class conversions:

$$J = \text{uint8}(255 * \text{double}(I - m) / \text{double}(M - m));$$

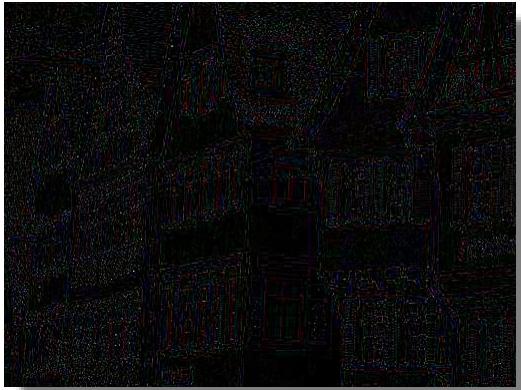
Without the `double` casts, the arithmetic is performed with 8 bits of precision, which yields incorrect results.



# More on image intensity scaling



Linear scaling is not always appropriate. If the image has some important features with very low intensity values and others that are large, the darker features may not be visible in a linearly scaled result. In such cases nonlinear scaling can perform better. The images below are differently scaled versions of  $D = \text{abs}(\mathbf{I} - \mathbf{J})$ ; where **I** is the image to the left and **J** was a quality-0 jpeg copy of itself.



`uint8(255 * (D-m) / (M-m))`



`L=log(D+1), uint8(255*L/max(L(:)))`



`uint8(D.^2)`



# Pixel Indexing in Matlab

“For” loops in Matlab are inefficient, whereas Matlab’s native indexing procedures are very fast.

Rather than

```
for r = 1:R
    for c = 1:C
        J(r,c,:) = IP_Function(I(r,c,:));
    end
end
```

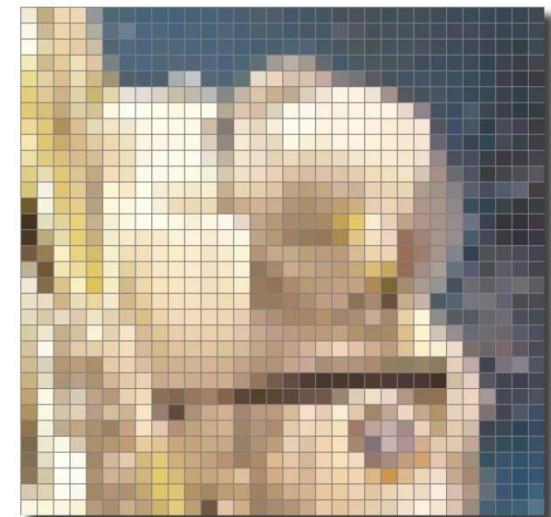
use, if possible

```
J = IP_Function(I);
```

But, sometimes that is not possible.

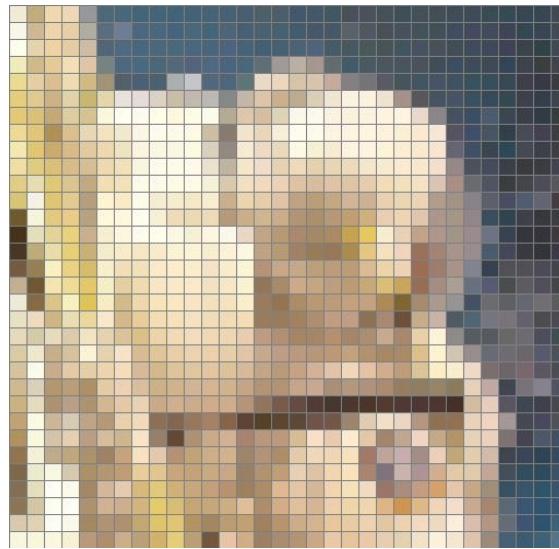
For example, if the output, **J**, is decimated with respect to the input, **I**, the above will not work (unless, of course, it is done within IP\_function).

“IP\_Function” is some arbitrary image processing function that you or someone else has written.





# Pixel Indexing in Matlab

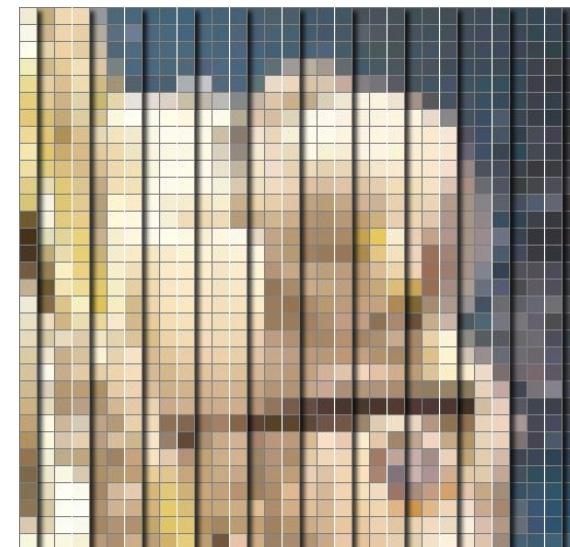


$r = 1:n:R;$

$I(r, :, :, :)$

$c = 1:n:C;$

$I(:, c, :, :)$



To decimate the above image by a factor of  $n$ , create a vector,  $\mathbf{r}$ , that contains the index of every  $n$ th row, and a similar vector,  $\mathbf{c}$ .

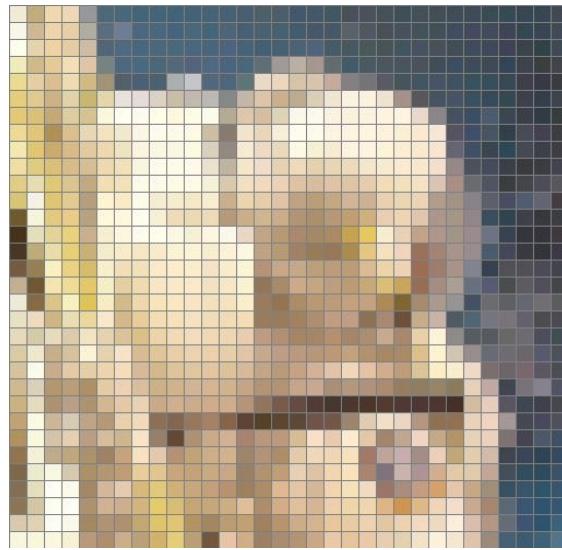
$r = [1 \ 4 \ 7 \ 10 \ 13 \ 16 \ 19 \ 22 \ 25 \ 28 \ 31]$

$c = [1 \ 4 \ 7 \ 10 \ 13 \ 16 \ 19 \ 22 \ 25 \ 28 \ 31]$



# Pixel Indexing in Matlab

This is called,  
'vectorizing'.

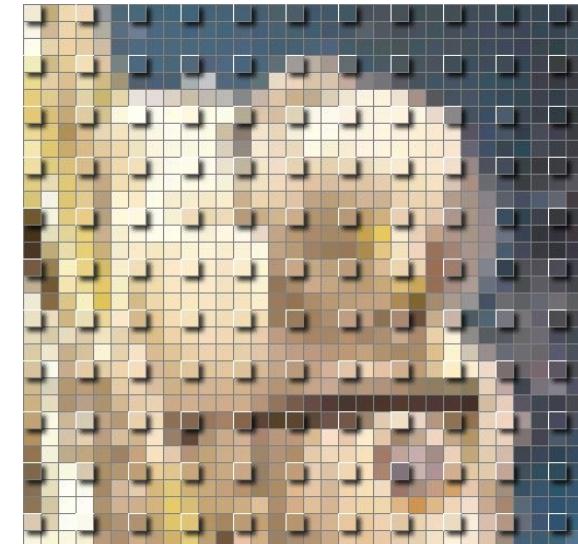
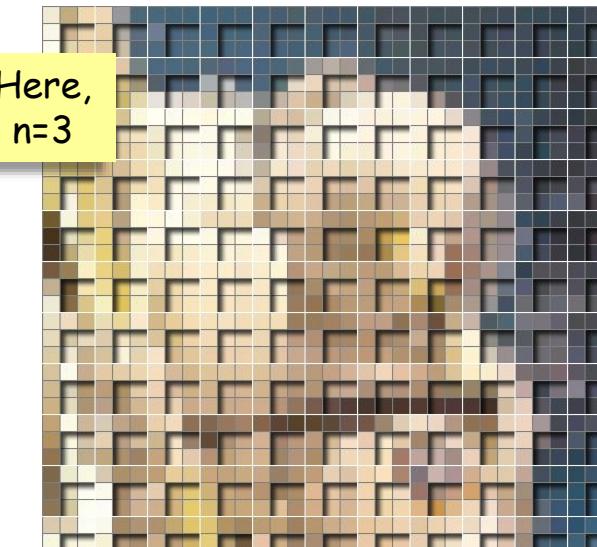


Take the  
pixels  
indexed  
by both  
 $r$  and  $c$ .



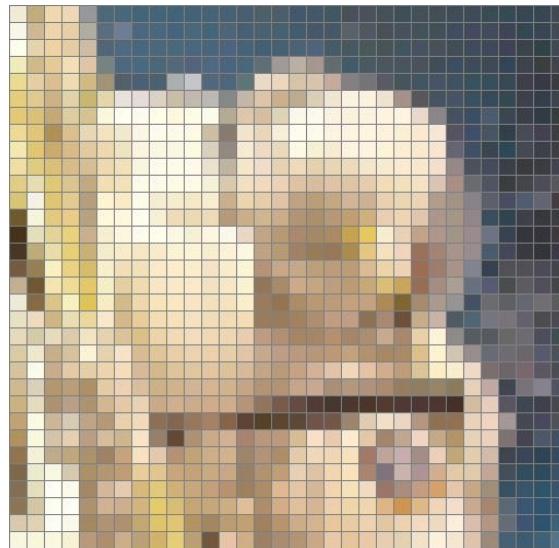
$I(r, c)$

Then, vectors **r** and **c** used as index arguments for image **I** select every  $n$ th column in every  $n$ th row.





# Pixel Indexing in Matlab

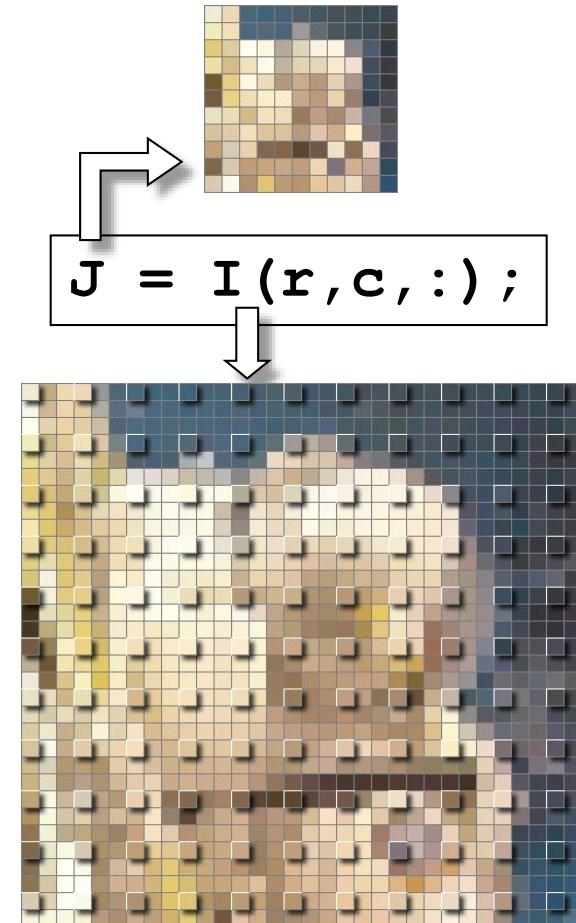


image,  $I$

$r = 1:n:R;$

$c = 1:n:C;$

Here,  
 $n=3$





# Pixel Indexing in Matlab

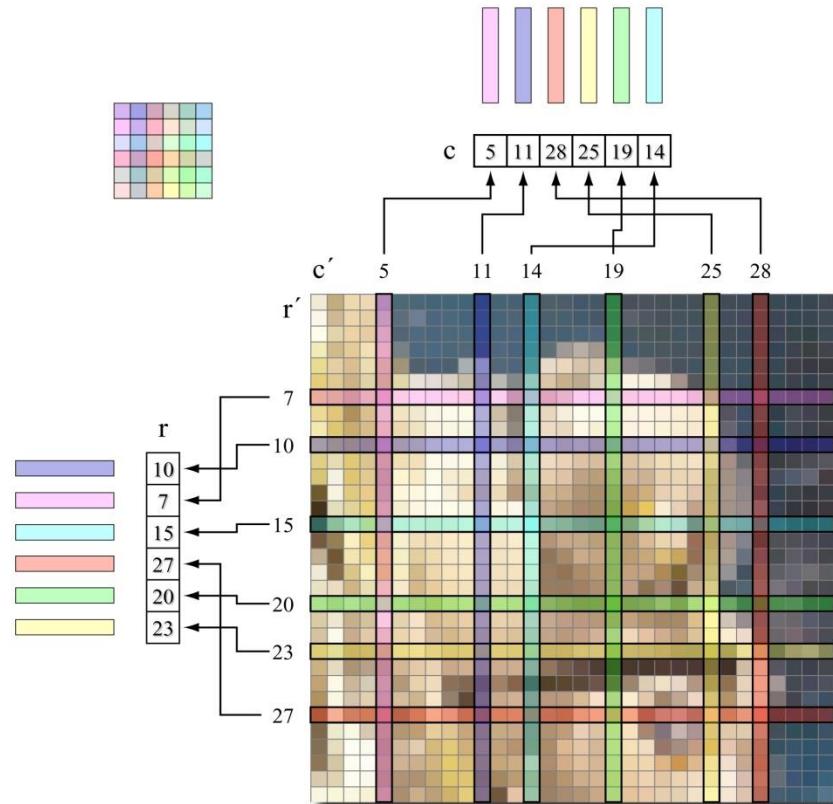
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r},\mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.





# Pixel Indexing in Matlab

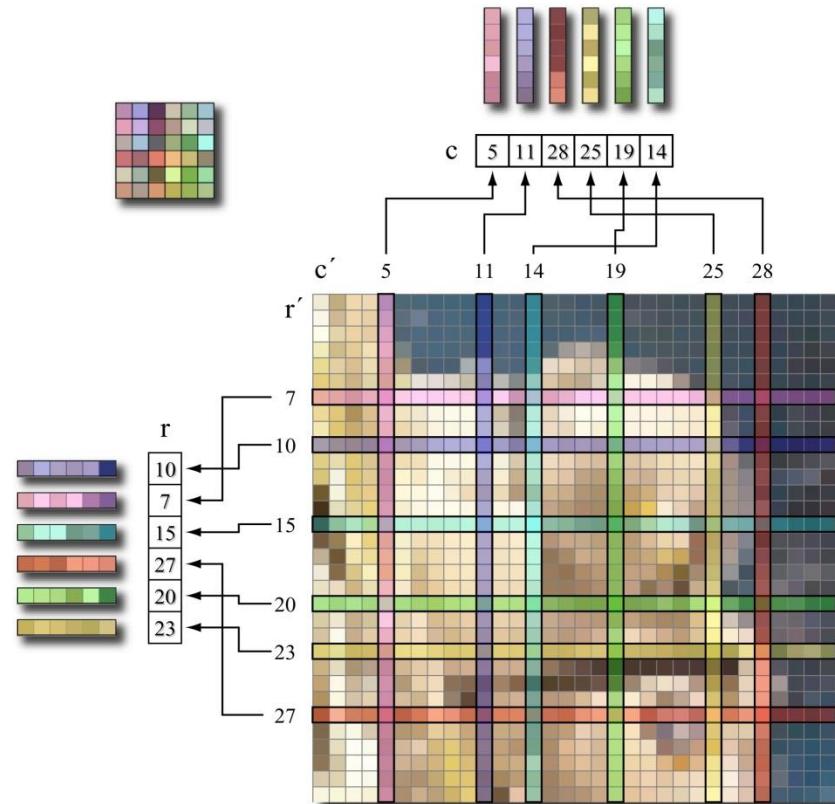
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r},\mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.





# Pixel Indexing in Matlab

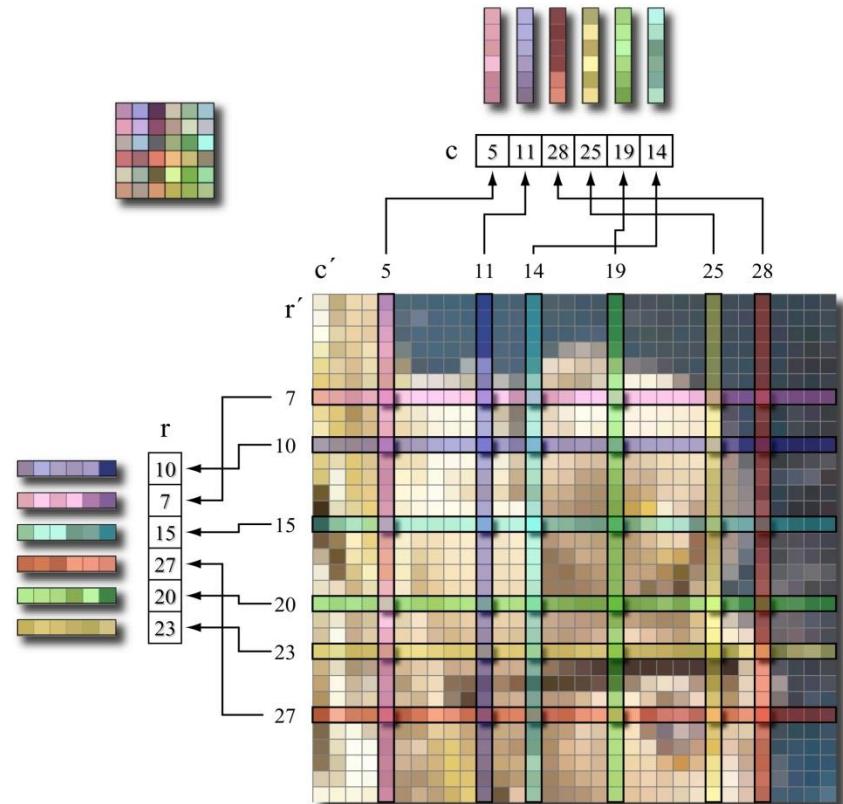
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r},\mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.





# Pixel Indexing in Matlab

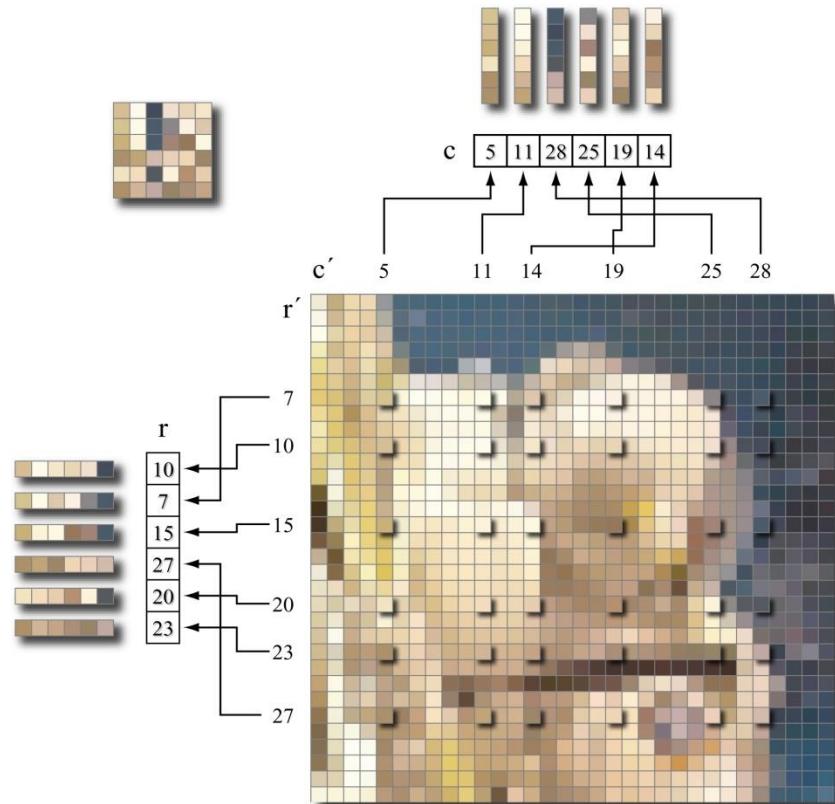
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r},\mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.





# Pixel Indexing in Matlab

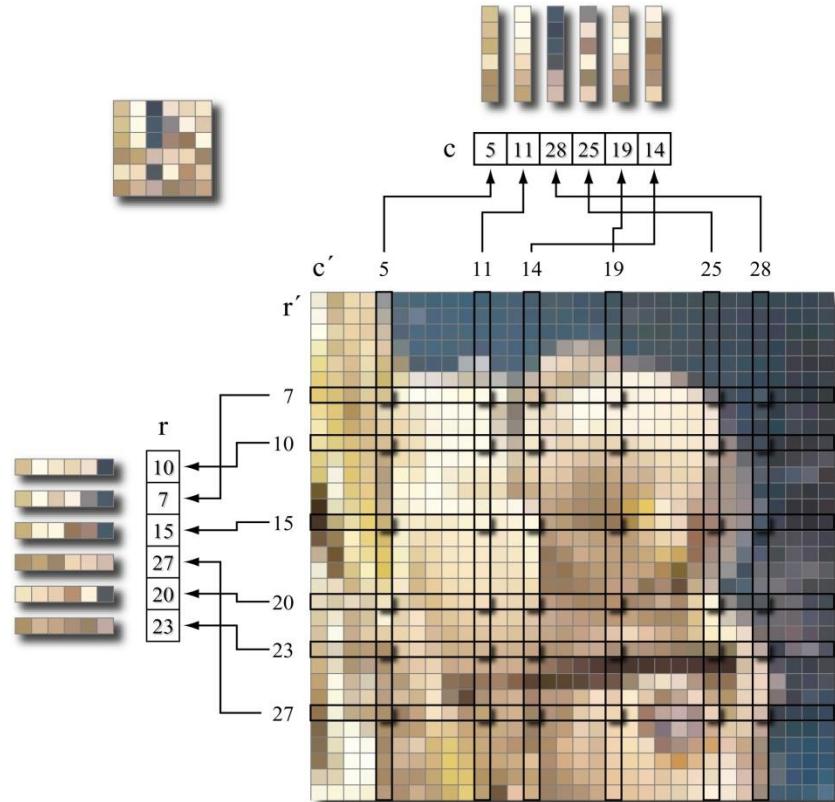
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r}, \mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.





# Pixel Indexing in Matlab

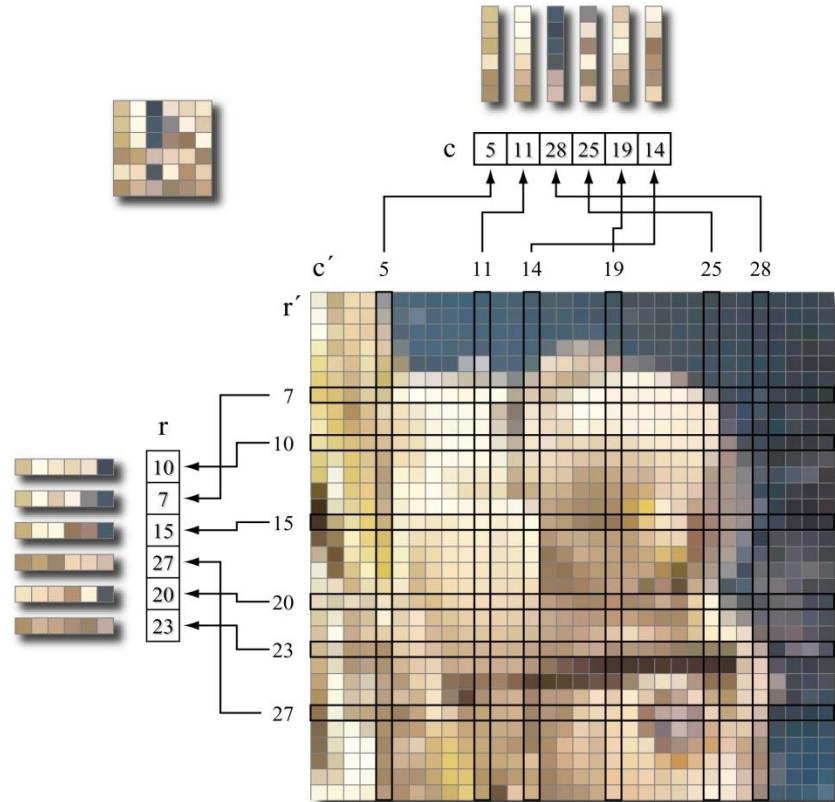
Indexing in Matlab is fully general.

If  $\mathbf{I}$  is  $R \times C \times B$ , vectors  $\mathbf{r}$  and  $\mathbf{c}$  can contain any numbers  $1 \leq r_k \leq R$  and  $1 \leq c_k \leq C$ .

The numbers can be in any order and can be repeated within  $\mathbf{r}$  and  $\mathbf{c}$ .

The result of  $\mathbf{I}(\mathbf{r},\mathbf{c})$  is an ordinal shuffling of the pixels from  $\mathbf{I}$  as indexed by  $\mathbf{r}$  and  $\mathbf{c}$ .

Whenever possible,  
avoid using 'for' loops;  
vectorize instead.

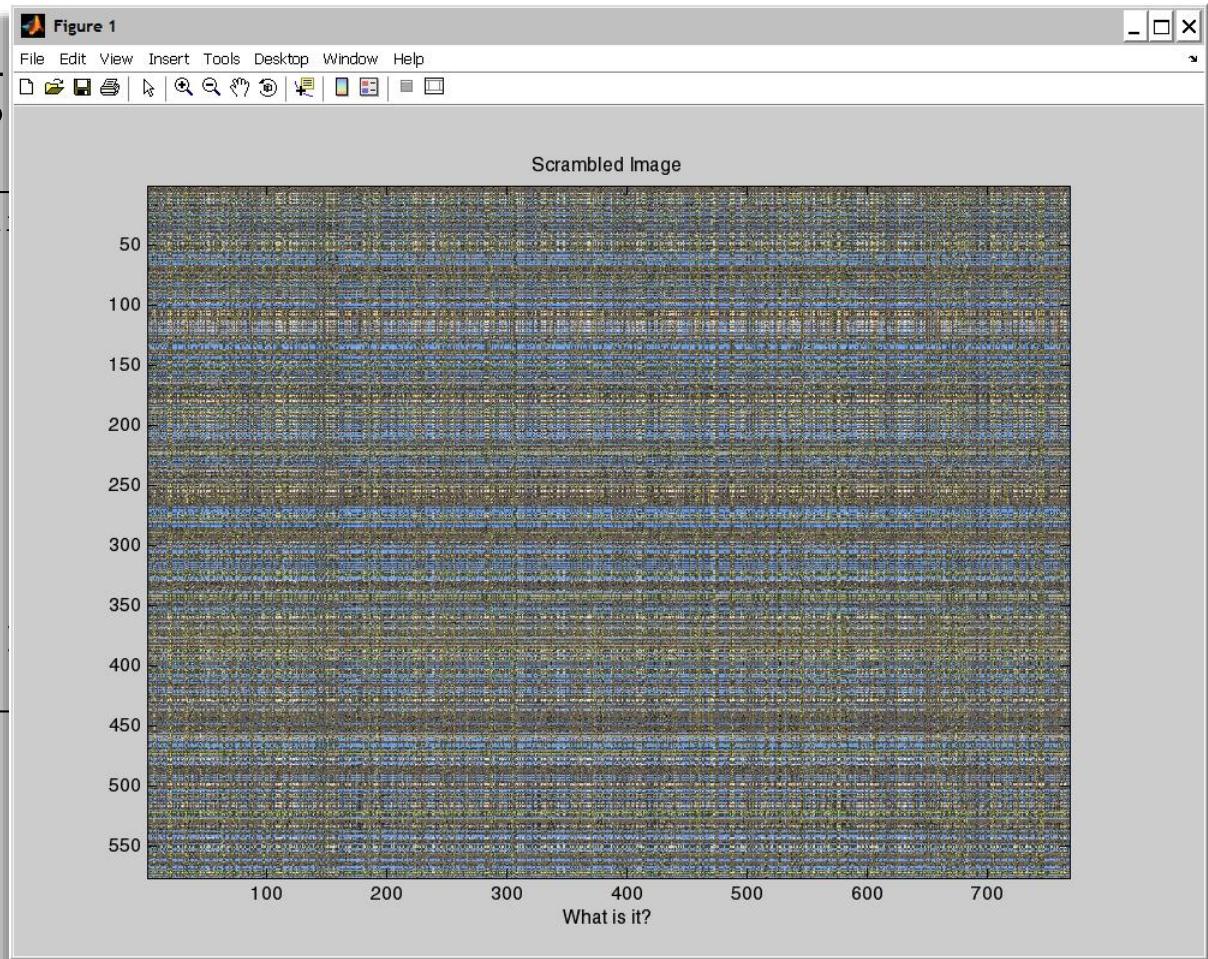




# Pixel Indexing

```
>> I = imread('Lawraa - Fl...  
>> size(I)  
ans =  
    576    768      3  
>> r = randperm(576);  
>> c = randperm(768);  
>> J = I(r,c,:);  
>> figure  
>> image(J)  
>> truesize  
>> title('Scrambled Image')  
>> xlabel('What is it?')
```

Fun (if you're an imaging geek) thing to try with Matlab indexing: Scramble an image!

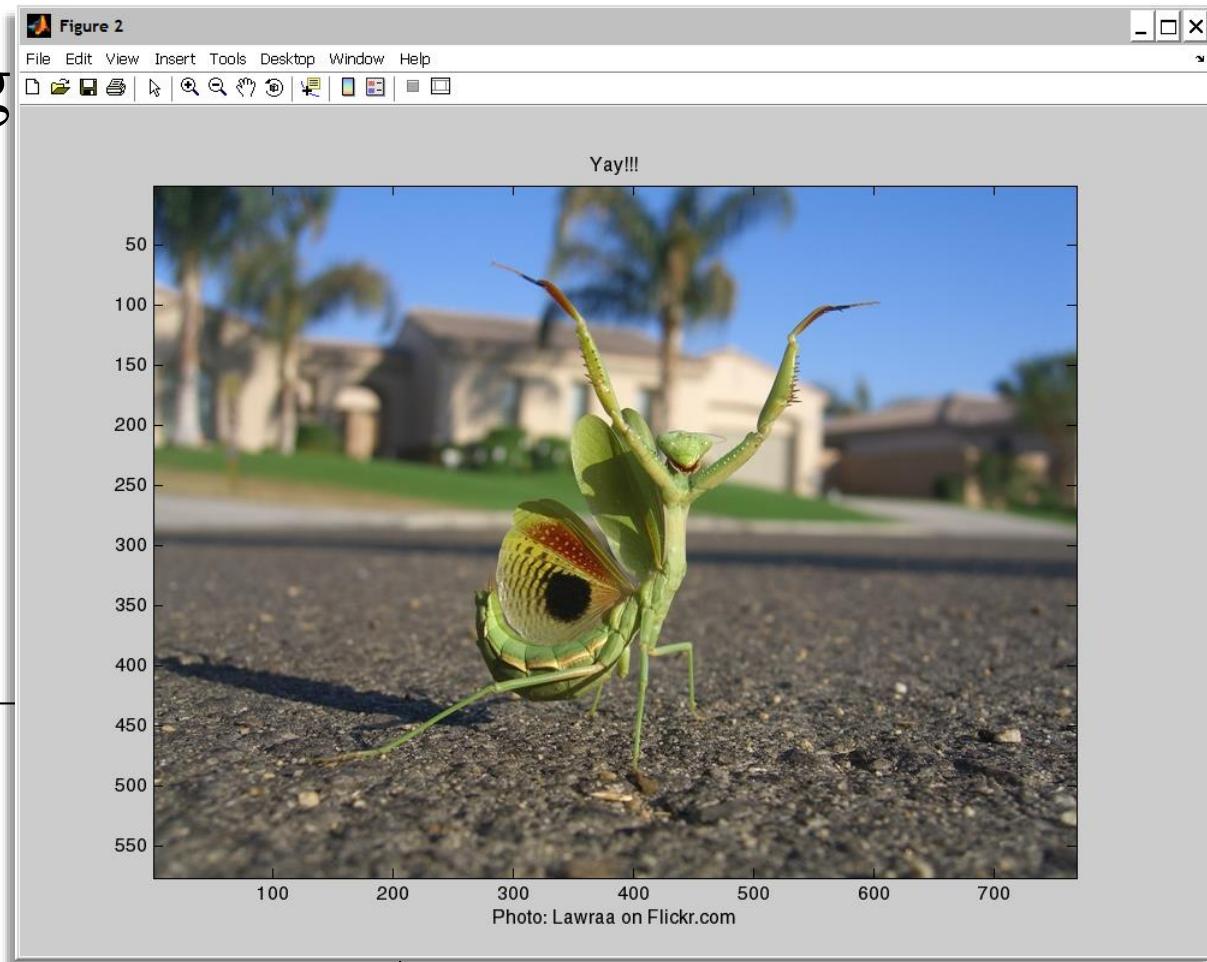




# Pixel Indexing

The image can be unscrambled using the row and column permutation vectors,  $r$  &  $c$ .

```
>> ...
>> xlabel('What is it?')
>> K(r,c,:)=J;
>> figure
>> image(K)
>> truesize
>> title('Yay!!!!')
>> xlabel('Photo: Lawraa on Flickr.com')
```





# Matlab Vectorization

Making a Matlab program run as fast as possible fast by avoiding loops is called “vectorization.” It is not something that Matlab does automatically. Rather it is a skill that a good programmer learns.

“tic” starts Matlab’s code timer. “toc” stops the timer and reports the elapsed time.

## Looped code

```
tic
for t = 1:1024
    y(t) = sin(2*pi*t/512);
end
toc
```

Elapsed time is 0.038923 seconds.

## Vectorized code

```
tic
y = sin(pi*(0:(1/256):2));
toc
```

Elapsed time is 0.002019 seconds.

$0.038923 / 0.002019 = 19.2784 \Rightarrow$  vectorized code is 20 times faster than the looped!

There is a web page on vectorization at the MathWorks:

[http://www.mathworks.com/help/matlab/matlab\\_prog/vectorization.html](http://www.mathworks.com/help/matlab/matlab_prog/vectorization.html)



# Matlab Vectorization

Example: compute a distance matrix, the matrix of squared pair-wise distances between two sets of vectors.

$$\mathbf{D}^2(\mathbf{A}, \mathbf{B}) = \begin{bmatrix} \|\mathbf{a}_1 - \mathbf{b}_1\|^2 & \cdots & \|\mathbf{a}_1 - \mathbf{b}_N\|^2 \\ \vdots & \ddots & \vdots \\ \|\mathbf{a}_1 - \mathbf{b}_N\|^2 & \cdots & \|\mathbf{a}_N - \mathbf{b}_N\|^2 \end{bmatrix}$$

$$\mathbf{A}_{M \times N} = [\mathbf{a}_1 \ \cdots \ \mathbf{a}_N], \quad \mathbf{B}_{M \times N} = [\mathbf{b}_1 \ \cdots \ \mathbf{b}_N],$$

$$\mathbf{a}_m, \mathbf{b}_n \in \mathbb{R}^M, \quad a_{m,i}, b_{n,j} \in \mathbb{R},$$

$$\mathbf{a}_m = \begin{bmatrix} a_{m,1} \\ \vdots \\ a_{m,M} \end{bmatrix}, \quad \mathbf{b}_n = \begin{bmatrix} b_{m,1} \\ \vdots \\ b_{m,M} \end{bmatrix}.$$

$$\mathbf{D}_{m,n}^2(\mathbf{A}, \mathbf{B})$$

$$= \|\mathbf{a}_m - \mathbf{b}_n\|^2$$

$$= \sum_{k=1}^M (a_{m,k} - b_{n,k})^2$$

$$= (a_{m,1}^2 + \cdots + a_{m,N}^2)$$

$$+ (b_{m,1}^2 + \cdots + b_{m,N}^2)$$

$$- 2(a_{m,1}b_{n,1} + \cdots + a_{m,N}b_{n,N})$$

$$= \|\mathbf{a}_m\|^2 + \|\mathbf{b}_n\|^2 - 2\mathbf{a}_m^T \mathbf{b}_n.$$



# Matlab Vectorization

Looped code

```
tic
for r = 1:1000
    for c = (r+1):1000
        diff = A(:,r)-B(:,c);
        D(r,c) = diff'*diff;
    end
end
D = D + D';
toc
```

Elapsed time is  
4.893803 seconds.

Example: compute  
a distance matrix.

`bsxfun(function,A,B)`  
applies the binary operation  
specified by `function`, to  
matrices A and B.

Vectorized code

```
tic
D = bsxfun(@plus,sum(A.^2)',bsxfun(@plus,sum(B.^2),-2*(A'*B)));
toc
```

Elapsed time is  
0.009223 seconds.

$4.893803 / 0.009223 = 530.609 \Rightarrow$  vectorized code is 530 times faster than the looped!